

Data Foundation

& The ML Framework

Understanding Data and How Machines Learn

Nipun Batra | IIT Gandhinagar

Recap: What Did We See Last Week?

AI is everywhere:

- AlphaFold solving protein folding
- AI writing code, music, art
- Self-driving cars, medical diagnosis

The key insight:

Machine Learning = Learning patterns from DATA

Today: What IS this data? How does learning actually work?

Today's Learning Goals

By the end of this lecture, you will be able to:

1. **Distinguish** traditional programming from machine learning
2. **Identify** the three learning paradigms (supervised, unsupervised, RL)
3. **Differentiate** classification from regression problems
4. **Explain** why train/test split is essential
5. **Apply** the sklearn API pattern to any ML problem

Part 1: The ML Framework

How Machines Learn

The Big Question

Every AI/ML system answers **one fundamental question**:

"Given some input, what should the output be?"


| Input | Output | System |
|----------------------------|------------------|------------------|
| Email text | Spam or Not Spam | Spam Filter |
| Image | "Cat" or "Dog" | Image Classifier |
| "The capital of France is" | "Paris" | Language Model |
| House features | Price (\$) | Price Predictor |

Traditional Programming vs ML

 center

The Paradigm Shift

Traditional Programming



Explicit programming

Machine Learning

Why This Matters

| Scenario | Traditional | ML |
|----------------------|-----------------------|---------------------------|
| Spam changes tactics | Rewrite rules | Retrain on new examples |
| 1000 new categories | 1000 new rule sets | Same algorithm, new data |
| Complex patterns | Impossible to specify | Model discovers them |
| Human bias | Encoded in rules | (Can still exist in data) |

ML excels when patterns are complex or rules are hard to specify explicitly.

Example: Spam Detection

Traditional Approach:

```
def is_spam(email):
    if "FREE" in email:
        return True
    if "winner" in email:
        return True
    if "click here" in email:
        return True
    if sender not in contacts:
        if num_links > 5:
            return True
    # ... 1000 more rules
    return False
```

What about "Fr33" or "w1nner"?

What if rules conflict?

The Power of Generalization

Training data has: "FREE money!!!"

Model learns: Unusual capitalization + exclamation marks + money words = suspicious

New email: "EARN ca\$h NOW!!!" → **Spam** (never seen before, but pattern matches)

Good ML models learn the **underlying pattern**, not just memorize examples.

Three Learning Paradigms

 center

Paradigm Comparison

| Aspect | Supervised | Unsupervised | Reinforcement |
|----------|-----------------|-------------------|------------------|
| Data | X + y (labels) | X only | States + Actions |
| Goal | Predict labels | Find structure | Maximize reward |
| Feedback | Correct answers | None | Reward signals |
| Example | Spam detection | Customer segments | Game playing |

Supervised Learning: The Teacher Analogy

SUPERVISED LEARNING

Student (Model)

"Is this spam?"

Teacher (Training Data)



"Yes, that's spam."
"No, that's legitimate."
"Yes, that's spam."
...

After many examples, student learns the patterns!

Supervised Learning: Two Types

 center

Classification: Discrete Outputs

Goal: Assign input to one of K categories

| Type | K | Example |
|-------------|----------|------------------------------|
| Binary | 2 | Spam / Not Spam |
| Multi-class | $K > 2$ | Cat, Dog, Bird, Fish |
| Multi-label | Multiple | [Action, Comedy] for a movie |

```
# Binary: One probability  
model.predict_proba(email) # → [0.15, 0.85] = 85% spam
```

```
# Multi-class: K probabilities (sum to 1)  
model.predict_proba(image) # → [0.70, 0.20, 0.05, 0.05]  
#          Cat    Dog   Bird   Fish
```

Classification Examples

| Task | Input | Classes | Real-World Use |
|-----------------------|--------------------|--------------------------------|--------------------------|
| Spam Detection | Email text | Spam, Not Spam | Gmail, Outlook |
| Medical Diagnosis | Symptoms, tests | Disease A, B, Healthy | Hospital systems |
| Image Recognition | Photo pixels | 1000 ImageNet classes | Google Photos |
| Sentiment Analysis | Review text | Positive, Negative, Neutral | Brand monitoring |
| Fraud Detection | Transaction | Fraud, Legitimate | Credit card companies |
| Face Recognition | Face image | Person 1, 2, ..., N | iPhone unlock |

Regression: Continuous Outputs

Goal: Predict a numerical value

```
# Output can be ANY number
model.predict(house_features)    # → 425,000.00
model.predict(face_image)         # → 27.3 (years old)
model.predict(stock_data)         # → 152.47 (price)
```

Classification: "Which bucket?" | Regression: "How much?"

Regression Examples

| Task | Input | Output | Range | Real-World Use |
|--------------------|-----------------|------------|----------------|------------------|
| House Pricing | Size, location | Price (\$) | \$100K - \$10M | Zillow, Redfin |
| Age Estimation | Face image | Years | 0 - 100 | Age verification |
| Demand Forecasting | History, season | Units | 0 - ∞ | Amazon inventory |
| Energy Prediction | Weather, time | kWh | 0 - ∞ | Power grid |
| Stock Prediction | Historical data | Price | 0 - ∞ | Trading |
| Salary Estimation | Resume features | Salary | \$0 - \$1M | LinkedIn |

Quick Check: Classification or Regression?

| Task | Answer |
|------------------------------------|--------------------------------------|
| "Will it rain tomorrow?" | Classification (Yes/No) |
| "How many mm of rain?" | Regression (continuous) |
| "What genre is this movie?" | Classification (Action, Comedy, ...) |
| "What rating will user give?" | Could be both! (1-5 stars) |
| "Which digit is written?" | Classification (0-9) |
| "How confident is the prediction?" | Regression (0.0 - 1.0) |

The In-Between: Ordinal Data

Some data is **ordered categories**:

| Rating | As Classification | As Regression |
|--------|-------------------|---------------|
| ★ | Class 0 | 1.0 |
| ★★ | Class 1 | 2.0 |
| ★★★ | Class 2 | 3.0 |
| ★★★★ | Class 3 | 4.0 |
| ★★★★★ | Class 4 | 5.0 |

Both approaches can work! Regression might predict 3.7 stars.

ML Tasks: The Decision Flowchart

 center

Computer Vision Task Hierarchy

 center

NLP Task Hierarchy

| Task | Input → Output | Example | Complexity |
|----------------|------------------------|---|------------|
| Classification | Text → Category | "Great movie!" → Positive | ★ |
| NER | Text → Tagged entities | "[Sundar Pichai]_PERSON visited [Google]_ORG" | ★★ |
| Seq2Seq | Sequence → Sequence | English → French | ★★★ |
| Generation | Prompt → Text | "Write a poem..." → Poem | ★★★★★ |

Part 2: Understanding Data

The Fuel for Machine Learning

"Data is the New Oil"

"Data is the new oil. Like oil, data is valuable, but if unrefined it cannot really be used."

— Clive Humby (2006)

| Oil Industry | ML Industry |
|----------------------|------------------------|
| Extract crude oil | Collect raw data |
| Refine into gasoline | Clean and process data |
| Powers engines | Powers models |
| Pollution issues | Bias issues |

What IS Data in ML?

 center

Anatomy of a Dataset

Features (X)

| | sqft | beds | baths | garage | year |
|---------|------|------|-------|--------|------|
| Row 1 → | 1500 | 3 | 2 | Yes | 1990 |
| Row 2 → | 2000 | 4 | 3 | Yes | 2005 |
| Row 3 → | 1200 | 2 | 1 | No | 1975 |
| Row 4 → | 1800 | 3 | 2 | Yes | 2010 |

Label (y)

| price |
|---------|
| 300,000 |
| 450,000 |
| 200,000 |
| 350,000 |

n_samples = 4 (rows)

n_features = 5 (columns in X)

Features: The Inputs

Features = Information about each example

Good Features

- Relevant to prediction
- Measurable/computable
- Available at prediction time
- Not too correlated with each other

Feature Examples

| Domain | Features |
|--------|----------------------|
| House | sqft, beds, zip code |
| Email | word counts, sender |

Types of Features

 center

Feature Type Details

| Type | Values | Example | Encoding |
|-------------|---------------|-------------------------|------------------|
| Numerical | Any number | Age: 25, Price: \$50.99 | Use directly |
| Categorical | Unordered set | Color: Red, Blue, Green | One-hot encoding |
| Binary | 2 values | Has garage: Yes/No | 0 or 1 |
| Ordinal | Ordered set | Size: S < M < L < XL | Integer encoding |
| Text | String | "Great product!" | Embedding |
| Date/Time | Timestamp | 2024-01-15 | Extract features |

One-Hot Encoding Example

Problem: Models need numbers, but "Red" isn't a number!

Solution: Create binary columns for each category

| Color | is_Red | is_Blue | is_Green |
|-------|--------|---------|----------|
| Red | 1 | 0 | 0 |
| Blue | 0 | 1 | 0 |
| Green | 0 | 0 | 1 |
| Blue | 0 | 1 | 0 |

```
pd.get_dummies(df['color']) # Does this automatically!
```

Labels: The Outputs

Label = What we want to predict

| Supervised Task | Label Type | Examples |
|-----------------------|--------------------|-------------------------|
| Binary classification | 0 or 1 | spam/not spam |
| Multi-class | Integer (0 to K-1) | digit (0-9) |
| Regression | Float | price (\$) |
| Multi-label | Binary vector | [action, comedy, drama] |

Unsupervised learning has NO labels! That's what makes it "unsupervised."

Types of Data Structures

 center

Tabular Data (Most Common)

```
import pandas as pd

df = pd.DataFrame({
    'sqft': [1500, 2000, 1200],
    'beds': [3, 4, 2],
    'price': [300000, 450000, 200000]
})

#      sqft  beds   price
# 0    1500      3  300000
# 1    2000      4  450000
# 2    1200      2  200000

X = df[['sqft', 'beds']] # Features
y = df['price']          # Labels
```

Image Data

```
import numpy as np

# Grayscale image: Height × Width
mnist_digit = np.zeros((28, 28)) # 784 pixels

# Color image: Height × Width × 3 (RGB)
photo = np.zeros((224, 224, 3)) # 150,528 values

# Batch of images: Batch × Height × Width × Channels
batch = np.zeros((32, 224, 224, 3)) # 32 images
```

| | | | | | |
|-----|---|-----|-----|-----|---|
| 0 | 0 | 0 | 23 | 155 | 0 |
| 0 | 0 | 89 | 254 | 254 | 0 |
| 0 | 0 | 155 | 254 | 178 | 0 |
| ... | | | | | |

← Each cell = pixel brightness
(0 = black, 255 = white)

Text Data


Raw text needs preprocessing:

```
text = "I love this movie! It's great."  
  
# Step 1: Tokenize (split into words/pieces)  
tokens = ["I", "love", "this", "movie", "!", "It", "'s", "great", "."]  
  
# Step 2: Convert to numbers (vocabulary index)  
indices = [23, 156, 45, 892, 2, 56, 78, 234, 3]  
  
# Step 3: (Optional) Convert to embeddings  
embeddings = model.embed(indices) # Shape: (9, 768)
```

Time Series Data

```
# Stock prices over time
dates = pd.date_range('2024-01-01', periods=100)
prices = [100.0, 101.5, 99.8, 102.3, ...] # 100 values

# Key characteristic: ORDER MATTERS!
# Shuffling destroys the patterns
```



How Much Data Do You Need?

 center

The Data Scaling Laws

| Model Complexity | Minimum Data | Sweet Spot | Diminishing Returns |
|------------------------|--------------|------------|---------------------|
| Linear Regression | 50 | 500 | 5,000 |
| Decision Tree | 100 | 1,000 | 10,000 |
| Random Forest | 500 | 5,000 | 50,000 |
| Neural Network (small) | 1,000 | 10,000 | 100,000 |
| Deep Learning | 10,000 | 100,000 | 1,000,000+ |
| LLMs | 1B tokens | 1T tokens | 10T+ tokens |

More data almost always helps, but there are diminishing returns. Quality > Quantity!

Data Quality Issues



Missing Values

```
# Original data
df = pd.DataFrame({
    'age': [25, None, 35, 42, None],
    'salary': [50000, 60000, None, 80000, 55000]
})


# Option 1: Drop rows with missing values
df.dropna() # Lose 3 rows!

# Option 2: Fill with mean/median
df['age'].fillna(df['age'].mean()) # Fill with 34

# Option 3: Fill with mode (categorical)
df['color'].fillna(df['color'].mode()[0])
```


Outliers

Normal distribution:



$$\mu = 50$$

With outlier:



$$\mu = 50$$

$$\text{outlier} = 500$$

Mean without outlier: 50

Mean with outlier: 95 ← Heavily skewed!

Detection: Z-score > 3, IQR method, visual inspection

Class Imbalance

Scenario: Fraud detection (1% fraud, 99% legitimate)

| Approach | Model Prediction | Accuracy |
|-------------|------------------|----------|
| Naive model | "All legitimate" | 99%! |
| Smart model | Tries to detect | 97% |

The 99% model is USELESS! It never detects fraud.

Solutions:

- Oversample minority class (SMOTE)
- Undersample majority class
- Use class weights
- Different metrics (precision, recall, F1)

The Data Lifecycle



Data Lifecycle in Code

```
# 1. COLLECT
df = pd.read_csv('raw_data.csv')

# 2. CLEAN
df = df.dropna() # Handle missing
df = df[df['age'] < 120] # Remove outliers

# 3. EXPLORE
df.describe() # Statistics
df.hist() # Visualize

# 4. TRANSFORM
df['log_price'] = np.log(df['price']) # Transform
X = pd.get_dummies(df[features]) # Encode

# 5. SPLIT
X_train, X_test, y_train, y_test = train_test_split(X, y)

# 6. USE
model.fit(X_train, y_train)
```

Part 3: Train/Test Split

The Most Important Concept

The Exam Analogy

 center

Two Study Strategies

Strategy A: Memorize

Q: "What is $2+3$?"

A: "5" (memorized)

Q: "What is $2+4$?"

A: "???" (never seen!)

Result: Fails on new questions

Strategy B: Learn

Q: "What is $2+3$?"

A: "5" (understands addition)

Q: "What is $2+4$?"

A: "6" (applies principle)

What is Overfitting?

 center

Overfitting in Detail

Definition: Model performs well on training data but poorly on new data

| Metric | Overfitting | Good Fit |
|-------------------|-------------|----------|
| Training Accuracy | 99% | 92% |
| Test Accuracy | 60% | 90% |
| Gap | 39% | 2% |

The model memorized the training data!

The Train/Test Split

 center

Why Split Works

THE RULES

1. Model trains ONLY on training data
2. Model NEVER sees test data during training
3. After training, evaluate on test data
4. Test performance = Expected real-world performance

The Golden Rule

NEVER PEEK AT TEST DATA!

If you use test data for:

- Choosing which model to use → **Data leakage**
- Tuning hyperparameters → **Data leakage**
- Feature selection → **Data leakage**

Your accuracy estimate will be **too optimistic** and your model will **fail in production**.

Train/Test Split in Code

```
from sklearn.model_selection import train_test_split

# The sacred split
X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.2,          # 20% for testing
    random_state=42,        # For reproducibility
    stratify=y              # Keep class proportions (classification)
)

print(f"Training samples: {len(X_train)}") # 800
print(f"Test samples: {len(X_test)}")      # 200

# NOW: Only touch X_train, y_train until final evaluation
```

Choosing Split Ratio

| Dataset Size | Train % | Test % | Reasoning |
|------------------|---------|--------|--------------------------|
| < 1,000 | 70% | 30% | Need enough test samples |
| 1,000 - 10,000 | 80% | 20% | Standard split |
| 10,000 - 100,000 | 90% | 10% | Plenty of test data |
| > 100,000 | 95% | 5% | Even 5% is thousands |

With huge datasets, even a small percentage gives reliable estimates.

The Three-Way Split

For model selection, you need **three** sets:



center

Three-Way Split Explained

| Set | Purpose | When Used |
|------------------|----------------------|---------------------------------|
| Training (60%) | Learn parameters | During <code>model.fit()</code> |
| Validation (20%) | Tune hyperparameters | Choosing model, settings |
| Test (20%) | Final evaluation | Once, at the very end |

```
# First split: separate test
X_temp, X_test, y_temp, y_test = train_test_split(X, y, test_size=0.2)

# Second split: separate validation
X_train, X_val, y_train, y_val = train_test_split(X_temp, y_temp, test_size=0.25)
# 0.25 of 0.8 = 0.2 of total
```

Common Mistakes

| Mistake | Why It's Bad | Fix |
|----------------------|--------------------------|----------------------------------|
| Training on all data | Can't detect overfitting | Always split first |
| Peeking at test data | Optimistic estimates | Lock away test data |
| Tuning on test | Indirect training | Use validation set |
| Small test set | High variance | Use at least 20% |
| No random state | Non-reproducible | Set <code>random_state=42</code> |
| Data leakage | False confidence | Check processing order |

Part 4: The ML Recipe

Putting It All Together

The Universal ML Recipe

 center

Step 1: Get Data

```
import pandas as pd
from sklearn.datasets import load_iris

# Option 1: Load from file
df = pd.read_csv('houses.csv')
X = df[['sqft', 'beds', 'baths']]
y = df['price']

# Option 2: Use sklearn datasets
iris = load_iris()
X, y = iris.data, iris.target

# Option 3: Create manually
X = [[1500, 3, 2], [2000, 4, 3], [1200, 2, 1]]
y = [300000, 450000, 200000]
```

Step 2: Prepare Data

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

# 1. Split first! (before any processing)
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# 2. Scale features (fit ONLY on train!)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train) # fit + transform
X_test_scaled = scaler.transform(X_test) # only transform!
```

Always fit scaler on training data only! Otherwise: data leakage.

Step 3: Choose a Model

 center

Model Complexity Ladder

| Model | Complexity | Interpretability | When to Use |
|-----------------|------------|------------------|----------------------------|
| Linear/Logistic | ★ | High | Start here, baseline |
| Decision Tree | ★★ | High | Need explanations |
| Random Forest | ★★★ | Medium | General purpose |
| XGBoost | ★★★★ | Low | Competitions |
| Neural Network | ★★★★★ | Very Low | Images, text, lots of data |

Always start simple! Only add complexity if needed.

Step 4: Train (Fit)

```
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeClassifier

# For regression
reg_model = LinearRegression()
reg_model.fit(X_train, y_train)

# For classification
clf_model = DecisionTreeClassifier(max_depth=5)
clf_model.fit(X_train, y_train)

# What happens inside:
# 1. Model sees (X, y) pairs
# 2. Adjusts internal parameters
# 3. Minimizes prediction error
# 4. Stores learned patterns
```

What Happens During Training?

Iteration 1: Predictions: [350K, 400K, 250K]
Actual: [300K, 450K, 200K]
Error: Large!
Action: Adjust parameters ↓

Iteration 2: Predictions: [320K, 430K, 220K]
Actual: [300K, 450K, 200K]
Error: Smaller
Action: Keep adjusting...

...

Iteration N: Predictions: [305K, 445K, 198K]
Actual: [300K, 450K, 200K]
Error: Small enough!
Action: Stop, save parameters

Step 5: Evaluate

```
# Make predictions on TEST data (never seen before!)
predictions = model.predict(X_test)

# Compare predictions to actual values
from sklearn.metrics import mean_squared_error, accuracy_score

# Regression
rmse = np.sqrt(mean_squared_error(y_test, predictions))
print(f"RMSE: ${rmse:.0f}") # RMSE: $25,000

# Classification
accuracy = accuracy_score(y_test, predictions)
print(f"Accuracy: {accuracy:.1%}") # Accuracy: 94.5%
```

Step 6: Deploy

```
import joblib

# Save the trained model
joblib.dump(model, 'house_price_model.pkl')
joblib.dump(scaler, 'scaler.pkl')

# Later, in production...
model = joblib.load('house_price_model.pkl')
scaler = joblib.load('scaler.pkl')

# New house comes in
new_house = [[1800, 3, 2]] # sqft, beds, baths
new_house_scaled = scaler.transform(new_house)
predicted_price = model.predict(new_house_scaled)
print(f"Predicted price: ${predicted_price[0]:,.0f}")
```

The sklearn API Pattern

 center

The Beauty of Consistent APIs

```
# ALL sklearn models follow the same pattern!

# Linear Regression
from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(X_train, y_train)
model.predict(X_test)

# Random Forest
from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier()
model.fit(X_train, y_train)
model.predict(X_test)

# Neural Network
from sklearn.neural_network import MLPClassifier
model = MLPClassifier()
model.fit(X_train, y_train)
model.predict(X_test)

# Same 3 methods: fit(), predict(), score()
```

Complete Example: Classification

```
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report

# 1. Load data
iris = load_iris()
X, y = iris.data, iris.target

# 2. Split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# 3. Train
model = DecisionTreeClassifier(max_depth=3)
model.fit(X_train, y_train)

# 4. Evaluate
y_pred = model.predict(X_test)
print(f"Accuracy: {accuracy_score(y_test, y_pred):.1%}")
print(classification_report(y_test, y_pred, target_names=iris.target_names))
```

Complete Example: Regression

```
from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
import numpy as np

# 1. Load data
housing = fetch_california_housing()
X, y = housing.data, housing.target

# 2. Split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# 3. Train
model = LinearRegression()
model.fit(X_train, y_train)

# 4. Evaluate
y_pred = model.predict(X_test)
print(f"RMSE: ${np.sqrt(mean_squared_error(y_test, y_pred))*100000:.0f}")
print(f"R2 Score: {r2_score(y_test, y_pred):.3f}")
```

Part 5: Course Roadmap

What's Coming Next

Your Journey Through AI

| Week | Topic | Big Question | You'll Learn |
|------|---------------------|-------------------------|--------------------------|
| 1 | Introduction | What can AI do? | Motivation, capabilities |
| 2 | Data Foundation | What is ML? | Framework, data, split |
| 3 | Supervised Learning | How do algorithms work? | LR, Trees, KNN |
| 4 | Model Selection | How to choose? | CV, tuning, ensembles |
| 5 | Neural Networks | What is deep learning? | Backprop, PyTorch |
| 6 | Computer Vision | How do machines see? | CNNs, YOLO |
| 7 | Language Models | How do LLMs work? | Transformers |
| 8 | Generative AI | How do machines create? | Diffusion, APIs |

Skills You'll Build

 center

Key Takeaways

Framework

1. ML learns from DATA
2. Three paradigms: Supervised, Unsupervised, RL
3. Classification vs Regression

Data

4. Features (X) + Labels (y)
5. Quality > Quantity
6. Train/Test split is SACRED

Practice

7. The sklearn pattern: `fit() → predict() → score()`

Ready to Build!

Next: Supervised Learning Deep Dive

Lab this week: Your first ML models with sklearn

"In God we trust. All others must bring data."

— W. Edwards Deming

Questions?