

# **From Language Model to Assistant**

## **The Complete Journey**

### **SFT, RLHF, and Modern AI Systems**

**Nipun Batra | IIT Gandhinagar**

# The Journey Complete

| Week | Topic               | The Question                      |
|------|---------------------|-----------------------------------|
| 1-2  | Foundations         | How do machines learn?            |
| 3-4  | Supervised Learning | How do we predict?                |
| 5    | Neural Networks     | What makes deep learning special? |
| 6    | Computer Vision     | How do machines see?              |
| 7    | Language Models     | How do machines understand text?  |
| 8    | LM → Assistant      | <b>How do we make it helpful?</b> |

# Last Week's Cliffhanger

We learned:

- LLMs predict the next token
- Transformers enable long-range attention
- Temperature controls creativity

**But we left with a puzzle:**

A model that predicts text well is NOT the same as a helpful assistant!

# The ChatGPT Problem

| What We Have                            | What We Want                              |
|---|---|
| Great text completion                   | Following instructions                    |
| "Paris is the capital of..." → "France" | "What's the capital of France?" → "Paris" |
| Random poem continuation                | "Write a poem about..."                   |
| Any code that fits                      | Working, safe code                        |

**Today: How do we bridge this gap?**

# Today's Agenda

1. **The Problem** - Why base models aren't enough
2. **SFT** - Teaching models to follow instructions
3. **RLHF** - Learning from human preferences
4. **DPO** - A simpler alternative
5. **The Full Pipeline** - Pre-training → SFT → Alignment
6. **Image Generation** - Quick overview
7. **AI Ethics & Future** - Responsible AI

# **Part 1: The Base Model Problem**

**Why Prediction Isn't Enough**

# Base Model Behavior

A model trained on web text learns to complete web text:

Prompt: "How do I make a cake?"

Base model might respond:

How do I make a cake?

I'm looking for a simple chocolate cake recipe. My mom used to make one but I lost the recipe card. Any help would be appreciated!

Posted by CakeLover92 on Reddit, March 2019

It's completing forum posts, not answering questions!

# More Base Model Problems

| Prompt                    | Base Model Response    | What We Want |
|---------------------------|------------------------|--------------|
| "Write Python code to..." | # TODO: implement this | Working code |
| "Is this email spam?"     | Continues the email    | Yes/No       |
| "Summarize this article"  | Writes more article    | Summary      |

Base models are trained to PREDICT, not to HELP.

# The Training Data Problem

**Base model training data:** The entire internet

- Wikipedia, Reddit, Stack Overflow
- Books, news articles, papers
- Code repositories, forums, social media

**What's missing:**

- Explicit instruction-following examples
- Feedback on what makes a "good" response
- Alignment with human values

## **Part 2: Supervised Fine-Tuning (SFT)**

**Teaching Models to Follow Instructions**

# SFT: The Key Idea

Collect examples of good instruction-following:

Instruction: "What is the capital of France?"

Response: "The capital of France is Paris."

Instruction: "Write a haiku about rain"

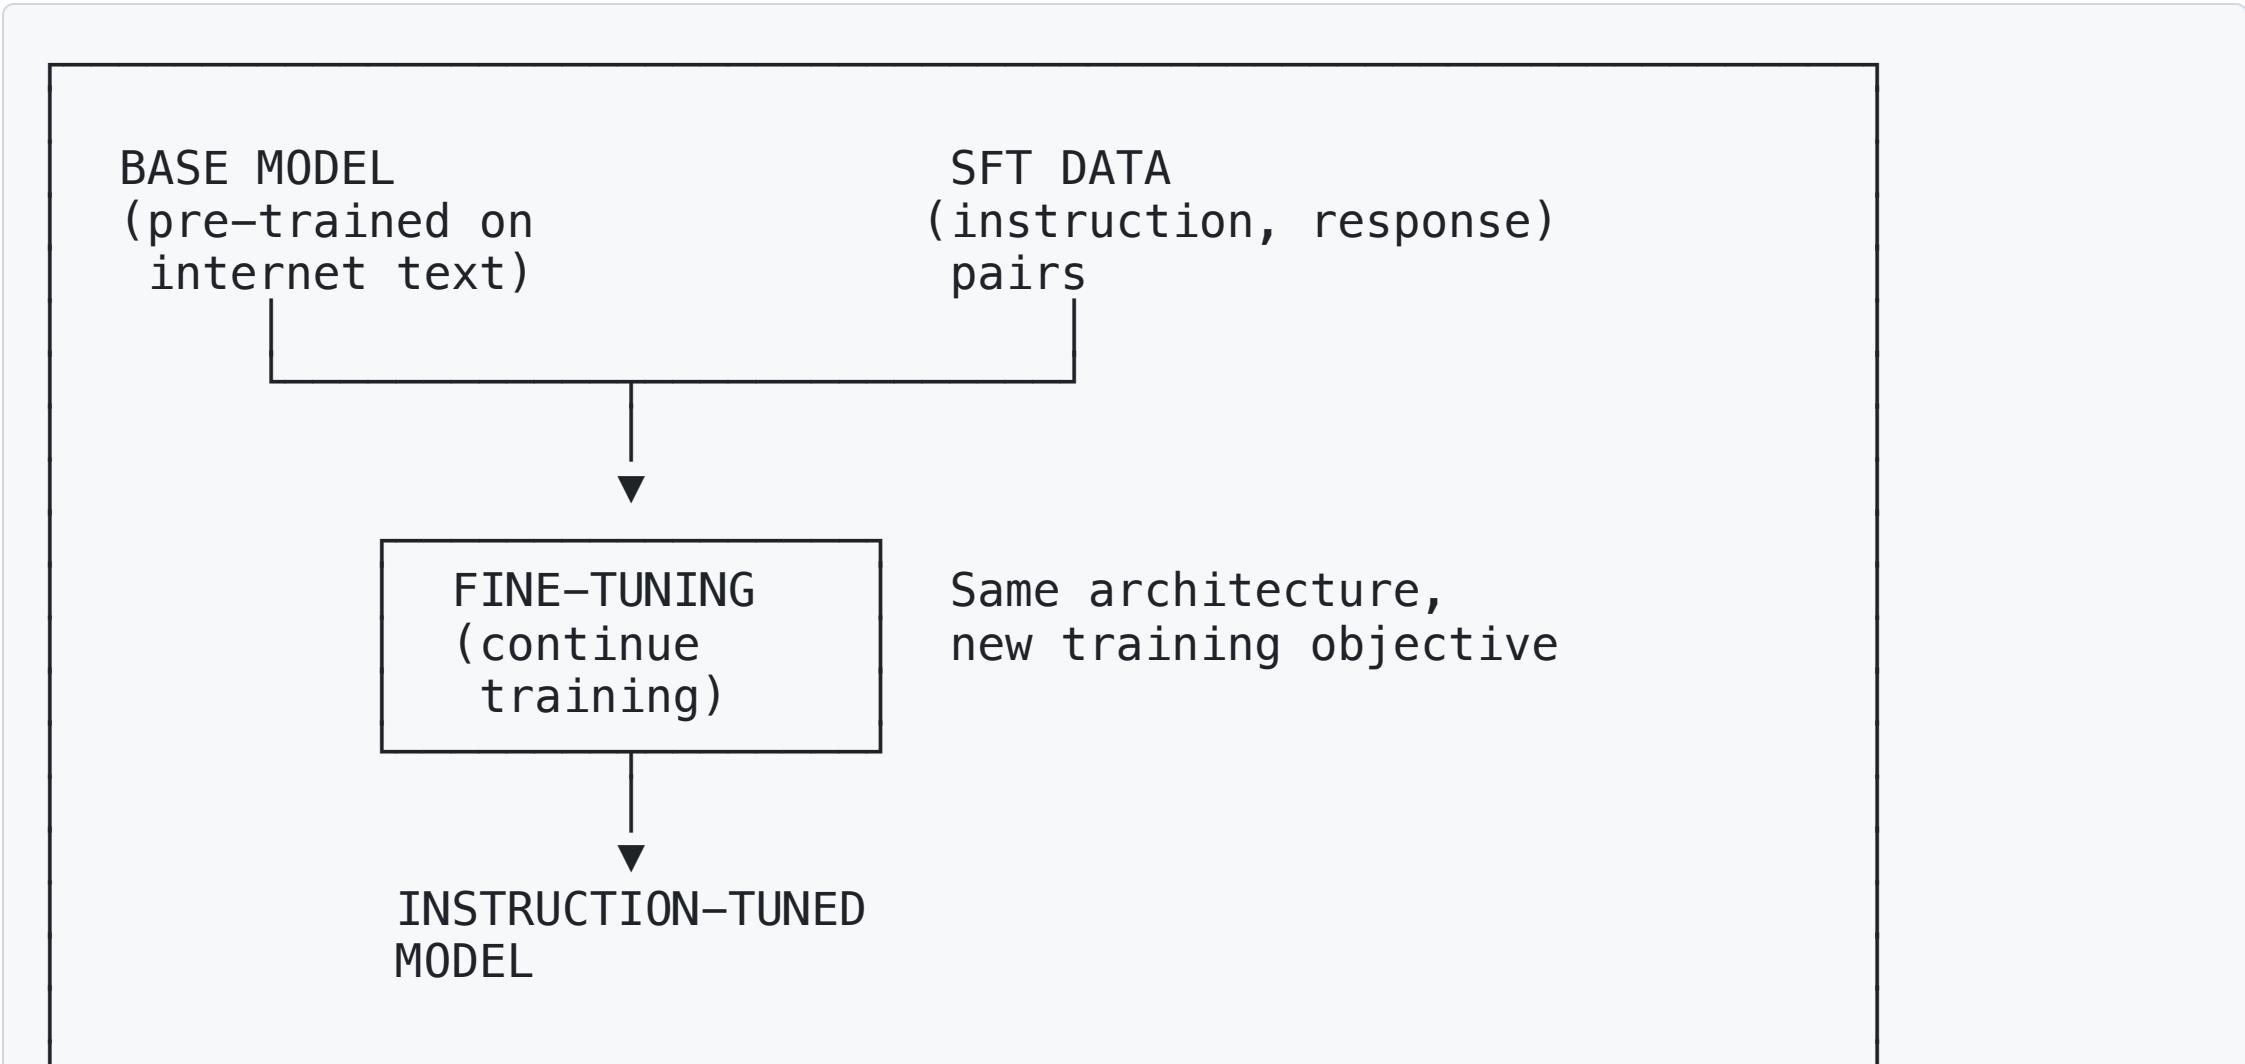
Response: "Gentle drops descend  
Washing away dusty streets  
Earth drinks and is whole"

Instruction: "Translate 'hello' to Spanish"

Response: "Hola"

Then fine-tune the base model on these examples!

# SFT Training Process



# Where Does SFT Data Come From?

| Source             | Method                            | Examples                   |
|--------------------|-----------------------------------|----------------------------|
| Human annotators   | Write ideal responses             | Expensive, high quality    |
| Existing datasets  | Convert Q&A datasets              | FLAN, Natural Instructions |
| Model distillation | GPT-4 generates for smaller model | Alpaca, Vicuna             |
| User feedback      | Real conversations                | ChatGPT usage data         |

# Famous SFT Datasets

| Dataset     | Size         | Method                            |
|-------------|--------------|-----------------------------------|
| FLAN        | 1,800+ tasks | Compiled from NLP datasets        |
| InstructGPT | ~100K        | Human annotators                  |
| Alpaca      | 52K          | GPT-3 generated                   |
| ShareGPT    | 90K+         | User-shared ChatGPT conversations |
| Dolly       | 15K          | Databricks employees              |

# SFT Training Example

```
from transformers import AutoModelForCausalLM, Trainer

# Load base model
model = AutoModelForCausalLM.from_pretrained("meta-llama/Llama-2-7b")

# Prepare SFT data
# Format: "### Instruction:\n{instruction}\n### Response:\n{response}"
train_data = load_dataset("sft_data.json")

# Fine-tune
trainer = Trainer(
    model=model,
    train_dataset=train_data,
    args=TrainingArguments(
        output_dir="../sft_model",
        num_train_epochs=3,
        per_device_train_batch_size=4,
    )
)
trainer.train()
```

# SFT Results

| Model                | Before SFT | After SFT |
|----------------------|------------|-----------|
| Follows instructions | 20%        | 85%       |
| Appropriate format   | 15%        | 90%       |
| Helpful responses    | 30%        | 75%       |

**SFT makes a HUGE difference!**

But there's still a problem...

# The Limits of SFT

SFT teaches **WHAT** to say, but not **WHICH** response is **BEST**.

**Prompt:** "Write a story about a dog"

**Response A:** (boring but correct)

**Response B:** (creative and engaging)

**Response C:** (grammatically perfect but dull)

**SFT can't distinguish between acceptable responses!**

We need a way to learn from **preferences**.

## **Part 3: RLHF**

### **Learning from Human Preferences**

# RLHF: The Key Insight

Instead of "this is the right answer"...

**Let humans rank responses from best to worst!**

Prompt: "Explain quantum physics simply"

Response A: [Technical jargon, hard to follow]

Response B: [Clear analogy with everyday objects]

Response C: [Accurate but dry explanation]

Human ranking: B > C > A

**Train the model to generate responses like B!**

# RLHF: Three Steps

## STEP 1: Collect Comparisons

---

Human ranks model responses: "A is better than B"

## STEP 2: Train Reward Model

---

Learn to predict human preferences  
reward(response) → score

## STEP 3: Optimize with RL

---

Fine-tune LLM to maximize reward model score  
(using PPO or similar RL algorithm)

# Step 1: Collect Human Preferences

Annotators compare pairs of responses:

| Prompt              | Response A         | Response B       | Preference       |
|---------------------|--------------------|------------------|------------------|
| "What's 2+2?"       | "The answer is 4." | "4"              | A (more helpful) |
| "Tell a joke"       | [Long, funny joke] | [Short, unfunny] | A                |
| "Write code for..." | [Working code]     | [Buggy code]     | A                |

Need thousands of comparisons!

## Step 2: Train Reward Model

**Input:** Prompt + Response

**Output:** Scalar score (how good is this?)

```
class RewardModel(nn.Module):
    def __init__(self, base_model):
        self.base = base_model
        self.head = nn.Linear(hidden_size, 1)

    def forward(self, prompt, response):
        hidden = self.base(prompt + response)
        score = self.head(hidden[:, -1]) # Score from last token
        return score
```

**Training objective:** reward(preferred) > reward(rejected)

## Step 3: RL Fine-tuning

Use PPO (Proximal Policy Optimization) to maximize reward:

```
# Simplified RLHF loop
for prompt in prompts:
    # Generate response from current policy
    response = model.generate(prompt)

    # Get reward
    reward = reward_model(prompt, response)

    # Update model to increase reward
    # (with KL penalty to stay close to SFT model)
    loss = -reward + beta * KL_divergence(model, sft_model)
    loss.backward()
    optimizer.step()
```

# Why the KL Penalty?

**Problem:** Without constraint, model might exploit reward model

Example reward hack:

Reward model likes "helpful" responses

- Model learns to add "I hope this helps!" to everything
- Gets high reward, but responses get worse

**Solution:** Stay close to the SFT model (KL divergence penalty)

# RLHF Results: InstructGPT

OpenAI's InstructGPT paper (2022):

| Metric           | Base GPT-3 | SFT | RLHF       |
|------------------|------------|-----|------------|
| Human preference | 22%        | 33% | <b>71%</b> |
| Truthfulness     | 34%        | 47% | <b>68%</b> |
| Less harmful     | 44%        | 61% | <b>84%</b> |

RLHF is what made GPT-3 → ChatGPT!

# RLHF Challenges

| Challenge      | Why It's Hard                  |
|----------------|--------------------------------|
| Expensive      | Need many human annotations    |
| Slow           | RL training is unstable        |
| Reward hacking | Model exploits reward model    |
| Alignment tax  | Sometimes hurts raw capability |

Is there a simpler alternative?

## **Part 4: DPO**

### **A Simpler Path to Alignment**

# DPO: Direct Preference Optimization

**Skip the reward model entirely!**

**Key insight:** We can derive a loss function directly from preferences

Traditional RLHF:

Preferences → Reward Model → RL Training → Aligned Model

DPO:

Preferences → Direct Fine-tuning → Aligned Model

# DPO Loss Function

For each comparison (prompt, preferred response, rejected response):

$$\mathcal{L} = -\log \sigma \left( \beta \log \frac{\pi(y_w|x)}{\pi_{ref}(y_w|x)} - \beta \log \frac{\pi(y_l|x)}{\pi_{ref}(y_l|x)} \right)$$

In simple terms:

- Increase probability of preferred responses
- Decrease probability of rejected responses
- Stay close to reference model

# DPO in Practice

```
from trl import DPOTrainer

# Prepare preference data
# Format: {"prompt": ..., "chosen": ..., "rejected": ...}
train_data = load_preferences_dataset()

# Simple training!
trainer = DPOTrainer(
    model=sft_model,
    ref_model=sft_model, # Reference for KL penalty
    train_dataset=train_data,
    beta=0.1, # Temperature parameter
)
trainer.train()
```

Much simpler than RLHF!

# DPO vs RLHF

| Aspect             | RLHF            | DPO                 |
|--------------------|-----------------|---------------------|
| Reward model       | Required        | Not needed          |
| Training stability | Unstable (RL)   | Stable (supervised) |
| Compute            | 3 models needed | 1 model + reference |
| Hyperparameters    | Many (PPO)      | Few (just $\beta$ ) |
| Performance        | Strong          | Comparable          |

DPO is becoming the preferred method!

## **Part 5: The Complete Pipeline**

**Pre-training → SFT → Alignment**

# The Full Journey

## STAGE 1: Pre-training

---

Data: Trillions of tokens (web, books, code)

Objective: Predict next token

Compute: Thousands of GPUs, months

Result: "Base model" – knows language, facts

↓

## STAGE 2: Supervised Fine-Tuning (SFT)

---

Data: ~100K instruction-response pairs

Objective: Learn to follow instructions

Compute: Hours to days

Result: "Instruction model" – follows commands

↓

## STAGE 3: Alignment (RLHF or DPO)

---

Data: Human preference comparisons

Objective: Be helpful, harmless, honest

Compute: Days

Result: "AI Assistant" – ChatGPT, Claude, etc.

# Real World Examples

| Model            | Pre-training    | SFT               | Alignment                |
|------------------|-----------------|-------------------|--------------------------|
| GPT-4            | Huge web corpus | OpenAI annotators | RLHF                     |
| Claude           | Web + books     | Anthropic         | RLHF + Constitutional AI |
| Llama 2 Chat     | 2T tokens       | Public datasets   | RLHF                     |
| Mistral Instruct | Web             | Public datasets   | DPO                      |
| Gemma Instruct   | Google data     | Instruction data  | SFT only                 |

# The Alignment Tax

**Trade-off:** Alignment can slightly reduce raw capability

| Task               | Base Model | Aligned Model |
|--------------------|------------|---------------|
| Trivia questions   | 82%        | 80%           |
| Code completion    | 76%        | 74%           |
| Math problems      | 68%        | 65%           |
| <b>Helpfulness</b> | 30%        | <b>90%</b>    |
| <b>Safety</b>      | 40%        | <b>95%</b>    |

Worth it for real-world use!

# Two Paradigms: Discriminative vs Generative

## Discriminative Models

$$P(\text{label}|\text{input})$$

- Classification
- Regression
- Everything we learned weeks 1-6

## Generative Models

$$P(\text{data}) \text{ or } P(\text{data}|\text{prompt})$$

- Text generation (LLMs)
- Image generation (Diffusion)
- This is where AI is today

# The Landscape

| Domain | Generative Tool    | What It Creates     |
|--------|--------------------|---------------------|
| Text   | ChatGPT, Claude    | Essays, code, poems |
| Images | DALL-E, Midjourney | Any image from text |
| Audio  | Suno, ElevenLabs   | Music, voices       |
| Video  | Sora, Runway       | Video clips         |
| 3D     | DreamFusion        | 3D models           |
| Code   | Copilot, Cursor    | Working programs    |

## **Part 2: Image Generation**

**From GANs to Diffusion**

# A Brief History

| Year | Model                   | Key Innovation                  |
|------|-------------------------|---------------------------------|
| 2014 | GANs                    | Generator vs Discriminator game |
| 2020 | VQVAE                   | Discrete image tokens           |
| 2021 | DALL-E                  | Text-to-image at scale          |
| 2022 | Stable Diffusion        | Open-source, diffusion models   |
| 2023 | DALL-E 3, Midjourney v5 | Photorealistic quality          |
| 2024 | Flux, SD3               | Even better quality             |

# GANs: The Generator-Discriminator Game

## Generator

- Creates fake images
- Tries to fool discriminator
- Gets better at faking

## Discriminator

- Tells real from fake
- Tries to catch generator
- Gets better at detecting

**Both improve until generated images are indistinguishable from real!**

# Diffusion Models: The New King

Idea: Learn to denoise images

Training:

Real image → Add noise → Noisy image



Model learns to reverse this!

Generation:

Pure noise → Denoise → Denoise → ... → Final image

# Diffusion: Step by Step

| Step | Image State   | What Happens             |
|------|---------------|--------------------------|
| 0    | Pure noise    | Start with random pixels |
| 1    | Mostly noise  | Model removes some noise |
| 2    | Shapes emerge | Structure appears        |
| ...  | ...           | ...                      |
| 50   | Clear image   | Final result             |

Each step removes a little noise!

# Text-to-Image: How It Works

Input: "A photo of a cat wearing a hat on Mars"



- 1. Text Encoder (CLIP)
- 2. Diffusion Model (guided by text)
- 3. VAE Decoder

→ Text embedding  
→ Denoising  
→ Final image



Output: Image of a cat in a hat on Mars!

# Using Image Generation

```
from openai import OpenAI

client = OpenAI()

response = client.images.generate(
    model="dall-e-3",
    prompt="A serene Japanese garden with cherry blossoms",
    size="1024x1024",
    quality="standard",
    n=1,
)

image_url = response.data[0].url
```

# Prompt Engineering for Images

| Bad<br>Prompt | Good Prompt   |
|---------------|---|
| "cat"         | "A fluffy orange tabby cat sleeping on a velvet cushion, soft lighting, photorealistic"       |
| "landscape"   | "Misty mountain landscape at sunrise, oil painting style, dramatic clouds, warm golden light" |

**Key elements:** Subject, style, lighting, composition, quality modifiers

## **Part 3: Multimodal AI**

**Text + Images + More**

# What is Multimodal?

**Modality** = Type of data (text, image, audio, video)

**Multimodal** = Understanding/generating multiple types

| Model  | Modalities   |
|--------|--|
| GPT-4V | Text input + Image input → Text output   |
| DALL-E | Text input → Image output  |
| Gemini | Text + Image + Audio → Text  |
| GPT-4o | Text + Image + Audio  Text + Audio |

# Vision-Language Models

**Input:** Image + Text question

**Output:** Text answer

```
response = client.chat.completions.create(  
    model="gpt-4-vision-preview",  
    messages=[{  
        "role": "user",  
        "content": [  
            {"type": "text", "text": "What's in this image?"},  
            {"type": "image_url", "image_url": {"url": image_url}}  
        ]  
    }]  
)
```

# Use Cases

| Task                 | Input            | Output         |
|----------------------|------------------|----------------|
| Image Captioning     | Photo            | Description    |
| Visual QA            | Photo + Question | Answer         |
| OCR + Understanding  | Document image   | Extracted info |
| Code from Screenshot | UI mockup        | Working code   |

## **Part 4: Using LLM APIs**

**Building with AI**

# The OpenAI API Pattern

```
from openai import OpenAI

client = OpenAI()

response = client.chat.completions.create(
    model="gpt-4",
    messages=[
        {"role": "system", "content": "You are a helpful assistant."},
        {"role": "user", "content": "Explain quantum computing"}
    ],
    temperature=0.7
)

print(response.choices[0].message.content)
```

# Message Roles

| Role      | Purpose        | Example                  |
|-----------|----------------|--------------------------|
| system    | Set behavior   | "You are a Python tutor" |
| user      | User input     | "How do I read a file?"  |
| assistant | Model response | "You can use open()..."  |

```
messages = [
    {"role": "system", "content": "Be concise."},
    {"role": "user", "content": "What is Python?"},
    {"role": "assistant", "content": "A programming language."},
    {"role": "user", "content": "What's it used for?"}
]
```

# Key Parameters

| Parameter                | Controls         | Range                               |
|--------------------------|------------------|-------------------------------------|
| <b>temperature</b>       | Randomness       | 0.0 (deterministic) to 2.0 (random) |
| <b>max_tokens</b>        | Response length  | 1 to context limit                  |
| <b>top_p</b>             | Nucleus sampling | 0.0 to 1.0                          |
| <b>frequency_penalty</b> | Repetition       | -2.0 to 2.0                         |

# Prompt Engineering Basics

| Technique     | Example                                       |
|---------------|---|
| Be specific   | "Write a 3-paragraph summary" not "Summarize" |
| Give examples | "Format: Name: X, Age: Y"                     |
| Role-play     | "You are an expert data scientist..."         |
| Step-by-step  | "Think through this step by step"             |

# Building Applications

```
def analyze_sentiment(text):
    response = client.chat.completions.create(
        model="gpt-4",
        messages=[
            {"role": "system", "content": """
                Analyze sentiment of the text.
                Return JSON: {"sentiment": "positive/negative/neutral",
                              "confidence": 0.0-1.0}
            """, "role": "user", "content": text}
        ],
        temperature=0
    )
    return json.loads(response.choices[0].message.content)
```

## **Part 5: Fine-tuning**

### **Customizing Models**

# When to Fine-tune?

| Scenario              | Use...             |
|-----------------------|--------------------|
| General task          | Prompt engineering |
| Specific style/format | Fine-tuning        |
| Domain knowledge      | RAG (Retrieval)    |
| Custom behavior       | Fine-tuning        |

# Fine-tuning Overview

## 1. PREPARE DATA

- Format: `{"messages": [{"role": ..., "content": ...}]}`
- Need 50–1000+ examples

## 2. UPLOAD DATA

- Upload to OpenAI/Hugging Face

## 3. TRAIN

- Fine-tune on your data
- Usually takes minutes to hours

## 4. USE

- Call your custom model

# Fine-tuning with OpenAI

```
# 1. Upload training file
file = client.files.create(
    file=open("training_data.jsonl", "rb"),
    purpose="fine-tune"
)

# 2. Create fine-tuning job
job = client.fine_tuning.jobs.create(
    training_file=file.id,
    model="gpt-3.5-turbo"
)

# 3. Use your model
response = client.chat.completions.create(
    model="ft:gpt-3.5-turbo:org:custom-name:id",
    messages=[...]
)
```

# Hugging Face: Open Models

```
from transformers import AutoModelForCausalLM, AutoTokenizer

# Load model
model_name = "meta-llama/Llama-2-7b"
tokenizer = AutoTokenizer.from_pretrained(model_name)
model = AutoModelForCausalLM.from_pretrained(model_name)

# Generate
inputs = tokenizer("Hello, how are", return_tensors="pt")
outputs = model.generate(**inputs, max_length=50)
print(tokenizer.decode(outputs[0]))
```

# RAG: Retrieval-Augmented Generation

**Problem:** LLMs don't know your private data

**Solution:** Retrieve relevant documents, add to context

1. User asks question
2. Search your documents for relevant chunks
3. Add chunks to prompt
4. LLM answers using retrieved context

## **Part 6: The Future**

**What's Next?**

# Current Capabilities

| Task             | State             |
|------------------|-------------------|
| Text generation  | Excellent         |
| Code generation  | Very good         |
| Image generation | Excellent         |
| Video generation | Emerging          |
| Audio generation | Good              |
| Reasoning        | Improving rapidly |

# Emerging Trends

| Trend            | What It Means                     |
|------------------|-----------------------------------|
| Agents           | AI that takes actions, uses tools |
| Reasoning models | o1/o3 - think before answering    |
| Multimodal       | Seamless text/image/audio         |
| Smaller models   | Run on phones, edge devices       |
| Open weights     | Llama, Mistral, etc.              |

# AI Agents

**Traditional LLM:** Answer questions

**AI Agent:** Take actions!

```
# Agent can:  
# - Search the web  
# - Run code  
# - Send emails  
# - Book appointments  
# - Write and execute programs
```

# Reasoning Models (o1/o3)

**Standard LLM:** Immediate response

**Reasoning model:** Think step-by-step internally

| Model | Math Score | Science Score |
|-------|------------|---------------|
| GPT-4 | 52%        | 64%           |
| o1    | 83%        | 78%           |
| o3    | 91%        | 87%           |

# Challenges Ahead

| Challenge      | Why It Matters                  |
|----------------|---------------------------------|
| Hallucinations | Models make up facts            |
| Bias           | Reflects training data biases   |
| Alignment      | Ensuring helpful, safe behavior |
| Cost           | Training = millions of dollars  |
| Environment    | Massive energy consumption      |
| Jobs           | Automation concerns             |

# Responsible AI

| Principle      | Implementation              |
|----------------|-----------------------------|
| Transparency   | Disclose AI use             |
| Fairness       | Test for bias               |
| Privacy        | Don't train on private data |
| Safety         | Content filtering           |
| Accountability | Human oversight             |

# **Course Summary**

**What We Learned**

# Your AI Journey

| Week | You Learned                             |
|------|---|
| 1-2  | ML fundamentals, data, train/test       |
| 3    | Linear/Logistic Regression, Trees, KNN  |
| 4    | Cross-validation, Ensembles, Clustering |
| 5    | Neural networks, PyTorch                |
| 6    | CNNs, Object detection, YOLO            |
| 7    | Embeddings, Attention, Transformers     |
| 8    | Generative AI, APIs, Future             |

# The Core Ideas

1. **ML = Learning from data** (not explicit programming)
2. **Supervised learning** is most common
  - Classification (categories) vs Regression (numbers)
3. **Neural networks** can learn complex patterns
  - CNNs for images, Transformers for sequences
4. **Attention is all you need**
  - Modern AI is built on transformers
5. **Generative AI** creates new content
  - Text, images, audio, video

# The Skills You Built

| Skill           | Tools                    |
|-----------------|--------------------------|
| ML basics       | sklearn, pandas, numpy   |
| Deep learning   | PyTorch                  |
| Computer vision | CNNs, YOLO               |
| NLP             | Transformers, APIs       |
| Generative AI   | OpenAI API, Hugging Face |

# Where to Go Next

## Deepen Understanding

- Fast.ai courses
- Stanford CS229, CS231n
- Coursera/Udacity

## Build Projects

- Kaggle competitions
- Personal projects
- Open source contributions

## Stay Current

- arXiv papers

# Key Resources

| Resource         | What It Offers                    |
|------------------|-----------------------------------|
| Hugging Face     | Pre-trained models, datasets      |
| Papers With Code | Latest research + implementations |
| Kaggle           | Competitions, notebooks, data     |
| Fast.ai          | Practical deep learning course    |
| 3Blue1Brown      | Visual math intuition             |
| Andrej Karpathy  | Deep learning from scratch        |

# Key Takeaways

1. AI is pattern recognition at scale
2. Data is everything — garbage in, garbage out
3. Start simple — complex ≠ better
4. Evaluate properly — test set is sacred
5. AI is a tool — you decide how to use it

# Congratulations!

## You Now Understand Modern AI

*"The best way to predict the future is to create it."*

— Alan Kay

Go build something amazing!

Questions?