

# From Language Model to Assistant

## The Complete Journey

SFT, RLHF, and Modern AI Systems

Nipun Batra | IIT Gandhinagar

# The Journey Complete

Week	Topic	The Question
1 - 2	Foundations	How do machines learn?
3 - 4	Supervised Learning	How do we predict?
5	Neural Networks	What makes deep learning special?
6	Computer Vision	How do machines see?
7	Language Models	How do machines understand text?
8	LM → Assistant	How do we make it helpful?

# Last Week's Cliffhanger

We learned:

- LLMs predict the next token
- Transformers enable long-range attention
- Temperature controls creativity

**But we left with a puzzle:**

A model that predicts text well is NOT the same as a helpful assistant!

# The ChatGPT Problem

What We Have	What We Want
Great text completion	Following instructions
"Paris is the capital of..." → "France"	"What's the capital of France?" → "Paris"
Random poem continuation	"Write a poem about..."
Any code that fits	Working, safe code

**Today: How do we bridge this gap?**

# Today's Agenda

1. **The Problem** - Why base models aren't enough
2. **SFT** - Teaching models to follow instructions
3. **RLHF** - Learning from human preferences
4. **DPO** - A simpler alternative
5. **The Full Pipeline** - Pre-training → SFT → Alignment
6. **Image Generation** - Quick overview
7. **AI Ethics & Future** - Responsible AI

# Part 1: The Base Model Problem

Why Prediction Isn't Enough

# Base Model Behavior

A model trained on web text learns to complete web text:

**Prompt:** "How do I make a cake?"

**Base model might respond:**

How do I make a cake?

I'm looking for a simple chocolate cake recipe. My mom used to make one but I lost the recipe card. Any help would be appreciated!

Posted by CakeLover92 on Reddit, March 2019

It's completing forum posts, not answering questions!

# More Base Model Problems

Prompt	Base Model Response	What We Want
"Write Python code to..."	<code># TODO: implement this</code>	Working code
"Is this email spam?"	Continues the email	Yes/No
"Summarize this article"	Writes more article	Summary

Base models are trained to PREDICT, not to HELP.

# The Training Data Problem

**Base model training data:** The entire internet

- Wikipedia, Reddit, Stack Overflow
- Books, news articles, papers
- Code repositories, forums, social media

**What's missing:**

- Explicit instruction-following examples
- Feedback on what makes a "good" response
- Alignment with human values

# **Part 2: Supervised Fine-Tuning (SFT)**

Teaching Models to Follow Instructions

# SFT: The Key Idea

Collect examples of good instruction-following:

Instruction: "What is the capital of France?"

Response: "The capital of France is Paris."

Instruction: "Write a haiku about rain"

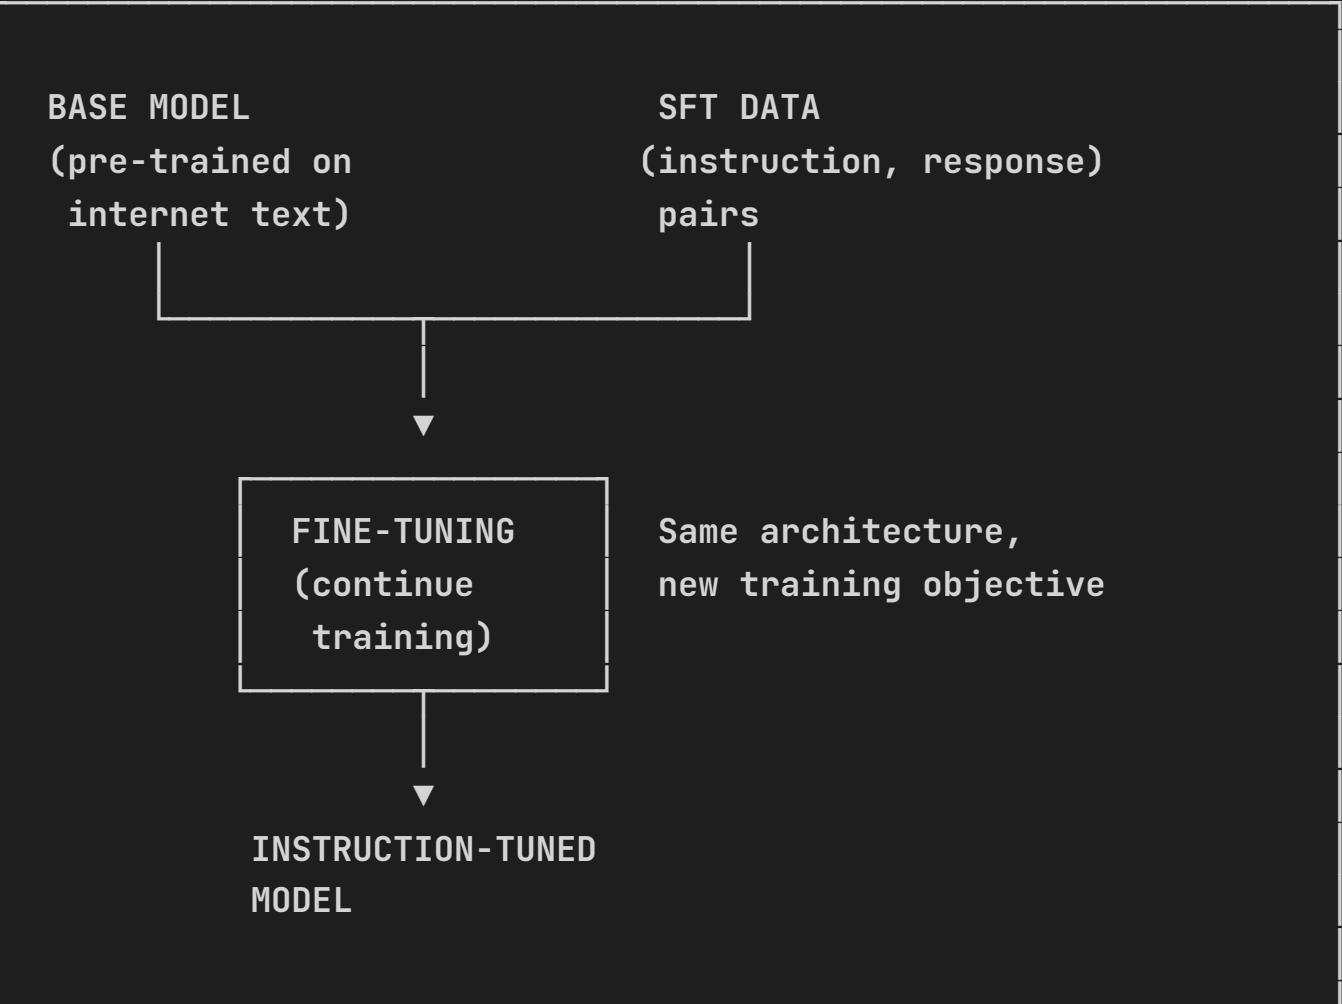
Response: "Gentle drops descend  
Washing away dusty streets  
Earth drinks and is whole"

Instruction: "Translate 'hello' to Spanish"

Response: "Hola"

Then fine-tune the base model on these examples!

# SFT Training Process



# Where Does SFT Data Come From?

Source	Method	Examples
Human annotators	Write ideal responses	Expensive, high quality
Existing datasets	Convert Q&A datasets	FLAN, Natural Instructions
Model distillation	GPT-4 generates for smaller model	Alpaca, Vicuna
User feedback	Real conversations	ChatGPT usage data

# Famous SFT Datasets

Dataset	Size	Method
FLAN	1,800+ tasks	Compiled from NLP datasets
InstructGPT	~100K	Human annotators
Alpaca	52K	GPT-3 generated
ShareGPT	90K+	User-shared ChatGPT conversations
Dolly	15K	Databricks employees

# SFT Training Example

```
from transformers import AutoModelForCausalLM, Trainer

# Load base model
model = AutoModelForCausalLM.from_pretrained("meta-llama/Llama-2-7b")

# Prepare SFT data
# Format: "### Instruction:\n{instruction}\n### Response:\n{response}"
train_data = load_dataset("sft_data.json")

# Fine-tune
trainer = Trainer(
    model=model,
    train_dataset=train_data,
    args=TrainingArguments(
        output_dir=".sft_model",
        num_train_epochs=3,
        per_device_train_batch_size=4,
    )
)
trainer.train()
```

# SFT Results

Model	Before SFT	After SFT
Follows instructions	20%	85%
Appropriate format	15%	90%
Helpful responses	30%	75%

**SFT makes a HUGE difference!**

But there's still a problem...

# The Limits of SFT

SFT teaches **WHAT** to say, but not **WHICH** response is **BEST**.

**Prompt:** "Write a story about a dog"

**Response A:** (boring but correct)

**Response B:** (creative and engaging)

**Response C:** (grammatically perfect but dull)

**SFT can't distinguish between acceptable responses!**

We need a way to learn from **preferences**.

# Part 3: RLHF

Learning from Human Preferences

# RLHF: The Key Insight

Instead of "this is the right answer"...

**Let humans rank responses from best to worst!**

Prompt: "Explain quantum physics simply"

Response A: [Technical jargon, hard to follow]

Response B: [Clear analogy with everyday objects]

Response C: [Accurate but dry explanation]

Human ranking: B > C > A

**Train the model to generate responses like B!**

# RLHF: Three Steps

## STEP 1: Collect Comparisons

---

Human ranks model responses: "A is better than B"

## STEP 2: Train Reward Model

---

Learn to predict human preferences  
reward(response) → score

## STEP 3: Optimize with RL

---

Fine-tune LLM to maximize reward model score  
(using PPO or similar RL algorithm)

# Step 1: Collect Human Preferences

Annotators compare pairs of responses:

Prompt	Response A	Response B	Preference
"What's 2+2?"	"The answer is 4."	"4"	A (more helpful)
"Tell a joke"	[Long, funny joke]	[Short, unfunny]	A
"Write code for..."	[Working code]	[Buggy code]	A

Need thousands of comparisons!

## Step 2: Train Reward Model

**Input:** Prompt + Response

**Output:** Scalar score (how good is this?)

```
class RewardModel(nn.Module):
    def __init__(self, base_model):
        self.base = base_model
        self.head = nn.Linear(hidden_size, 1)

    def forward(self, prompt, response):
        hidden = self.base(prompt + response)
        score = self.head(hidden[:, -1]) # Score from last token
        return score
```

**Training objective:** reward(preferred) > reward(rejected)

# Step 3: RL Fine-tuning

Use PPO (Proximal Policy Optimization) to maximize reward:

```
# Simplified RLHF loop
for prompt in prompts:
    # Generate response from current policy
    response = model.generate(prompt)

    # Get reward
    reward = reward_model(prompt, response)

    # Update model to increase reward
    # (with KL penalty to stay close to SFT model)
    loss = -reward + beta * KL_divergence(model, sft_model)
    loss.backward()
    optimizer.step()
```

# Why the KL Penalty?

**Problem:** Without constraint, model might exploit reward model

Example reward hack:

Reward model likes "helpful" responses

- Model learns to add "I hope this helps!" to everything
- Gets high reward, but responses get worse

**Solution:** Stay close to the SFT model (KL divergence penalty)

# RLHF Results: InstructGPT

OpenAI's InstructGPT paper (2022):

Metric	Base GPT-3	SFT	RLHF
Human preference	22%	33%	71%
Truthfulness	34%	47%	68%
Less harmful	44%	61%	84%

RLHF is what made GPT-3 → ChatGPT!

# RLHF Challenges

Challenge	Why It's Hard
Expensive	Need many human annotations
Slow	RL training is unstable
Reward hacking	Model exploits reward model
Alignment tax	Sometimes hurts raw capability

Is there a simpler alternative?

# Part 4: DPO

A Simpler Path to Alignment

# DPO: Direct Preference Optimization

Skip the reward model entirely!

**Key insight:** We can derive a loss function directly from preferences

Traditional RLHF:

Preferences → Reward Model → RL Training → Aligned Model

DPO:

Preferences → Direct Fine-tuning → Aligned Model

# DPO Loss Function

For each comparison (prompt, preferred response, rejected response):

$$\mathcal{L} = -\log \sigma \left( \beta \log \frac{\pi(y_w|x)}{\pi_{ref}(y_w|x)} - \beta \log \frac{\pi(y_l|x)}{\pi_{ref}(y_l|x)} \right)$$

In simple terms:

- Increase probability of preferred responses
- Decrease probability of rejected responses
- Stay close to reference model

# DPO in Practice

```
from trl import DPOTrainer

# Prepare preference data
# Format: {"prompt": ..., "chosen": ..., "rejected": ...}
train_data = load_preferences_dataset()

# Simple training!
trainer = DPOTrainer(
    model=sft_model,
    ref_model=sft_model, # Reference for KL penalty
    train_dataset=train_data,
    beta=0.1, # Temperature parameter
)
trainer.train()
```

Much simpler than RLHF!

# DPO vs RLHF

Aspect	RLHF	DPO
Reward model	Required	Not needed
Training stability	Unstable (RL)	Stable (supervised)
Compute	3 models needed	1 model + reference
Hyperparameters	Many (PPO)	Few (just $\beta$ )
Performance	Strong	Comparable

DPO is becoming the preferred method!

# Part 5: The Complete Pipeline

Pre-training → SFT → Alignment

# The Full Journey

## STAGE 1: Pre-training

Data: Trillions of tokens (web, books, code)  
Objective: Predict next token  
Compute: Thousands of GPUs, months  
Result: "Base model" - knows language, facts

↓

## STAGE 2: Supervised Fine-Tuning (SFT)

Data: ~100K instruction-response pairs  
Objective: Learn to follow instructions  
Compute: Hours to days  
Result: "Instruction model" - follows commands

↓

## STAGE 3: Alignment (RLHF or DPO)

Data: Human preference comparisons  
Objective: Be helpful, harmless, honest  
Compute: Days  
Result: "AI Assistant" - ChatGPT, Claude, etc.

# Real World Examples

Model	Pre-training	SFT	Alignment
GPT-4	Huge web corpus	OpenAI annotators	RLHF
Claude	Web + books	Anthropic	RLHF + Constitutional AI
Llama 2 Chat	2T tokens	Public datasets	RLHF
Mistral Instruct	Web	Public datasets	DPO
Gemma Instruct	Google data	Instruction data	SFT only

# The Alignment Tax

**Trade-off:** Alignment can slightly reduce raw capability

Task	Base Model	Aligned Model
Trivia questions	82%	80%
Code completion	76%	74%
Math problems	68%	65%
<b>Helpfulness</b>	30%	<b>90%</b>
<b>Safety</b>	40%	<b>95%</b>

**Worth it for real-world use!**

# Two Paradigms: Discriminative vs Generative

## Discriminative Models

$$P(\text{label}|\text{input})$$

- Classification
- Regression
- Everything we learned weeks 1-6

## Generative Models

$$P(\text{data}) \text{ or } P(\text{data}|\text{prompt})$$

- Text generation (LLMs)
- Image generation (Diffusion)
- This is where AI is today

# The Landscape

Domain	Generative Tool	What It Creates
Text	ChatGPT, Claude	Essays, code, poems
Images	DALL-E, Midjourney	Any image from text
Audio	Suno, ElevenLabs	Music, voices
Video	Sora, Runway	Video clips
3D	DreamFusion	3D models
Code	Copilot, Cursor	Working programs

# Part 2: Image Generation

From GANs to Diffusion

# A Brief History

Year	Model	Key Innovation
2014	GANs	Generator vs Discriminator game
2020	VQVAE	Discrete image tokens
2021	DALL-E	Text-to-image at scale
2022	Stable Diffusion	Open-source, diffusion models
2023	DALL-E 3, Midjourney v5	Photorealistic quality
2024	Flux, SD3	Even better quality

# GANs: The Generator-Discriminator Game

## Generator

- Creates fake images
- Tries to fool discriminator
- Gets better at faking

## Discriminator

- Tells real from fake
- Tries to catch generator
- Gets better at detecting

Both improve until generated images are indistinguishable from real!

# Diffusion Models: The New King

**Idea:** Learn to **denoise** images

**Training:**

Real image → Add noise → Noisy image



Model learns to reverse this!

**Generation:**

Pure noise → Denoise → Denoise → ... → Final image

# Diffusion: Step by Step

Step	Image State	What Happens
0	Pure noise	Start with random pixels
1	Mostly noise	Model removes some noise
2	Shapes emerge	Structure appears
...	...	...
50	Clear image	Final result

Each step removes a little noise!

# Text-to-Image: How It Works

Input: "A photo of a cat wearing a hat on Mars"

↓

- 1. Text Encoder (CLIP)
- 2. Diffusion Model (guided by text)
- 3. VAE Decoder

→ Text embedding  
→ Denoising  
→ Final image

↓

Output: Image of a cat in a hat on Mars!

# Using Image Generation

```
from openai import OpenAI

client = OpenAI()

response = client.images.generate(
    model="dall-e-3",
    prompt="A serene Japanese garden with cherry blossoms",
    size="1024x1024",
    quality="standard",
    n=1,
)

image_url = response.data[0].url
```

# Prompt Engineering for Images

Bad Prompt	Good Prompt
"cat"	"A fluffy orange tabby cat sleeping on a velvet cushion, soft lighting, photorealistic"
"landscape"	"Misty mountain landscape at sunrise, oil painting style, dramatic clouds, warm golden light"

**Key elements:** Subject, style, lighting, composition, quality modifiers

# Part 3: Multimodal AI

Text + Images + More

# What is Multimodal?

**Modality** = Type of data (text, image, audio, video)

**Multimodal** = Understanding/generating multiple types

Model	Modalities
GPT-4V	Text input + Image input → Text output
DALL-E	Text input → Image output
Gemini	Text + Image + Audio → Text
GPT-4o	Text + Image + Audio ➡ Text + Audio

# Vision-Language Models

**Input:** Image + Text question

**Output:** Text answer

```
response = client.chat.completions.create(  
    model="gpt-4-vision-preview",  
    messages=[{  
        "role": "user",  
        "content": [  
            {"type": "text", "text": "What's in this image?"},  
            {"type": "image_url", "image_url": {"url": image_url}}  
        ]  
    }]  
)
```

# Use Cases

Task	Input	Output
Image Captioning	Photo	Description
Visual QA	Photo + Question	Answer
OCR + Understanding	Document image	Extracted info
Code from Screenshot	UI mockup	Working code

# Part 4: Using LLM APIs

Building with AI

# The OpenAI API Pattern

```
from openai import OpenAI

client = OpenAI()

response = client.chat.completions.create(
    model="gpt-4",
    messages=[
        {"role": "system", "content": "You are a helpful assistant."},
        {"role": "user", "content": "Explain quantum computing"}
    ],
    temperature=0.7
)

print(response.choices[0].message.content)
```

# Message Roles

Role	Purpose	Example
system	Set behavior	"You are a Python tutor"
user	User input	"How do I read a file?"
assistant	Model response	"You can use open()..."

```
messages = [
    {"role": "system", "content": "Be concise."},
    {"role": "user", "content": "What is Python?"},
    {"role": "assistant", "content": "A programming language."},
    {"role": "user", "content": "What's it used for?"}
]
```

# Key Parameters

Parameter	Controls	Range
temperature	Randomness	0.0 (deterministic) to 2.0 (random)
max_tokens	Response length	1 to context limit
top_p	Nucleus sampling	0.0 to 1.0
frequency_penalty	Repetition	-2.0 to 2.0

# Prompt Engineering Basics

Technique	Example
Be specific	"Write a 3-paragraph summary" not "Summarize"
Give examples	"Format: Name: X, Age: Y"
Role-play	"You are an expert data scientist..."
Step-by-step	"Think through this step by step"

# Building Applications

```
def analyze_sentiment(text):
    response = client.chat.completions.create(
        model="gpt-4",
        messages=[
            {"role": "system", "content": """
                Analyze sentiment of the text.
                Return JSON: {"sentiment": "positive/negative/neutral",
                              "confidence": 0.0-1.0}
            """,
            {"role": "user", "content": text}
        ],
        temperature=0
    )
    return json.loads(response.choices[0].message.content)
```

# Part 5: Fine-tuning

Customizing Models

# When to Fine-tune?

Scenario	Use...
General task	Prompt engineering
Specific style/format	Fine-tuning
Domain knowledge	RAG (Retrieval)
Custom behavior	Fine-tuning

# Fine-tuning Overview

## 1. PREPARE DATA

- Format: {"messages": [{"role": "...", "content": ...}]}
  - Need 50-1000+ examples

## 2. UPLOAD DATA

- Upload to OpenAI/Hugging Face

## 3. TRAIN

- Fine-tune on your data
  - Usually takes minutes to hours

## 4. USE

- Call your custom model

# Fine-tuning with OpenAI

```
# 1. Upload training file
file = client.files.create(
    file=open("training_data.jsonl", "rb"),
    purpose="fine-tune"
)

# 2. Create fine-tuning job
job = client.fine_tuning.jobs.create(
    training_file=file.id,
    model="gpt-3.5-turbo"
)

# 3. Use your model
response = client.chat.completions.create(
    model="ft:gpt-3.5-turbo:org:custom-name:id",
    messages=[ ... ]
)
```

# Hugging Face: Open Models

```
from transformers import AutoModelForCausalLM, AutoTokenizer

# Load model
model_name = "meta-llama/Llama-2-7b"
tokenizer = AutoTokenizer.from_pretrained(model_name)
model = AutoModelForCausalLM.from_pretrained(model_name)

# Generate
inputs = tokenizer("Hello, how are", return_tensors="pt")
outputs = model.generate(**inputs, max_length=50)
print(tokenizer.decode(outputs[0]))
```

# RAG: Retrieval-Augmented Generation

**Problem:** LLMs don't know your private data

**Solution:** Retrieve relevant documents, add to context

1. User asks question
2. Search your documents for relevant chunks
3. Add chunks to prompt
4. LLM answers using retrieved context

# Part 6: The Future

What's Next?

# Current Capabilities

Task	State
Text generation	Excellent
Code generation	Very good
Image generation	Excellent
Video generation	Emerging
Audio generation	Good
Reasoning	Improving rapidly

# Emerging Trends

Trend	What It Means
Agents	AI that takes actions, uses tools
Reasoning models	o1/o3 - think before answering
Multimodal	Seamless text/image/audio
Smaller models	Run on phones, edge devices
Open weights	Llama, Mistral, etc.

# AI Agents

**Traditional LLM:** Answer questions

**AI Agent:** Take actions!

```
# Agent can:  
# - Search the web  
# - Run code  
# - Send emails  
# - Book appointments  
# - Write and execute programs
```

# Reasoning Models (o1/o3)

**Standard LLM:** Immediate response

**Reasoning model:** Think step-by-step internally

Model	Math Score	Science Score
GPT-4	52%	64%
o1	83%	78%
o3	91%	87%

# Challenges Ahead

Challenge	Why It Matters
Hallucinations	Models make up facts
Bias	Reflects training data biases
Alignment	Ensuring helpful, safe behavior
Cost	Training = millions of dollars
Environment	Massive energy consumption
Jobs	Automation concerns

# Responsible AI

Principle	Implementation
Transparency	Disclose AI use
Fairness	Test for bias
Privacy	Don't train on private data
Safety	Content filtering
Accountability	Human oversight

# Course Summary

What We Learned

# Your AI Journey

Week	You Learned
1 - 2	ML fundamentals, data, train/test
3	Linear/Logistic Regression, Trees, KNN
4	Cross-validation, Ensembles, Clustering
5	Neural networks, PyTorch
6	CNNs, Object detection, YOLO
7	Embeddings, Attention, Transformers
8	Generative AI, APIs, Future

# The Core Ideas

1. **ML = Learning from data** (not explicit programming)

2. **Supervised learning** is most common

- Classification (categories) vs Regression (numbers)

3. **Neural networks** can learn complex patterns

- CNNs for images, Transformers for sequences

4. **Attention is all you need**

- Modern AI is built on transformers

5. **Generative AI** creates new content

- Text, images, audio, video

# The Skills You Built

Skill	Tools
ML basics	sklearn, pandas, numpy
Deep learning	PyTorch
Computer vision	CNNs, YOLO
NLP	Transformers, APIs
Generative AI	OpenAI API, Hugging Face

# Where to Go Next

## Deepen Understanding

- Fast.ai courses
- Stanford CS229, CS231n
- Coursera/Udacity

## Build Projects

- Kaggle competitions
- Personal projects
- Open source contributions

## Stay Current

- arXiv papers
- AI newsletters
- Twitter/X AI community

## Specialize

- Computer Vision
- NLP
- Reinforcement Learning
- AI Safety

# Key Resources

Resource	What It Offers
Hugging Face	Pre-trained models, datasets
Papers With Code	Latest research + implementations
Kaggle	Competitions, notebooks, data
Fast.ai	Practical deep learning course
3Blue1Brown	Visual math intuition
Andrej Karpathy	Deep learning from scratch

# Key Takeaways

1. AI is pattern recognition at scale
2. Data is everything — garbage in, garbage out
3. Start simple — complex ≠ better
4. Evaluate properly — test set is sacred
5. AI is a tool — you decide how to use it

# Congratulations!

You Now Understand Modern AI

*"The best way to predict the future is to create it."*

— Alan Kay

Go build something amazing!

Questions?