

The Machine Learning Task Zoo

A Safari Through 40+ Real-World AI Problems

Nipun Batra · IIT Gandhinagar

What We'll Explore Today

THE ML TASK SAFARI MAP

COMPUTER VISION

Classification
Detection
Segmentation
Pose Estimation
Depth Estimation

AUDIO & SPEECH

Speech-to-Text
Text-to-Speech
Speaker Recognition

REINFORCEMENT

Game Playing
Robot Control

NATURAL LANGUAGE PROCESSING

Sentiment Analysis
Named Entity Recognition
Translation
Summarization
Question Answering

GENERATIVE MODELS

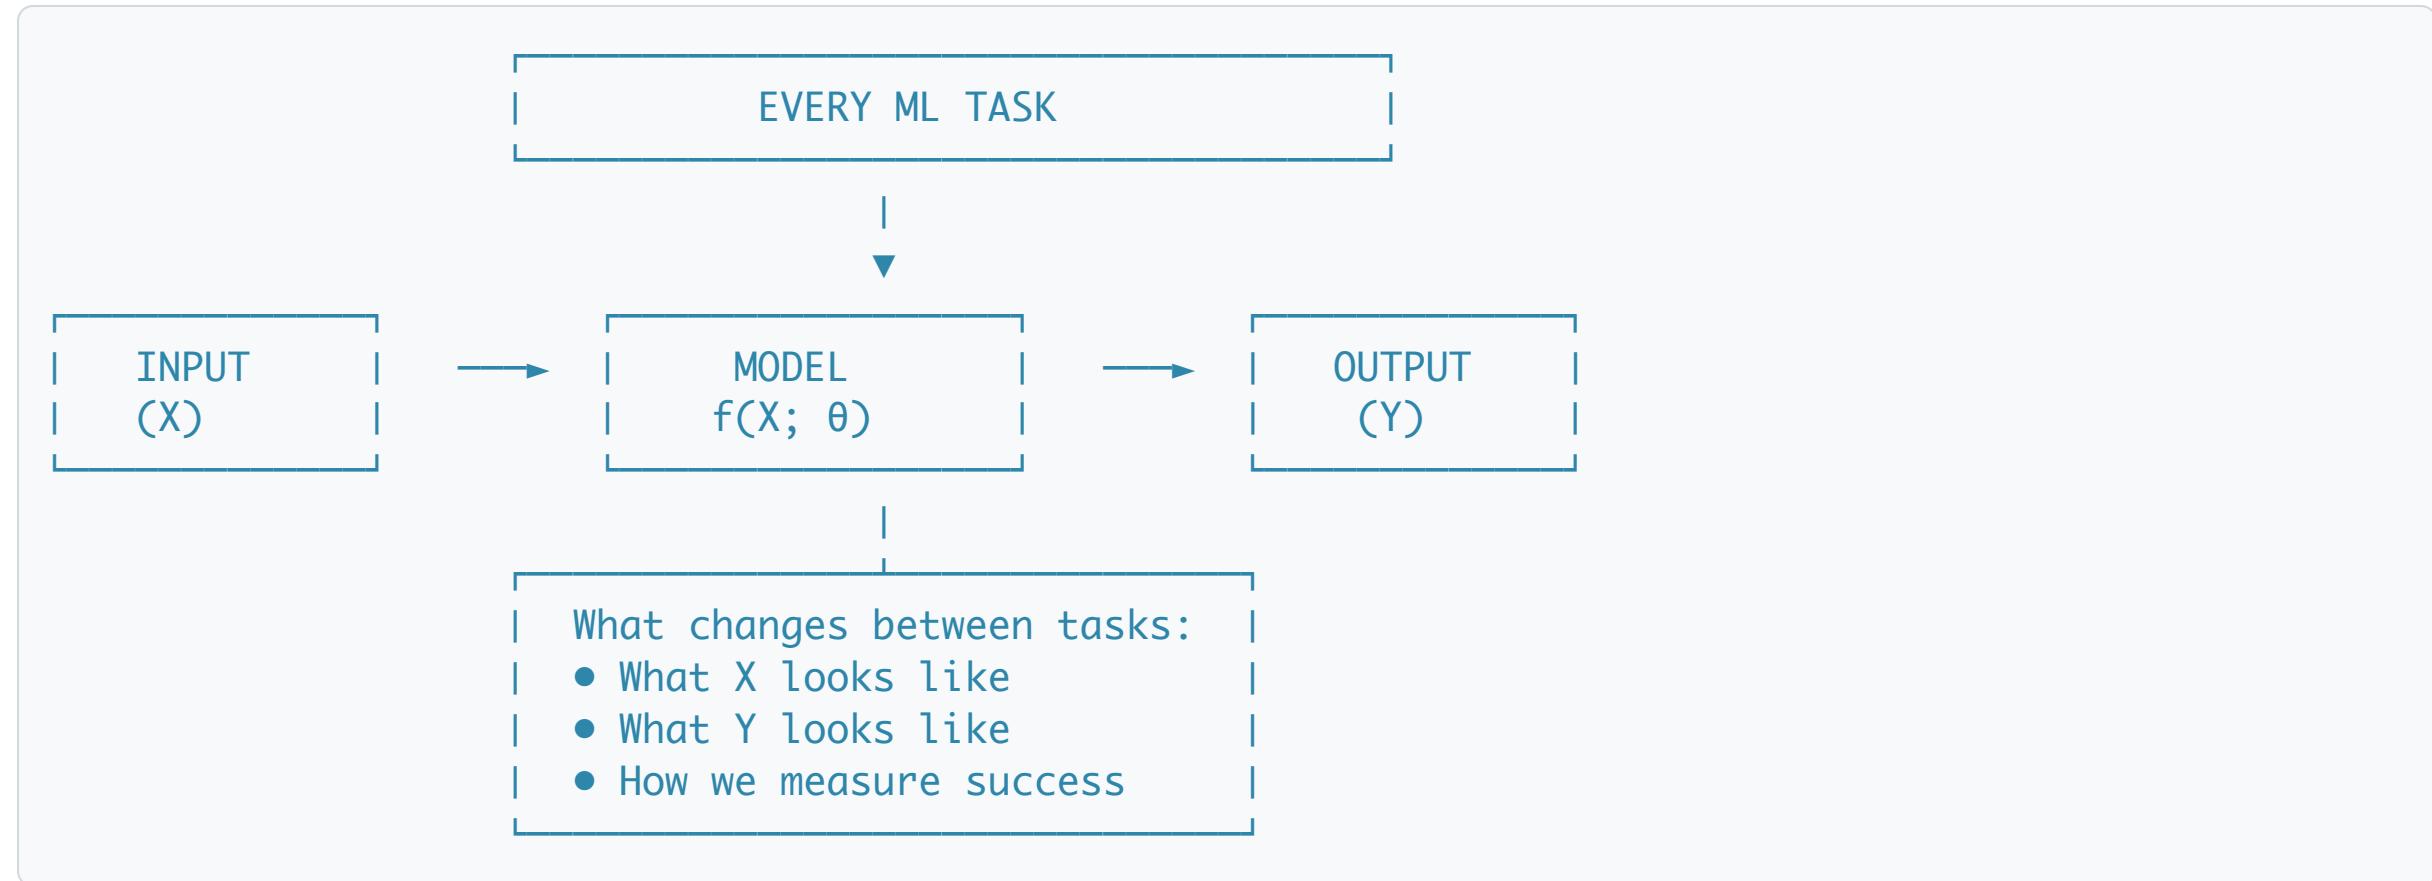
Image Generation
Text Generation
Video Generation

MULTIMODAL

Visual QA
Image Captioning

The Universal ML Recipe

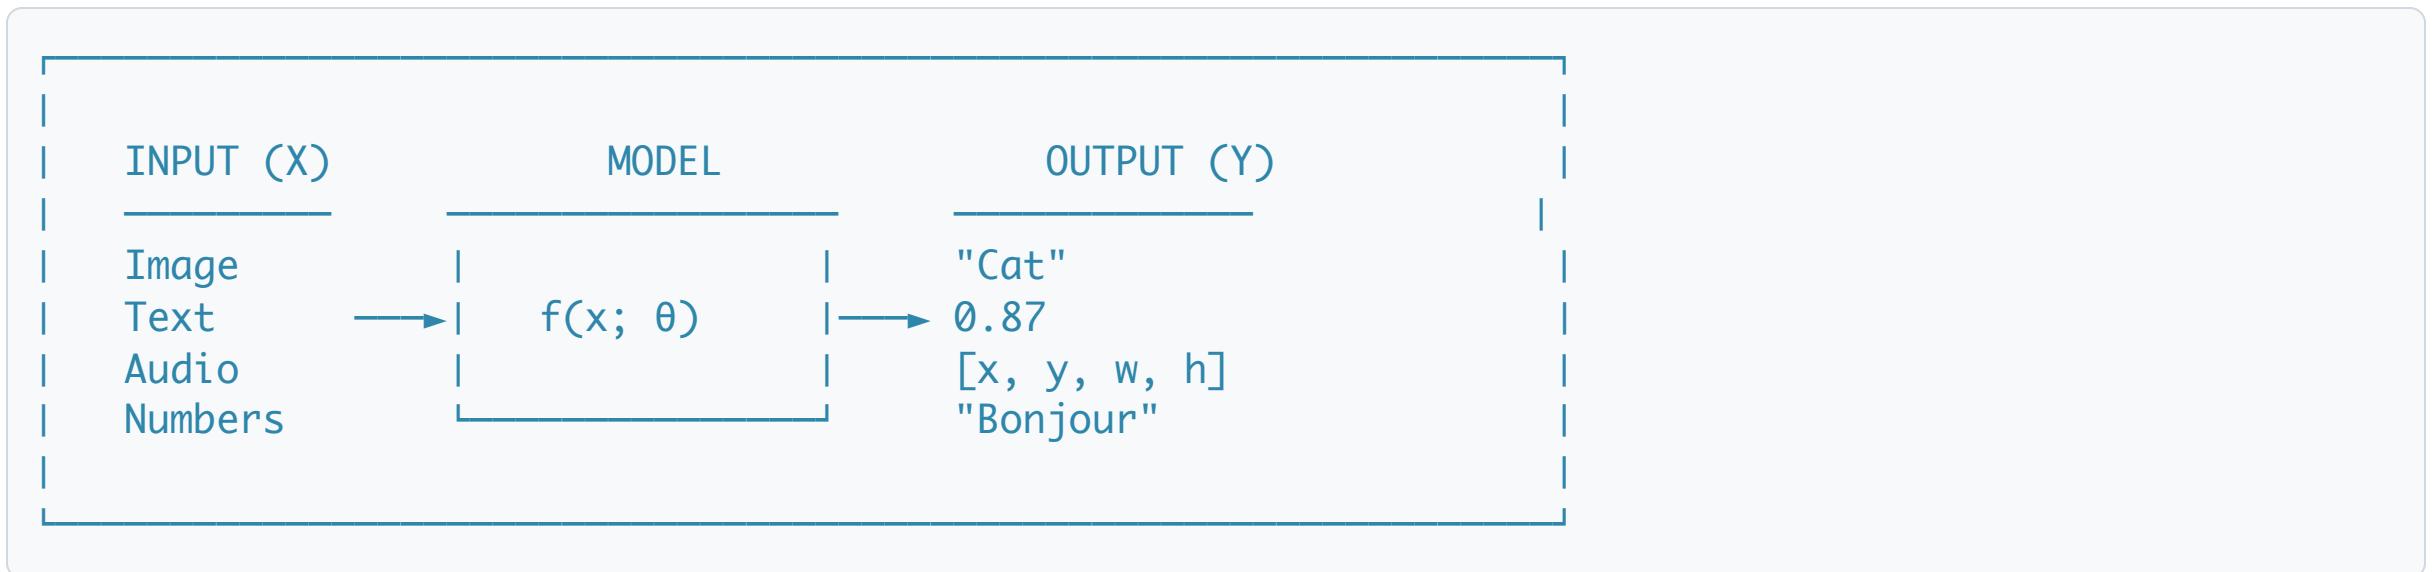
Before we dive into 40+ tasks, remember this simple pattern:



The same Transformer architecture powers ChatGPT, DALL-E, and self-driving cars!

How to Think About ML Tasks

Every task is defined by **what goes in** and **what comes out**:



The same model architecture can solve many different tasks — what changes is the data!

A Simple Classification

ML TASKS BY INPUT/OUTPUT

INPUT TYPE → OUTPUT TYPE

Image	→	Label	(Classification)
Image	→	Boxes	(Detection)
Image	→	Pixel Labels	(Segmentation)
Image	→	Text	(Captioning)
Text	→	Label	(Sentiment)
Text	→	Text	(Translation)
Audio	→	Text	(Speech-to-Text)
Text	→	Audio	(Text-to-Speech)
Numbers	→	Number	(Regression)
Numbers	→	Groups	(Clustering)
Noise	→	Image	(Generation)
Text	→	Image	(Text-to-Image)
Game State	→	Action	(RL)

Domain 1: Computer Vision

Teaching Machines to See

"A picture is worth a thousand words... to a neural network, it's worth millions of numbers!"

The Vision Task Hierarchy

LEVEL 1: Classification "There's a dog somewhere in this image"

Easiest: Just need the answer

LEVEL 2: Detection "There's a dog at position (x, y, w, h)"

Harder: Need to locate it with a box

LEVEL 3: Segmentation "These exact pixels belong to the dog"

Even harder: Pixel-perfect boundaries

LEVEL 4: Pose Estimation "Dog's head is at (x₁,y₁), legs at ..."

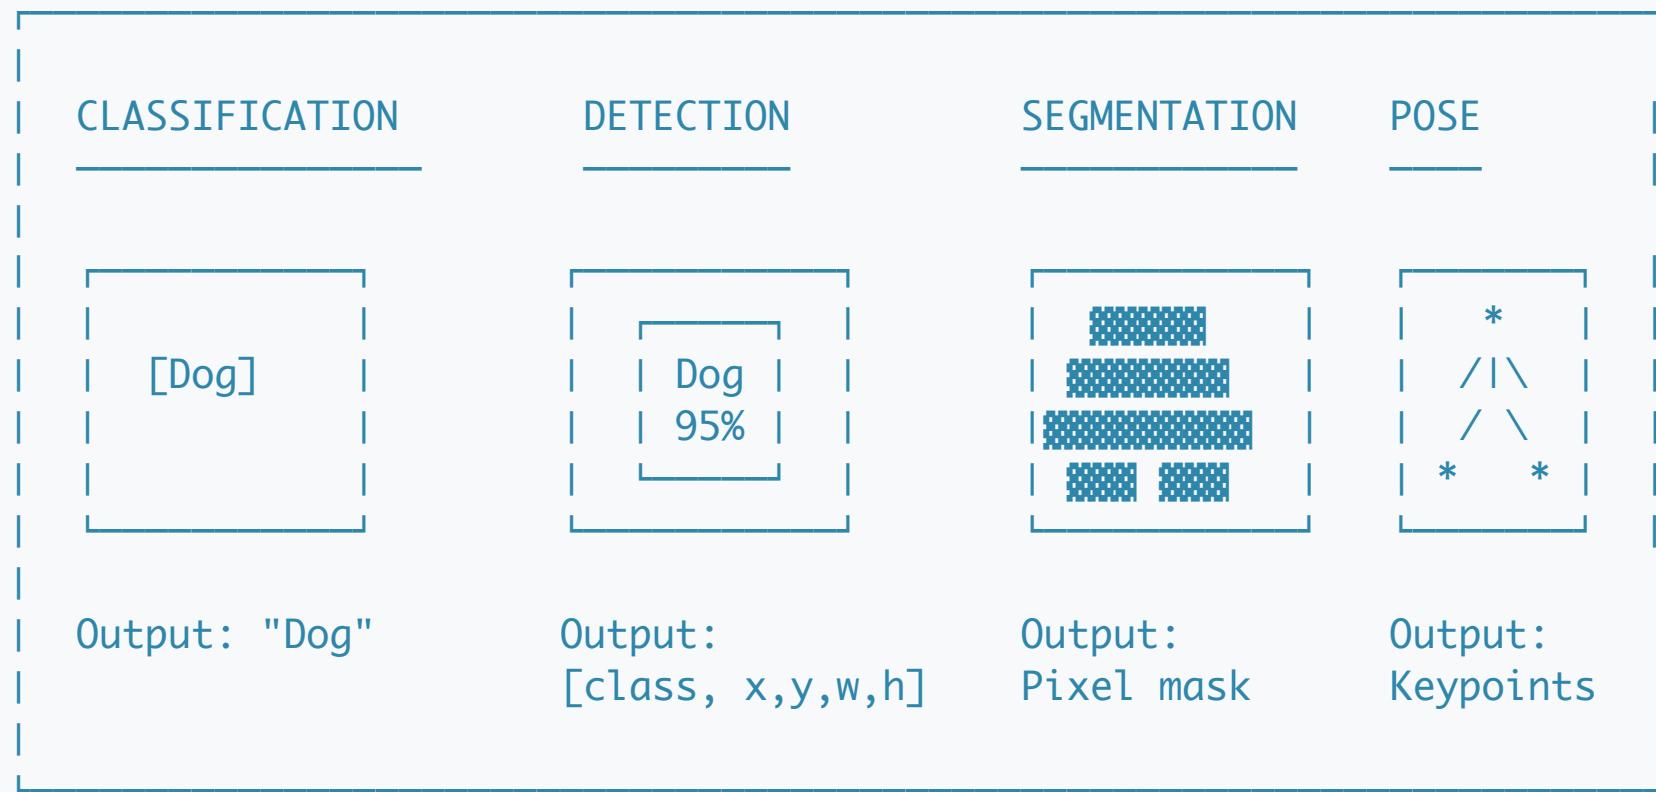
Hardest: Find specific body parts/keypoints

More precision, more data, more compute needed



Let's See This Visually

SAME IMAGE, DIFFERENT TASKS:



Task 1: Image Classification

What: Assign one label to an image.

[Photo of
a dog]

→ "Golden Retriever"
Confidence: 94.2%

Real-world uses:

- Google Photos auto-tagging
- Medical X-ray diagnosis
- Quality control in factories
- Plant disease detection

Example: MNIST Digits

Input: 28×28 grayscale image
Output: One of {0, 1, 2, ..., 9}



→ "7" (98.5%)

10 classes, 60K training images
The "Hello World" of computer vision!

The Math Behind Classification

Input Image (e.g., $224 \times 224 \times 3 = 150,528$ numbers)



Neural Network
(millions of
parameters)



Raw scores (logits)
[2.5, -1.2, 8.7, 0.3, ...] (one per class)



Softmax

← Converts to probabilities



[0.01, 0.00, 0.94, 0.05, ...] (sums to 1.0)



Prediction: Class 3 (94% confidence)

ImageNet: The Olympics of Image Classification

IMAGENET CHALLENGE

Dataset: 14 million images, 1000 classes

Classes include:

[Dog]	[Car]	[Music]	[Food]	[Home]
120 dog breeds!	Cars	Musical instr.	Foods	Objects

Year	Winner	Top-5 Error	Note
2010	Traditional ML	28.2%	Hand-crafted features
2012	AlexNet (CNN)	16.4%	Deep learning begins!
2015	ResNet	3.6%	Superhuman performance!
2020	ViT	1.0%	Transformers enter vision

Multi-Label Classification

Sometimes one label isn't enough!

SINGLE-LABEL

[Cat]

-> "Cat"
(one class)

Each image has
exactly ONE label

Use: Softmax
 Σ probabilities = 1

MULTI-LABEL

Cat + Dog
+ Couch

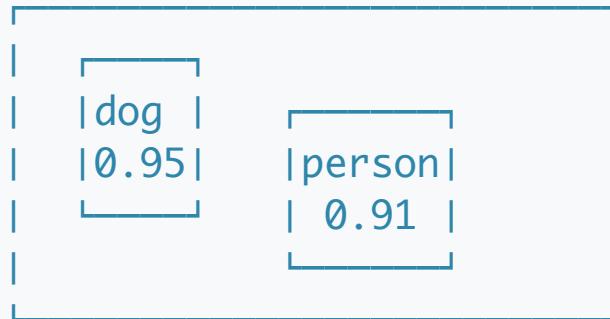
-> ["Cat", "Dog"]
["Couch"]
["Indoors"]

Each image can have
MULTIPLE labels

Use: Sigmoid (per class)
Each class: 0 to 1 independently

Task 2: Object Detection

What: Find objects AND locate them with boxes.



Output for each detection:

- Class label ("dog")
- Confidence score (0.95)
- Bounding box: (x, y, width, height)

Example: Self-Driving Car

Detections in one frame:

- └ Car at (120, 80) conf: 0.97
- └ Car at (400, 90) conf: 0.89
- └ Person at (300, 150) conf: 0.92
- └ Bicycle at (50, 200) conf: 0.88
- └ Traffic Light at (250, 20) conf: 0.99

Must process 30+ frames/second!

Detection vs Classification: Key Differences

CLASSIFICATION vs DETECTION

CLASSIFICATION

Input: One image
Output: One label

Assumes: Object fills
most of image

Architecture:
CNN → FC → Softmax

Example:
"Is this a cat or dog?"

Popular Models:
ResNet, EfficientNet

DETECTION

Input: One image
Output: List of (class, box)

Handles: Multiple objects,
any size, anywhere

Architecture:
CNN → Multiple detection heads

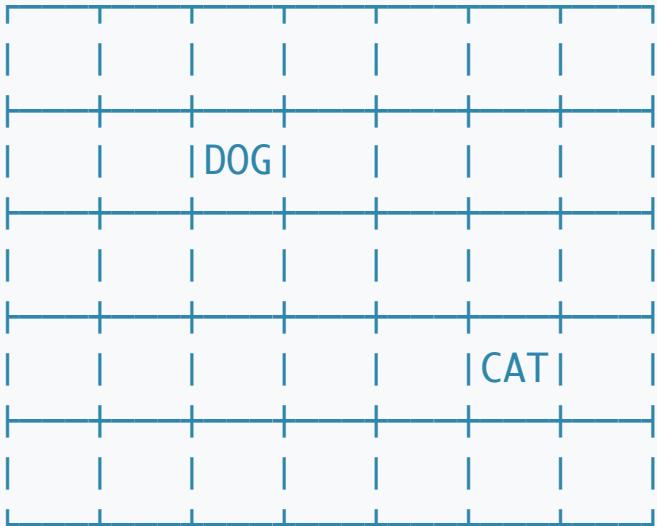
Example:
"Find all cats and dogs"

Popular Models:
YOLO, Faster R-CNN, DETR

How YOLO Works (Simplified)

"You Only Look Once" - Fast single-pass detection

Step 1: Divide image into grid (e.g., 7x7)



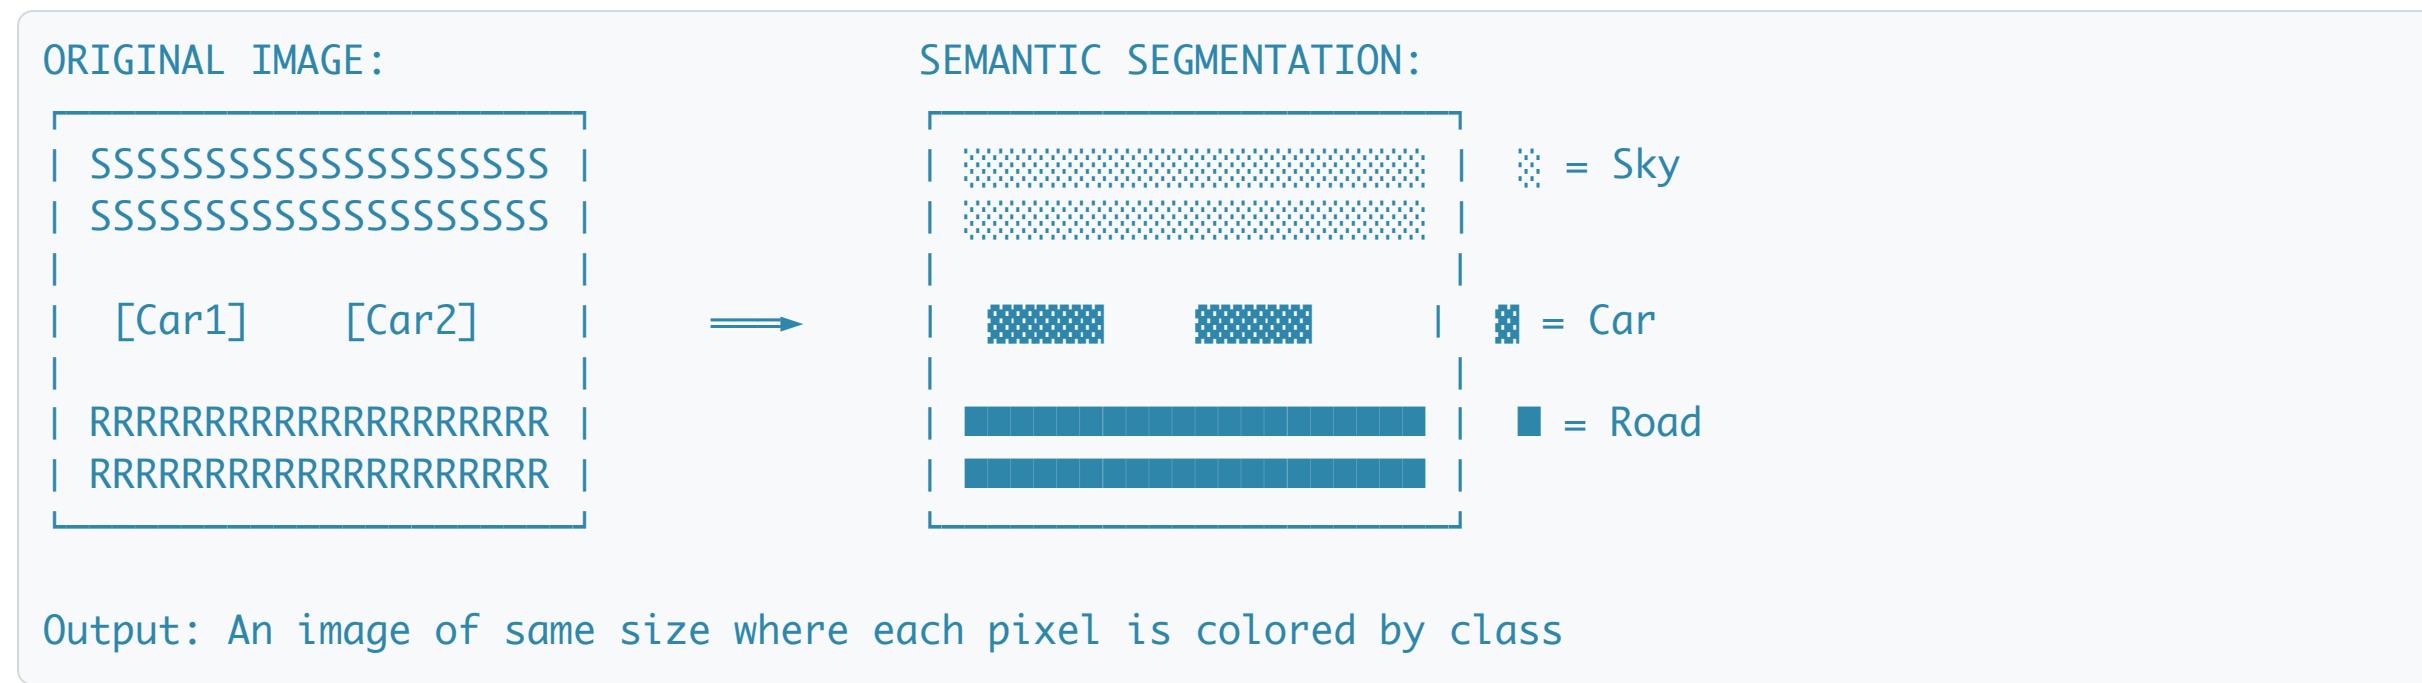
Each cell predicts:
* B bounding boxes
• Confidence scores
• C class probabilities

Step 2: Remove overlapping boxes (Non-Max Suppression)

Step 3: Output final detections

Task 3: Semantic Segmentation

What: Label every pixel with its class.

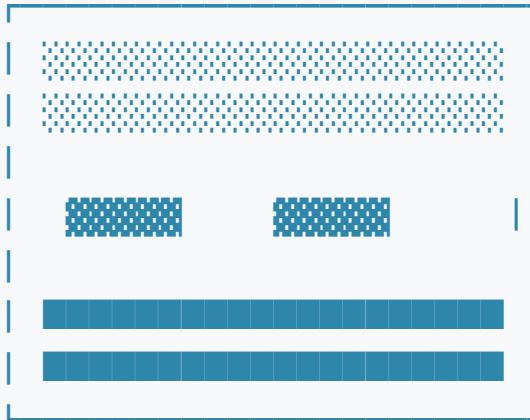


****Both cars have the same color** — semantic segmentation doesn't distinguish between instances of the same class!**

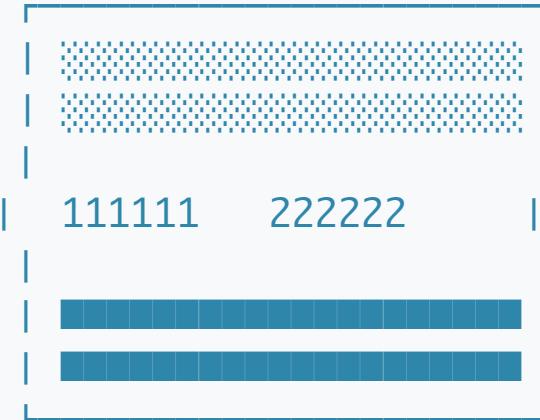
Task 4: Instance Segmentation

What: Label every pixel AND distinguish individual objects.

SEMANTIC SEGMENTATION:



INSTANCE SEGMENTATION:



vs

Both cars = same "Car" color
Can't tell them apart!

Car #1 = Blue, Car #2 = Green
Can track each car individually

Self-driving cars need instance segmentation — you must track which car is which to predict their movements!

Panoptic Segmentation: The Complete Picture

Combining everything: Semantic + Instance

PANOPTIC SEGMENTATION

Sky (stuff - no instances)



Car #1

Car #1

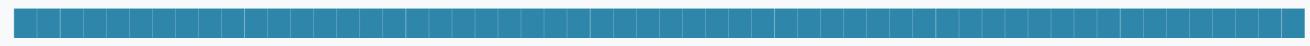
Person #1

Person #1

Car #2

Car #2

Road (stuff - no instances)

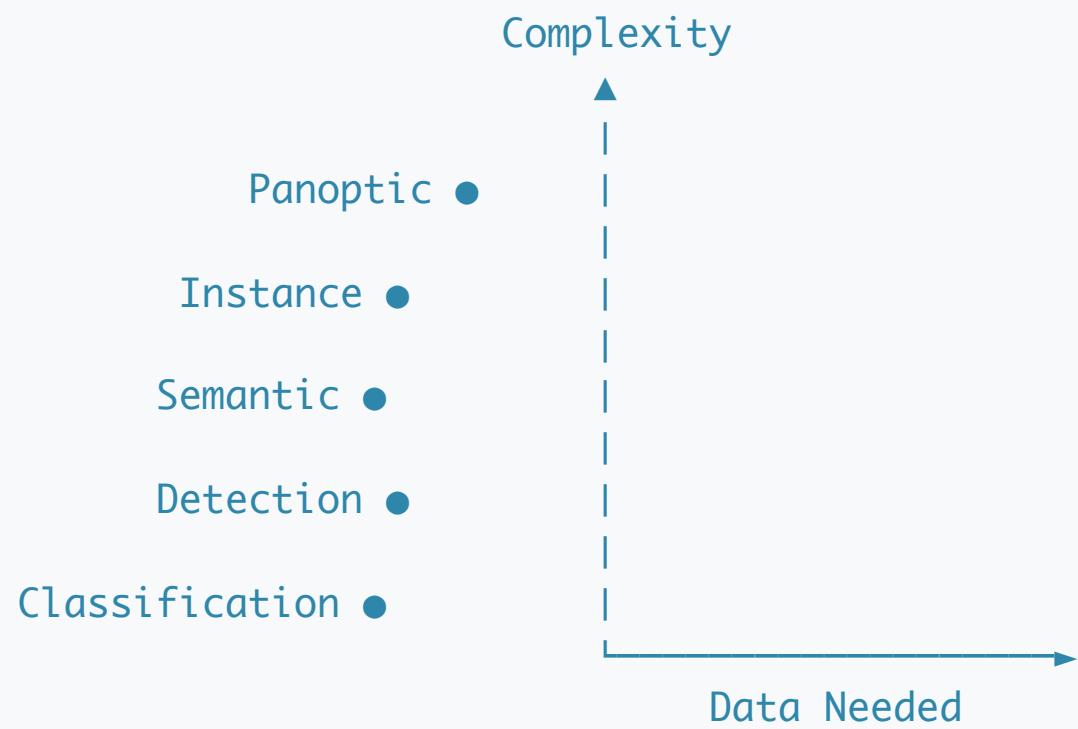


"Stuff" classes: sky, road, grass (don't count instances)

"Things" classes: cars, people (each instance gets unique ID)

Segmentation Summary

Task	What it outputs	Can count objects?	Use case
Semantic	Pixel classes	No	Land use mapping, medical imaging
Instance	Pixel + instance IDs	Yes	Object tracking, robotics
Panoptic	Both combined	Yes for "things"	Autonomous driving



Task 5: Pose Estimation

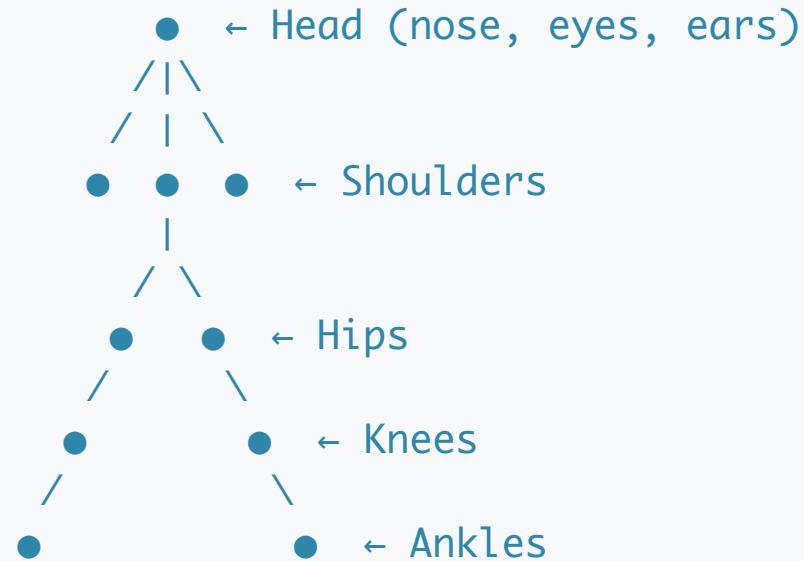
What: Find body keypoints (skeleton) of humans or animals.

Original Photo:

Person
standing
with
raised arm



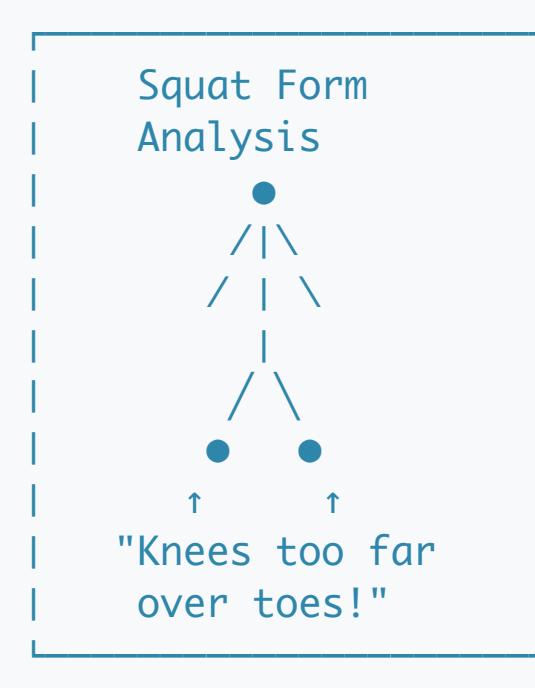
Detected Skeleton:



Output: 17 keypoints with (x, y, confidence) each

Pose Estimation: Real Applications

Fitness & Sports



Other Applications

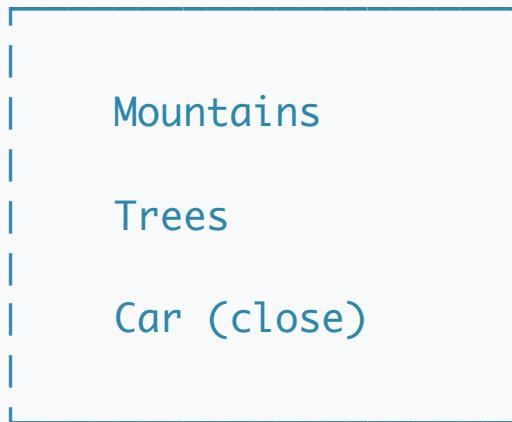
- Motion capture for movies/games
- Running form analysis
- Dance move recognition
- Sign language interpretation
- Controller-free gaming (Kinect)
- Fall detection for elderly

Apple Fitness+ uses pose estimation to analyze your workout form in real-time!

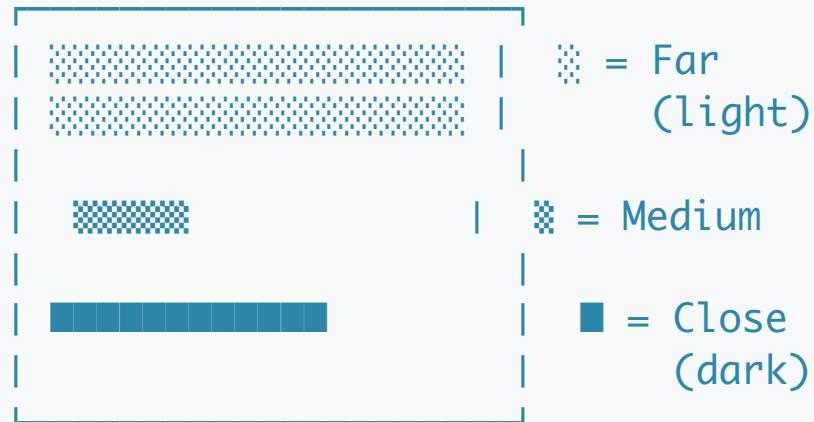
Task 6: Depth Estimation

What: Predict distance of each pixel from the camera.

RGB Image:



Depth Map:



-->

Output: Same-size image where pixel intensity = distance

One camera, 3D understanding! Traditional 3D sensing requires special hardware (LiDAR, stereo cameras), but AI can estimate depth from a single RGB image.

Depth Estimation: How It's Used

AR/VR

Place virtual
furniture in
your room!

[Sofa]
(knows it's
on the floor)

Portrait Mode

Blur background
based on depth

[Person sharp]
[Floor blurred]

Robotics

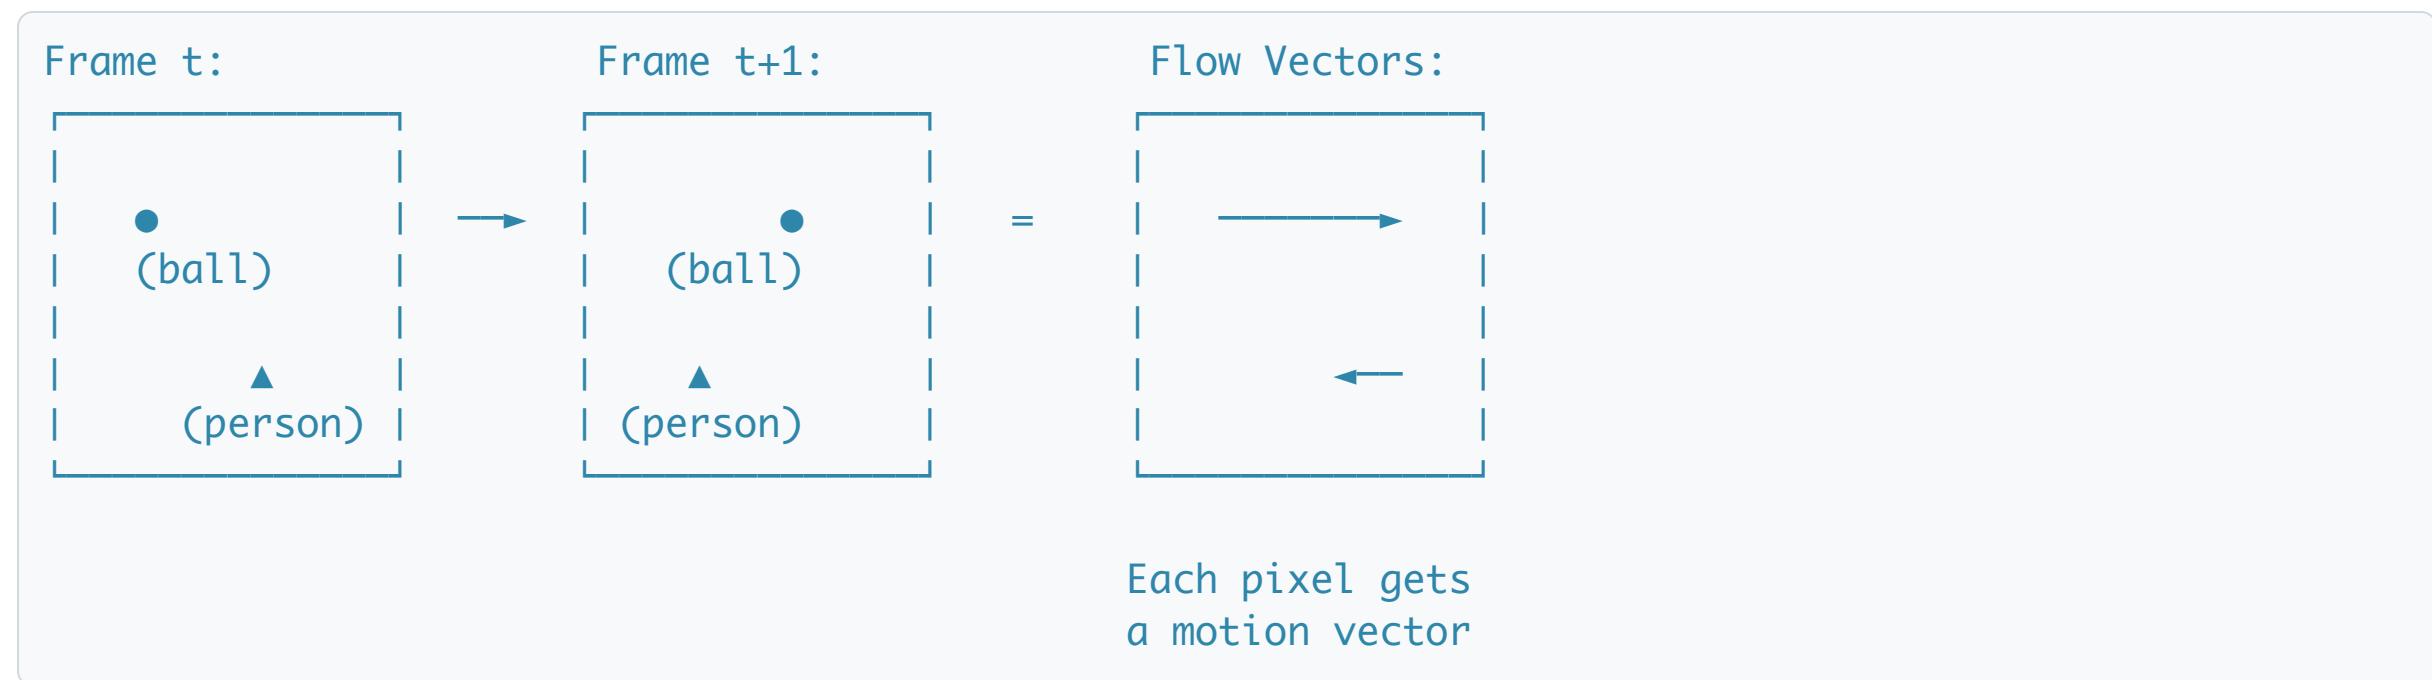
Navigate without
bumping into
objects

Robot --> Box
(knows distance)

The iPhone's **Portrait Mode** uses a combination of depth sensors AND neural network depth estimation!

Task 7: Optical Flow

What: Track how each pixel moves between video frames.



Key insight: Every pixel gets a (dx, dy) vector showing where it moved!

Optical Flow Applications

OPTICAL FLOW USE CASES

VIDEO COMPRESSION

Instead of storing every frame, store keyframes + motion vectors
Result: 10x smaller file sizes

ACTION RECOGNITION

"Running" = specific pattern of flow vectors
"Waving" = different pattern

AUTONOMOUS DRIVING

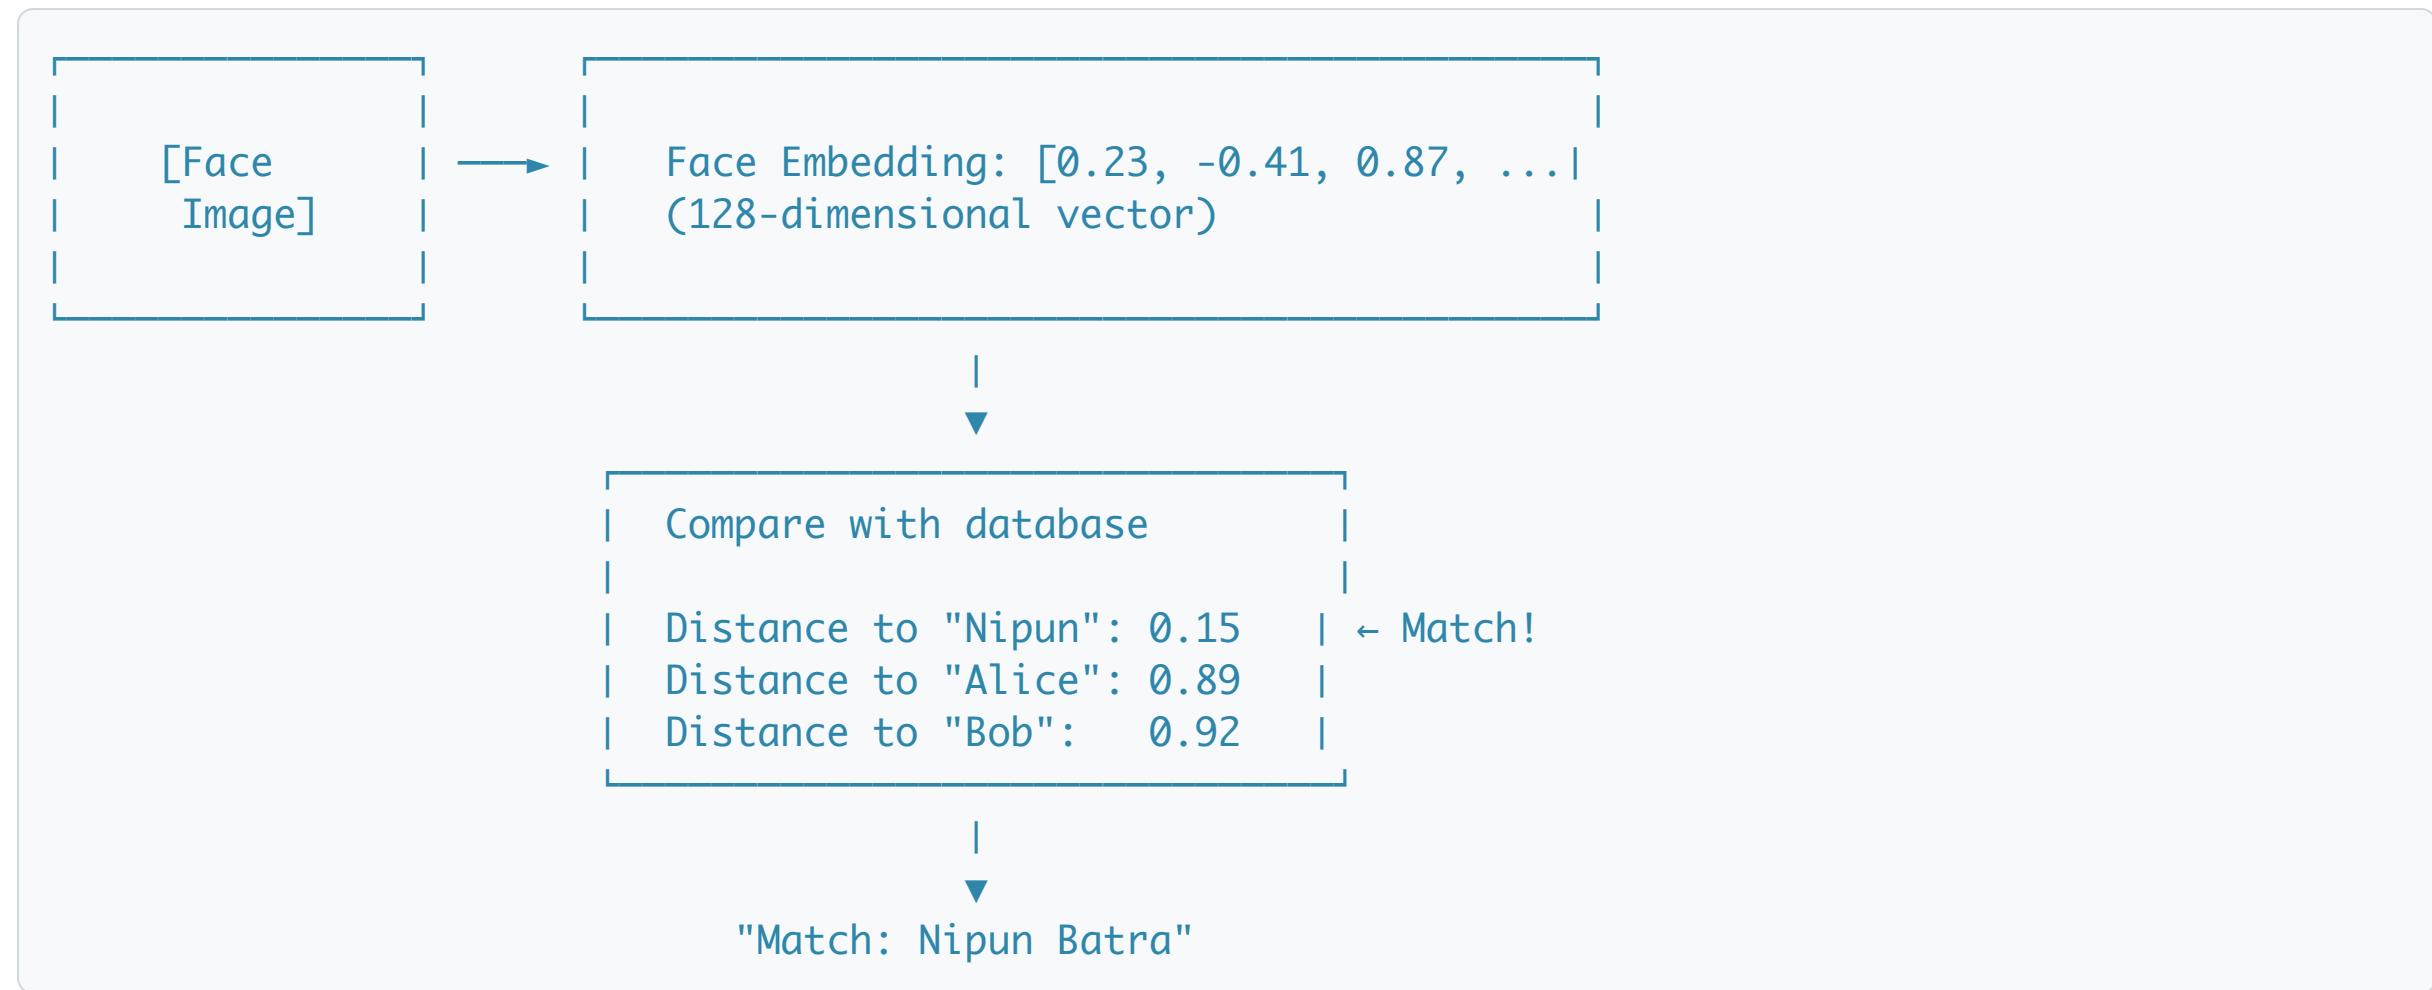
Objects moving towards you → collision warning
Everything moving left → you're turning right

VIDEO GAMES

Frame interpolation: turn 30fps into 60fps

Task 8: Face Recognition

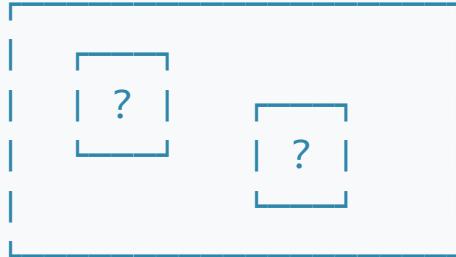
What: Identify WHO a face belongs to.



Face Detection ≠ Face Recognition

FACE DETECTION

Question: "Where are the faces?"

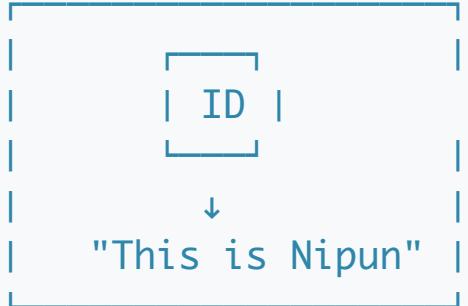


Output: Bounding boxes

Used BEFORE recognition
(find faces first)

FACE RECOGNITION

Question: "Who is this person?"



Output: Identity label

Requires database of known faces

Vision Tasks Summary

Task	Input	Output	Example Use
Classification	Image	Label	"Is this spam?"
Detection	Image	Boxes + labels	Self-driving cars
Segmentation	Image	Pixel mask	Medical imaging
Pose Estimation	Image	Keypoints	Fitness apps
Depth Estimation	Image	Depth map	AR furniture
Optical Flow	2 frames	Motion vectors	Video compression
Face Recognition	Face	Identity	Phone unlock

Domain 2: Natural Language Processing

Teaching Machines to Read & Write

"Language is the dress of thought." — Samuel Johnson

The NLP Task Landscape

NLP TASKS

UNDERSTANDING TASKS

- Sentiment Analysis
(Is this positive?)
- Named Entity Recognition
(Find names, places, dates)
- Question Answering
(Find the answer)
- Topic Classification
(Sports? Politics? Tech?)

GENERATION TASKS

- Text Generation
(Write like Shakespeare)
- Summarization
(Shorten this article)
- Translation
(English → Hindi)
- Paraphrasing
(Same meaning, new words)

MODERN LLMS CAN DO ALL OF THESE WITH A SINGLE MODEL!

Task 9: Sentiment Analysis

What: Classify text by emotion/opinion.

"This movie was absolutely amazing! Best film of the year. A masterpiece!"

▼
POSITIVE
(0.96)

Output Options:

Binary: Positive / Negative

3-class: Positive / Neutral / Negative

5-class: 1 to 5 stars

Continuous: -1.0 to +1.0
(very negative → very positive)

Sentiment Analysis: Real World

Brand Monitoring

Twitter Stream:

- + "Love the new iPhone!" -> Positive
- + "Battery dies so fast" -> Negative
- + "Just bought one!" -> Neutral
- + "Worst purchase ever" -> Negative
- + "Camera is incredible" -> Positive

Daily Sentiment: 67% positive

Customer Feedback

Support Tickets:

- | |
|------------------------|
| Urgent (Negative) ●●●● |
| Normal (Neutral) ●● |
| Praise (Positive) ● |

→ Route angry customers
to senior support!

Amazon analyzes millions of reviews using sentiment analysis to understand product reception!

The Challenge: Sarcasm & Context

WHY SENTIMENT IS HARD

SARCASM:

"Oh great, another software update. Just what I needed."

Words: positive ("great", "needed")

Meaning: NEGATIVE

NEGATION:

"This movie is not bad."

Contains "bad" → but overall POSITIVE

CONTEXT:

"The battery lasts forever" (phone review) → POSITIVE

"This movie lasts forever" (movie review) → NEGATIVE

MIXED:

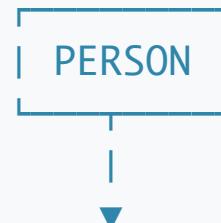
"The food was great but the service was terrible."

Overall? Positive? Negative? Neutral? Depends on priority!

Task 10: Named Entity Recognition (NER)

What: Find and label names, places, dates, organizations, etc.

Input: "Elon Musk announced that Tesla will open a factory
in Berlin by March 2025."



Output: "Elon Musk announced that Tesla will open a factory



in Berlin by March 2025."

NER: Entity Types

COMMON NER ENTITY TYPES

TYPE	EXAMPLES	COLOR CODE
PERSON	Elon Musk, Marie Curie	[Blue]
ORGANIZATION	Tesla, Google, UN	[Green]
LOCATION	Berlin, Mount Everest	[Yellow]
DATE	March 2025, last Tuesday	[Orange]
MONEY	\$5 million, €100	[Brown]
PERCENT	15%, three percent	[White]
TIME	3:30 PM, midnight	[Purple]
PRODUCT	iPhone 15, Model S	[Red]

Domain-specific:

- Medical: DISEASE, DRUG, SYMPTOM
- Legal: CASE_NUMBER, COURT, JUDGE
- Finance: TICKER, EXCHANGE, CURRENCY

NER: Real Applications

Search Engines

Query: "restaurants near Eiffel Tower"

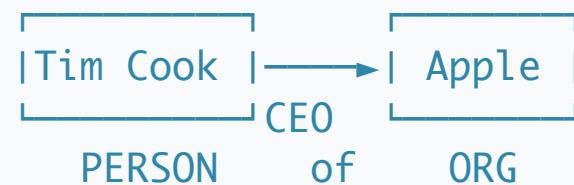
NER finds:

- LOCATION: "Eiffel Tower"
- Shows map of Paris
- Lists nearby restaurants

Knowledge Graphs

Text: "Tim Cook is the CEO of Apple"

Extracted:



Google Search uses NER to understand your queries and build its Knowledge Graph!

Task 11: Machine Translation

What: Convert text from one language to another.



Translation Challenges

WHY TRANSLATION IS HARD

WORD ORDER:

English: "I eat an apple" (Subject-Verb-Object)
Japanese: " はりんごを べる" (Subject-Object-Verb)

IDIOMS:

"It's raining cats and dogs" → Not about animals!
Must translate the MEANING, not words

CONTEXT:

"The bank is by the river"
- bank = financial institution? river bank?

CULTURAL CONCEPTS:

Some words have no direct translation
- Japanese " れ " (komorebi): sunlight filtering through trees

GENDER/FORMALITY:

"You" in English = tu/vous in French (formal vs informal)

Task 12: Text Summarization

Extractive: Pick important sentences verbatim.

Long Article:

- | Sentence 1 ← Keep
- | Sentence 2
- | Sentence 3
- | Sentence 4 ← Keep
- | Sentence 5
- | Sentence 6 ← Keep
- | ...

Just highlights!
No new words created.

Abstractive: Generate new text (paraphrase).

Long Article:

- | [Original 1000 words about climate change...]
- AI rewrites
- ▼

- | "Climate scientists warn that global temperatures..."
- | (100 words)

LLMs like GPT-4 and Claude do ****abstractive**** summarization — they truly understand and paraphrase!

Task 13: Question Answering

Extractive QA:

Find answer span in given text.

Context: "Albert Einstein was born in Ulm, Germany on March 14, 1879."

Question: "Where was Einstein born?"

Answer: "Ulm, Germany"



— Highlighted from the context text

Generative QA:

Generate free-form answer.

Question: "Explain quantum entanglement to a 5-year-old."

Answer: "Imagine you have two magic coins. When you flip one and it lands on heads, the other one ALWAYS lands on heads too, even if it's on the moon!"



— Created new text (not from any doc)

QA: The Evolution

QUESTION ANSWERING EVOLUTION

ERA 1: Rule-Based (1960s-2000s)

Keywords → Database lookup → Template answer
"Very brittle, only worked for specific domains"

ERA 2: Extractive (2016-2020)

BERT-style: "Find the answer IN the text"
Great for reading comprehension tasks

ERA 3: Generative (2020-present)

LLMs: Generate answers from learned knowledge
Can answer questions about ANYTHING
Can reason, explain, and elaborate

ERA 4: RAG (Retrieval-Augmented Generation)

LLM + Search = Best of both worlds
Accurate, up-to-date, with sources

Task 14: Text Generation (LLMs)

What: Predict and generate the next tokens, one at a time.

Prompt: "The secret to happiness is"



LLM

(GPT, Claude,
Gemini, etc.)



Token 1: "not" $P(\text{"not"} \mid \text{prompt}) = 0.15$



Token 2: "in" $P(\text{"in"} \mid \text{prompt} + \text{"not"}) = 0.42$



Token 3: "wealth" $P(\text{"wealth"} \mid \dots + \text{"in"}) = 0.23$



... continues until <END> token

Output: "The secret to happiness is not in wealth but in
meaningful connections with others."

How LLMs Generate Text

AUTOREGRESSIVE GENERATION

Step 1: "The" → predict next → "secret"
Step 2: "The secret" → predict next → "to"
Step 3: "The secret to" → predict next → "happiness"
...

Each step:

Probability Distribution Over Vocabulary

P("the")	= 0.02	
P("happiness")	= 0.25	
P("success")	= 0.18	
P("life")	= 0.12	
P("love")	= 0.08	
...	= ...	

Sample from this distribution (or take argmax)

NLP Tasks Summary

Task	Input	Output	Key Challenge
Sentiment	Text	Positive/Negative	Sarcasm, context
NER	Text	Entity spans + types	Ambiguous names
Translation	Text (lang A)	Text (lang B)	Word order, idioms
Summarization	Long text	Short text	Keeping key info
QA	Question + context	Answer	Finding relevant info
Generation	Prompt	Continued text	Coherence, factuality

Domain 3: Audio & Speech

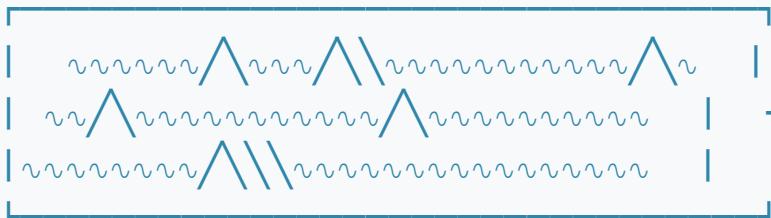
Teaching Machines to Hear

"The human voice is the most beautiful instrument of all." — Arvo Pärt

Task 15: Speech-to-Text (ASR)

What: Convert spoken audio to text.

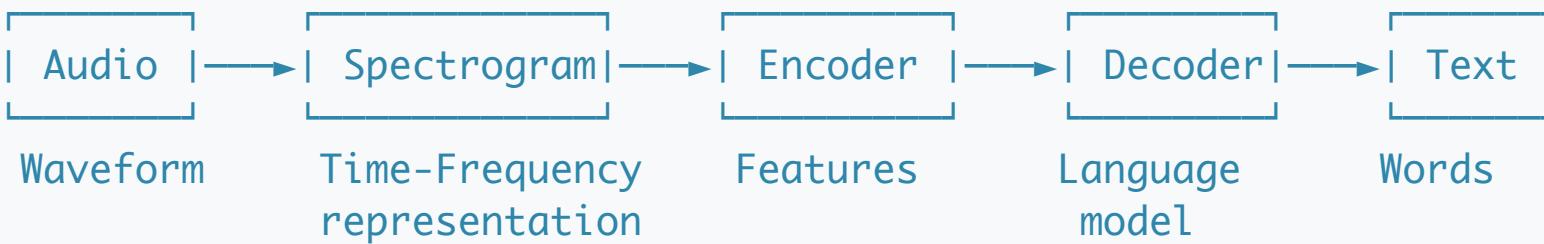
Audio Waveform:



Text Output:

"Hello, how are
you today?"

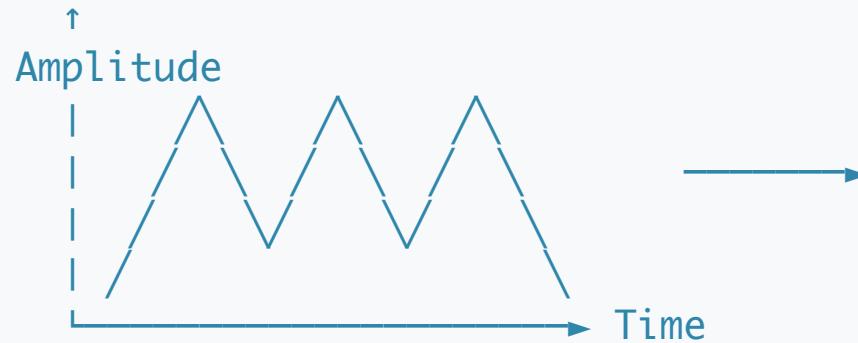
Pipeline:



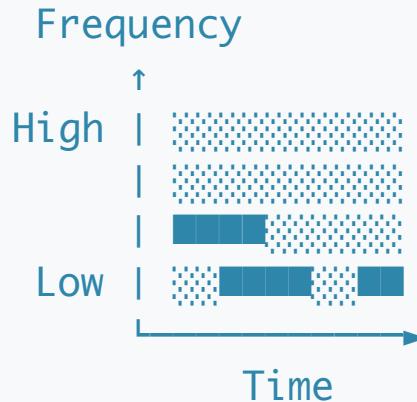
ASR: The Spectrogram

Converting sound to "images" that neural networks can process

Audio Wave:



Spectrogram:
(time → frequency "image")



Bright areas = loud frequencies at that moment
Pattern = unique "fingerprint" of each word!

Whisper by OpenAI can transcribe audio in 99 languages with near-human accuracy!

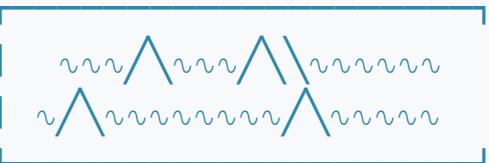
Task 16: Text-to-Speech (TTS)

What: Convert text to natural-sounding audio.

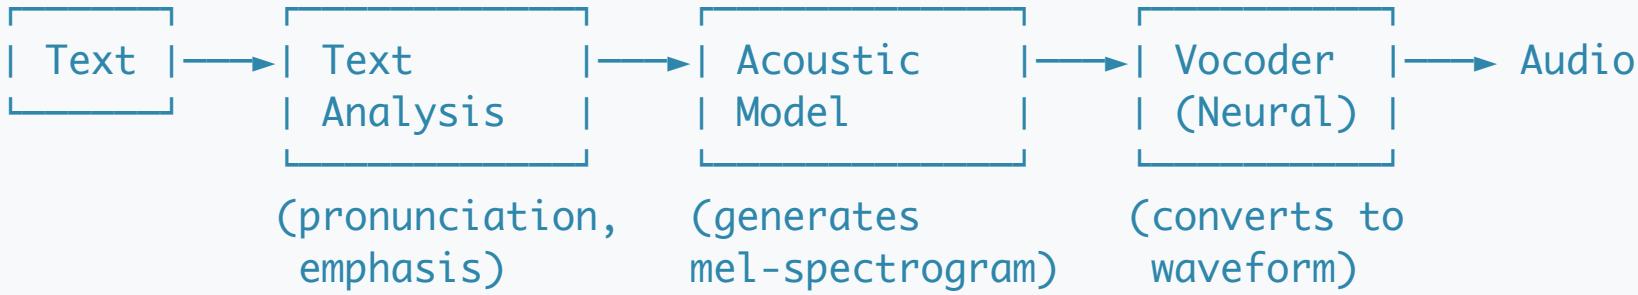
Text Input:

```
| "Welcome to the  
| future of AI.  
| This is exciting!"
```

Audio Output:



Modern TTS Pipeline:



TTS: Then vs Now

TTS EVOLUTION

1990s: Concatenative TTS

Splice together recorded phonemes

Result: Robotic, unnatural "The-wea-ther-to-day-is..."

2010s: Statistical Parametric TTS

HMM-based models, smoother but still artificial

2016: WaveNet (DeepMind)

Neural network generates audio sample by sample

Human-like quality, but VERY slow

2020s: Parallel Neural TTS

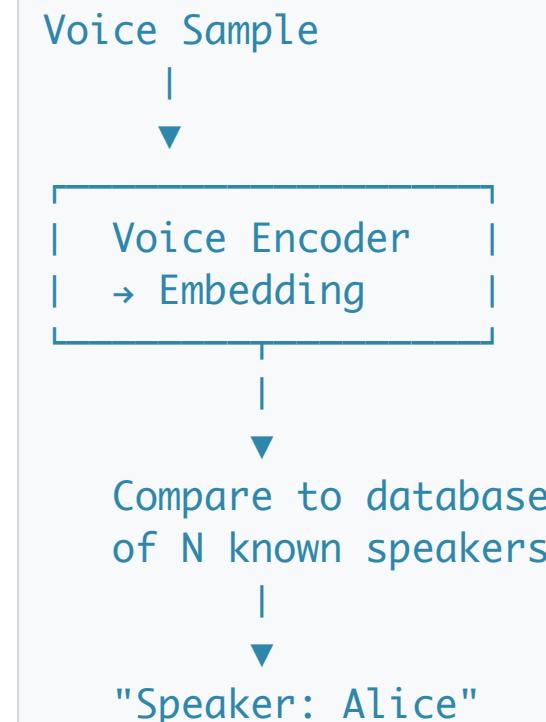
Real-time, expressive, can clone voices

Examples: ElevenLabs, Bark, XTTs

Task 17-18: Speaker Recognition

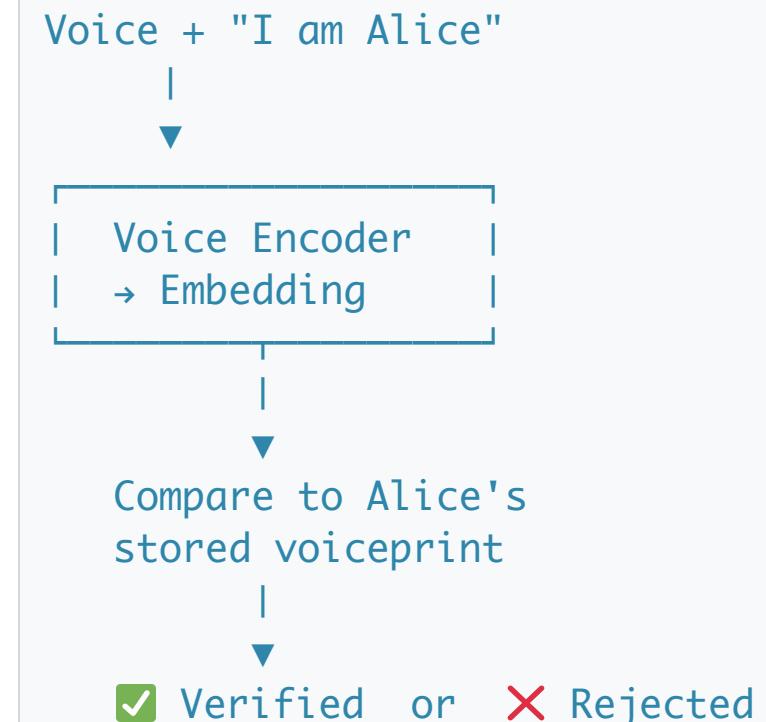
Speaker Identification:

Who is speaking? (1-of-N)



Speaker Verification:

Is this who they claim to be?



"Hey Siri" uses **speaker verification** — it only responds to the device owner's voice!

Domain 4: Unsupervised Learning

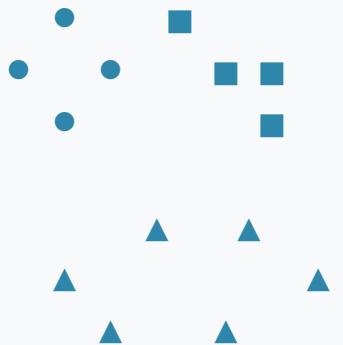
Finding Patterns Without Labels

"The goal is to find structure in chaos."

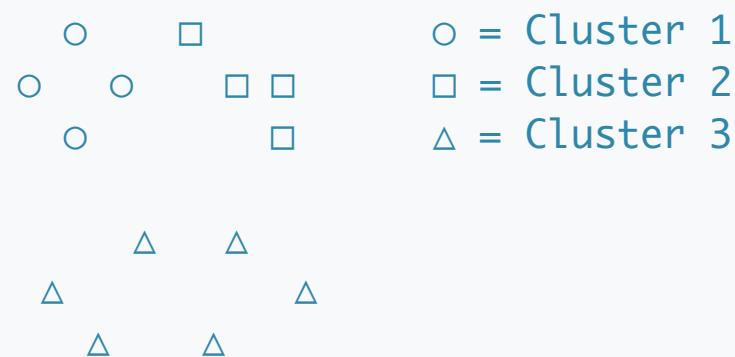
Task 19: Clustering

What: Group similar items together automatically (no labels needed!).

Before (unlabeled data):



After (3 clusters found):



○ = Cluster 1
□ = Cluster 2
△ = Cluster 3

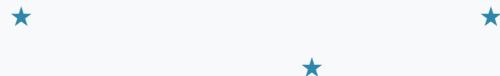
Algorithm (K-Means) figures out:

- There are 3 natural groups
- Which points belong to which group

Clustering: The K-Means Algorithm

K-MEANS: STEP BY STEP

Step 1: Pick K random "centroids" (cluster centers)



Step 2: Assign each point to nearest centroid

ooo near ★₁ □□□ near ★₂ △△△ near ★₃

Step 3: Move centroids to center of their points

★₁ moves to average of ooo

★₂ moves to average of □□□

★₃ moves to average of △△△

Step 4: Repeat steps 2-3 until centroids stop moving

Done! Points are now clustered.

Clustering: Real Applications

Customer Segmentation

Cluster 1: "VIPs"

- └ High spending
- └ Infrequent visits
- └ Premium products

Cluster 2: "Regulars"

- └ Medium spending
- └ Weekly visits
- └ Staple items

Cluster 3: "Bargain Hunters"

- └ Low spending
- └ Sale days only
- └ Discounted items

Image Compression

Original: 16 million colors

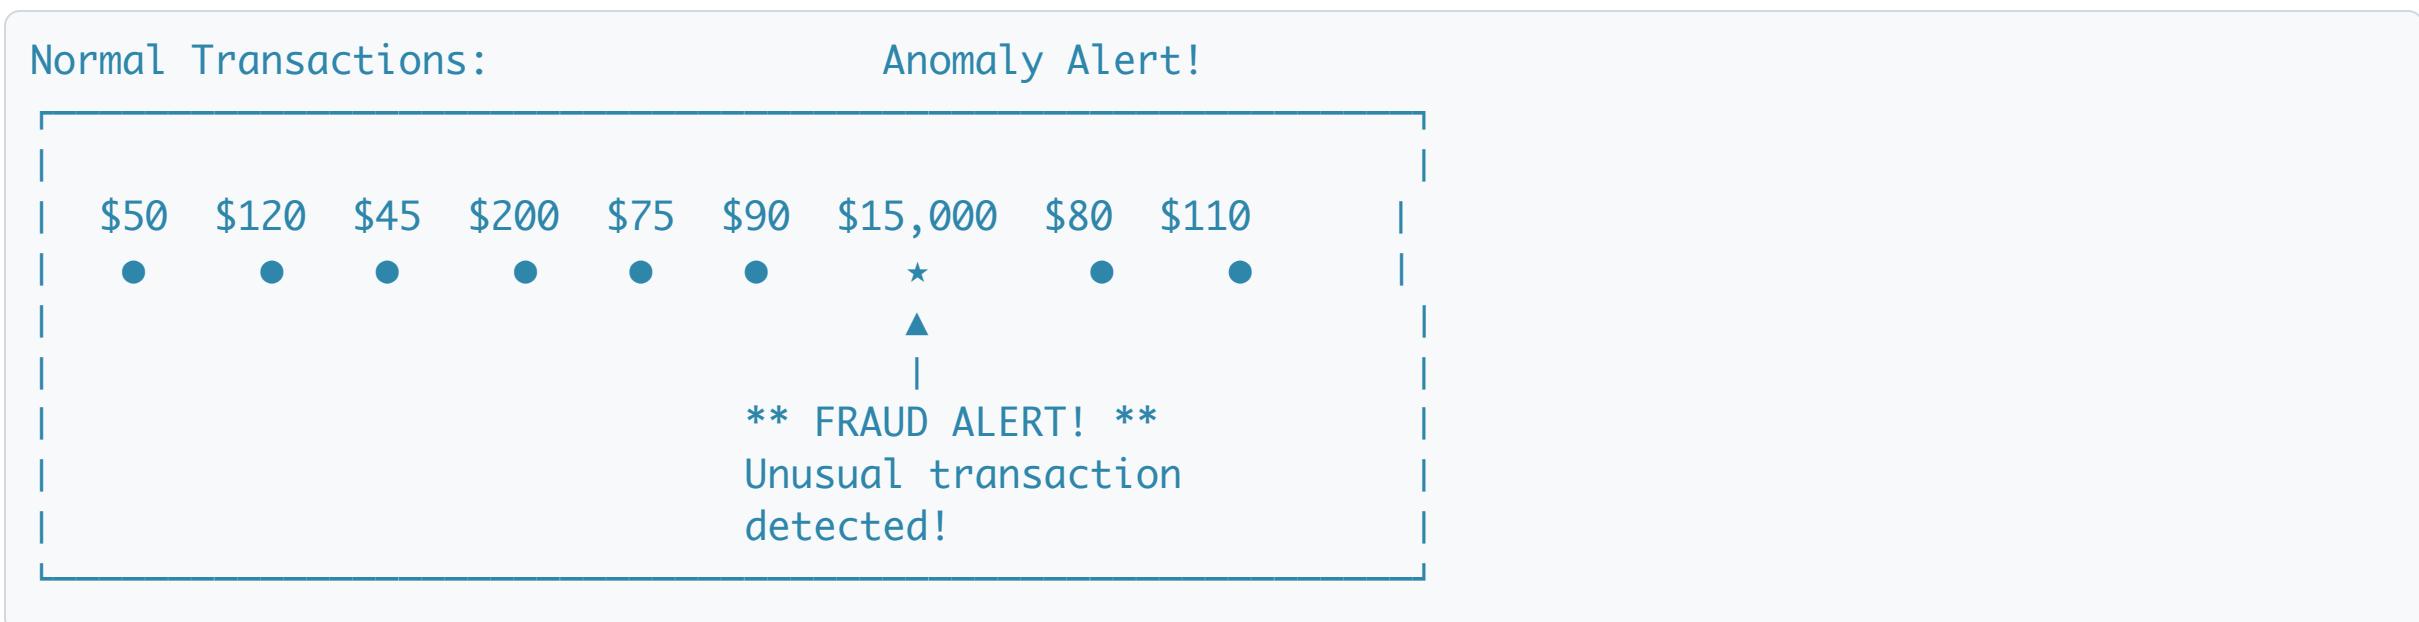
After K-Means (K=16):
Only 16 colors!

| [Color Palette] |
| and 8 more shades |

File size: 10x smaller
Quality: Still looks good!

Task 20: Anomaly Detection

What: Find the outliers / unusual patterns.



Anomaly detection is "learning what's normal, then flagging what's not."

Anomaly Detection: Methods

ANOMALY DETECTION APPROACHES

1. STATISTICAL

If value is > 3 standard deviations from mean \rightarrow Anomaly
Simple but assumes normal distribution

2. DISTANCE-BASED

If point is far from all other points \rightarrow Anomaly
Works for any shape of data

3. DENSITY-BASED

If point is in a low-density region \rightarrow Anomaly
Good for varying cluster sizes

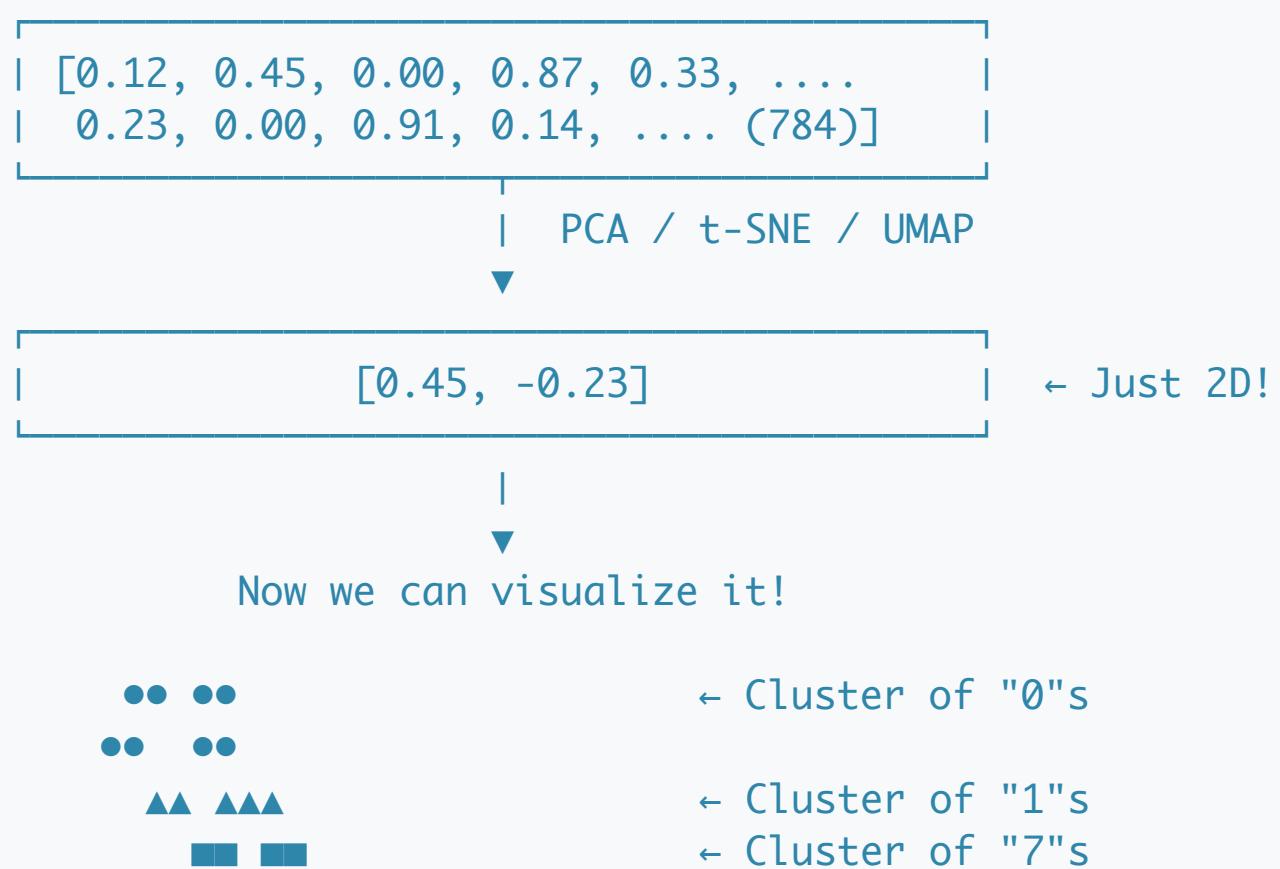
4. AUTOENCODER (Neural Network)

Train to reconstruct normal data
High reconstruction error \rightarrow Anomaly
Best for complex, high-dimensional data

Task 21: Dimensionality Reduction

What: Compress high-dimensional data while preserving structure.

Original: 784 dimensions (28x28 MNIST image)



Why Reduce Dimensions?

BENEFITS OF DIMENSIONALITY REDUCTION

1. VISUALIZATION

Can't plot 784D data, but can plot 2D!
Reveal clusters and patterns to humans

2. SPEED

ML on 10 features is 100x faster than 1000 features
Less computation, less memory

3. NOISE REMOVAL

Lower dimensions often capture signal, remove noise
Can improve model accuracy!

4. THE CURSE OF DIMENSIONALITY

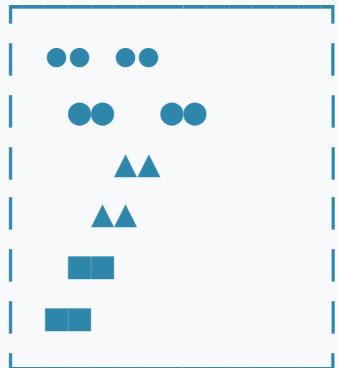
As dimensions \uparrow , distance between points \rightarrow same
Need exponentially more data in high dimensions

PCA vs t-SNE vs UMAP

Method	Speed	Preserves	Best For
PCA	Very fast	Global structure	Initial exploration
t-SNE	Slow	Local clusters	Visualizing clusters
UMAP	Fast	Both local + global	Best overall

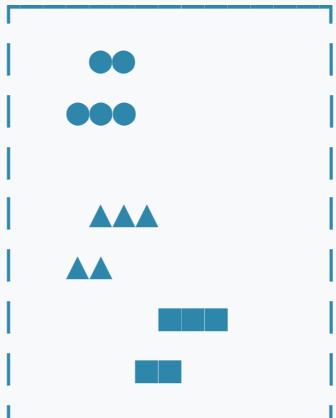
Same data, different methods:

PCA:



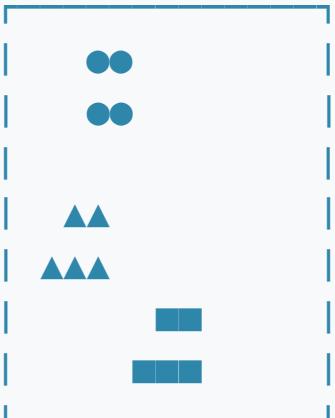
Linear projection

t-SNE:



Clusters tight

UMAP:



Clusters + structure

Domain 5: Generative Models

Creating New Content

"Creativity is just connecting things." — Steve Jobs

The Generative AI Revolution

GENERATIVE AI TIMELINE

2014: GANs invented
First realistic image generation

2020: GPT-3 launches
Text generation goes mainstream

2022: DALL-E 2, Stable Diffusion, Midjourney
Anyone can generate images from text

2022: ChatGPT launches
100M users in 2 months (fastest ever)

2023: GPT-4, Claude, Gemini
Multimodal: text + images + code

2024: Sora (video), Suno (music)
Generate any media type from text

Task 22: Image Generation

What: Create new images from noise, text, or other images.

Text-to-Image (Stable Diffusion, DALL-E, Midjourney):

Prompt: "A robot painting
a sunset, oil
painting style"



Generated Image:

[AI Generated Art]
[Beautiful AI
generated art]

Noise-to-Image (GANs, Diffusion):

Random Noise → Generator → Realistic Image
[$z \sim N(0,1)$]

How Diffusion Models Work

DIFFUSION: THE CORE IDEA

TRAINING (Forward process):

Take image → gradually add noise → pure noise

[Img] → [Img+noise] → [Img+noise] → [Img+noise] → [Noise]
Clean image Slight noise Medium noise Heavy noise Pure noise

GENERATION (Reverse process):

Start with noise → gradually denoise → clean image

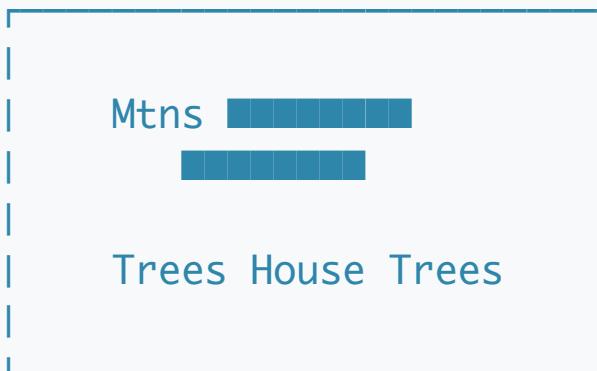
[Noise] → [Img+noise] → [Img+noise] → [Img+noise] → [Img]
Pure noise Heavy noise Medium noise Slight noise Clean image!

The model learns: "Given noisy image, predict the noise"

Task 23: Image Inpainting

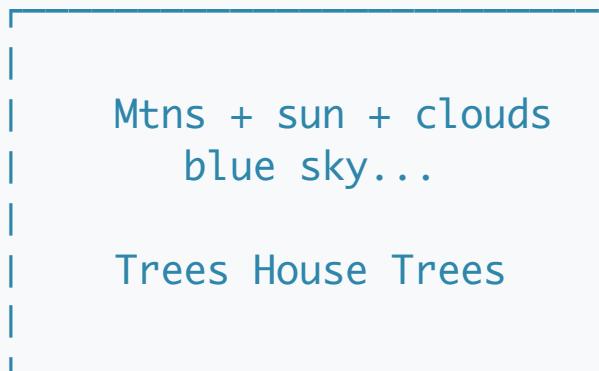
What: Fill in missing or masked regions intelligently.

Original with hole:



(user painted a mask)

Inpainted result:



(AI filled it in)

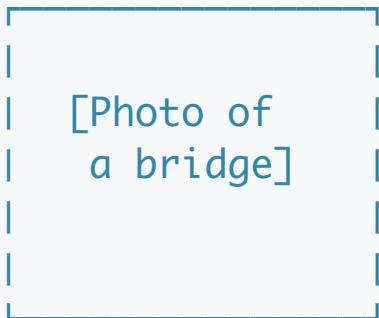
Uses:

- Remove unwanted objects (photobombers!)
- Restore damaged/old photos
- Extend image boundaries (outpainting)

Task 24: Style Transfer

What: Apply the artistic style of one image to the content of another.

Content Image:

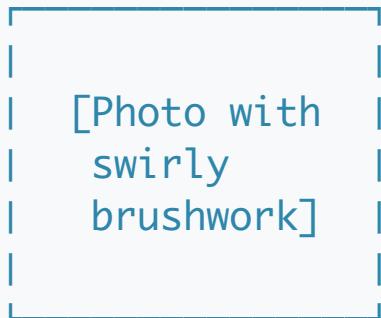


Style Image:



+

Result:



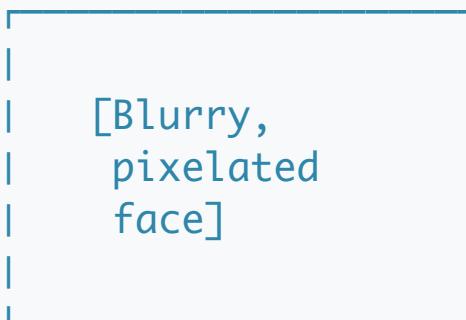
The model learns to separate:

- CONTENT: What objects are in the image (bridge, sky, water)
- STYLE: How they're rendered (brushstrokes, colors, texture)

Task 25: Super Resolution

What: Upscale low-resolution images while adding realistic detail.

Low Resolution (64x64):



High Resolution (512x512):

Traditional upscaling: Just makes pixels bigger (still blurry)

AI upscaling: Adds plausible detail that wasn't there!

****Ethical note:**** Super resolution "hallucinates" detail. The added details are plausible but not necessarily accurate. Not suitable for forensics or legal evidence!

Domain 6: Self-Supervised Learning

The Secret Sauce of Modern AI

"Give me a lever long enough and I shall move the world." — Archimedes

The Self-Supervised Revolution

THE LABELING BOTTLENECK

SUPERVISED LEARNING:

Need millions of labeled examples
Labeling is EXPENSIVE and SLOW
Limited by human annotation capacity

ImageNet: 14 million images
Cost: ~\$500,000 and years of work!

SELF-SUPERVISED LEARNING:

Create labels FROM the data itself
"Free" supervision from data structure
Can use BILLIONS of examples

GPT-3: 45 TB of text
Cost: Compute only (no labeling!)

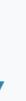
Task 26: Masked Language Modeling (BERT-style)

What: Predict the hidden word(s) — fill in the blank.

Training example:

Original: "The cat sat on the mat."

Masked: "The cat sat on the [MASK]."



Predictions: "mat" (0.45) ← Correct!

"floor" (0.22)

"couch" (0.15)

"bed" (0.08)

...

BERT was trained on **3.3 billion words** from Wikipedia and books, just playing fill-in-the-blank!

Task 27: Next Token Prediction (GPT-style)

What: Predict what comes next in a sequence.

Input: "The capital of France is"



GPT



Next token distribution:

"Paris" (0.89) ← Most likely

"the" (0.03)

"in" (0.02)

"a" (0.01)

...

GPT, Claude, Gemini, and all LLMs are trained with just this one objective — repeated trillions of times! The simplicity is the brilliance.

BERT vs GPT: Key Differences

BERT vs GPT

BERT (Bidirectional)

"The [MASK] sat on mat"

↑↓↑↓↑

Looks both directions

Best for:

- Understanding text
- Classification
- Question answering
- Named entity recognition

Can't generate text well
(doesn't predict in order)

GPT (Autoregressive)

"The cat sat on" → "the" → "mat"



Only looks backward

Best for:

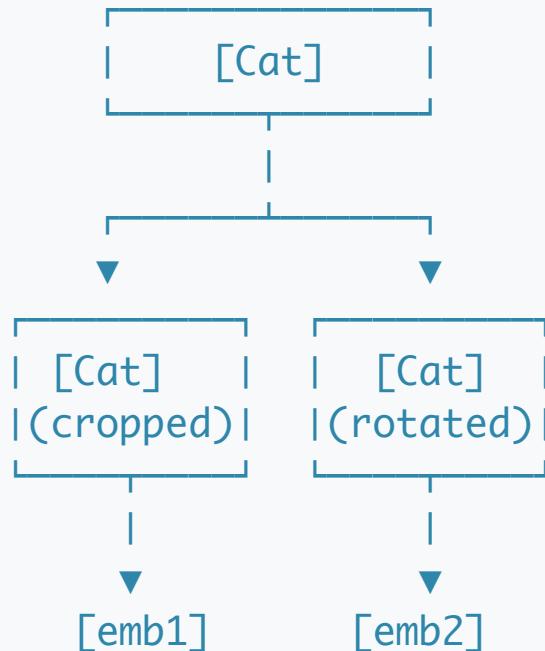
- Generating text
- Chat/dialogue
- Code completion
- Creative writing

Can generate, but slower
(one token at a time)

Task 28: Contrastive Learning

What: Learn that different views of the same image should have similar embeddings.

Original Image:



Create different "views"
(augmentations)

These should be SIMILAR!

Meanwhile: embeddings of DIFFERENT images should be DIFFERENT!

[cat_emb] ←————→ [dog_emb]
Push apart!

Contrastive Learning: The Big Picture

CONTRASTIVE LEARNING

Key Insight:

We don't need labels to learn good representations!
Just need to know: "These are the same" vs "These are different"

Training:

- Take image → create 2 augmented versions (positive pair)
- Other images in batch = negative pairs
- Learn: positives close, negatives far

Result:

An encoder that maps similar images to similar embeddings
Can then use these embeddings for ANY downstream task!
Often matches supervised learning with just 1% of labels!

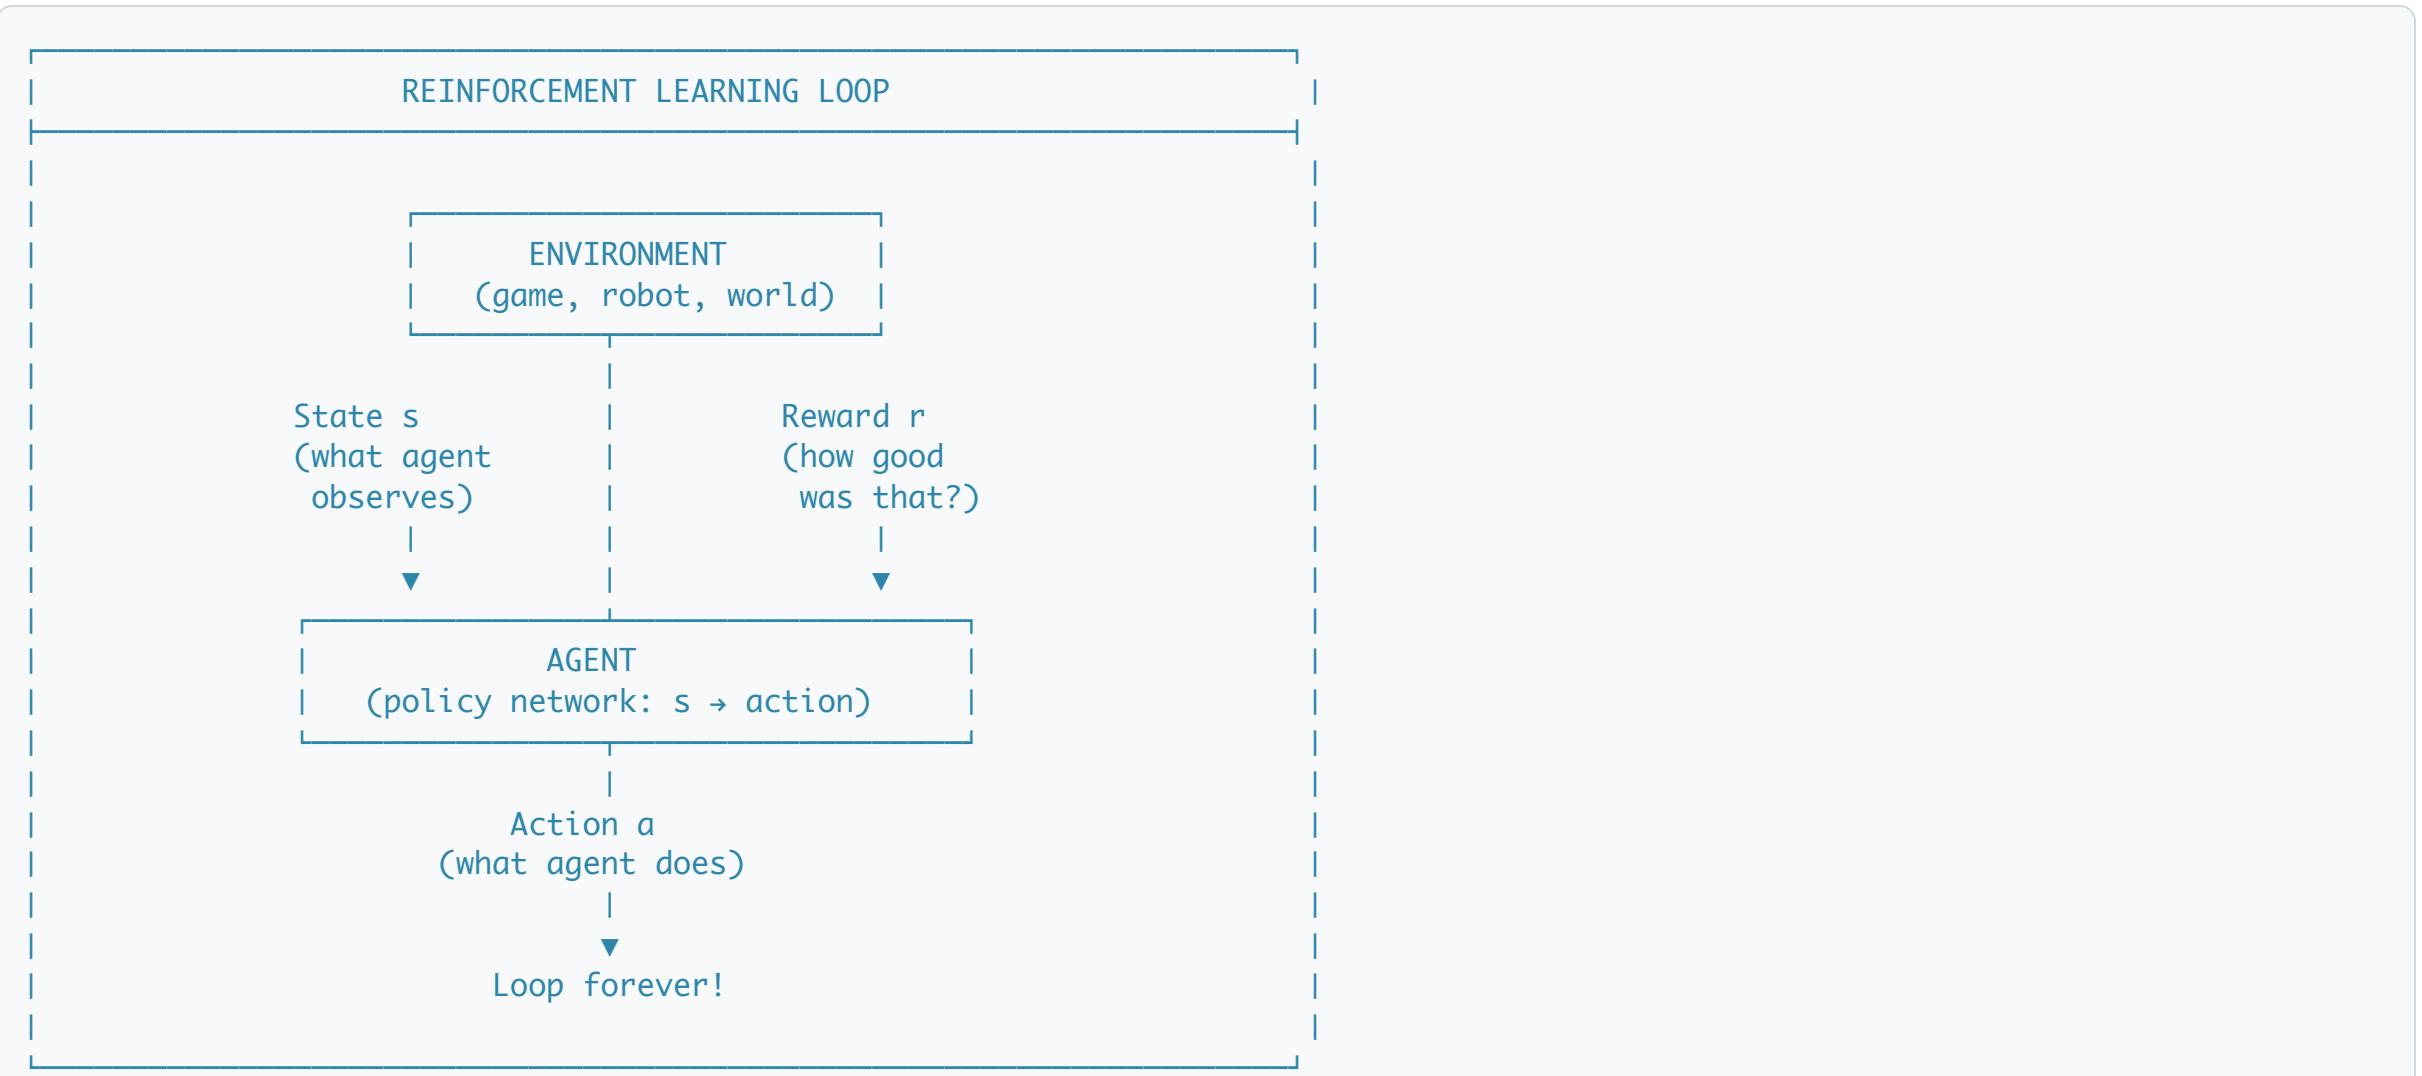
Famous methods: SimCLR, MoCo, CLIP

Domain 7: Reinforcement Learning

Learning by Doing

"Experience is the teacher of all things." — Julius Caesar

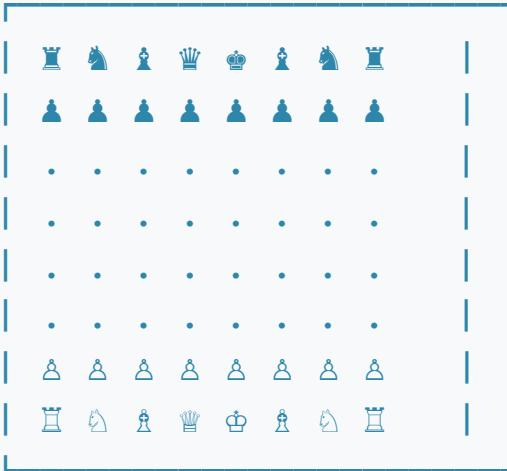
The RL Framework



Task 29: Game Playing

What: Learn optimal strategy through trial and error.

Game State (Chess):



Agent Decision:

Best move:
 $e2 \rightarrow e4$

Evaluation:
+0.3 pawns

AlphaGo/AlphaZero: Learned by playing MILLIONS of games against itself!
No human games needed – pure self-play!

RL Milestones in Games

RL GAME-PLAYING ACHIEVEMENTS

1992: TD-Gammon

Backgammon at world champion level
First major RL success!

2013: DQN (Atari)

Superhuman at 29 Atari games
Raw pixels as input!

2016: AlphaGo

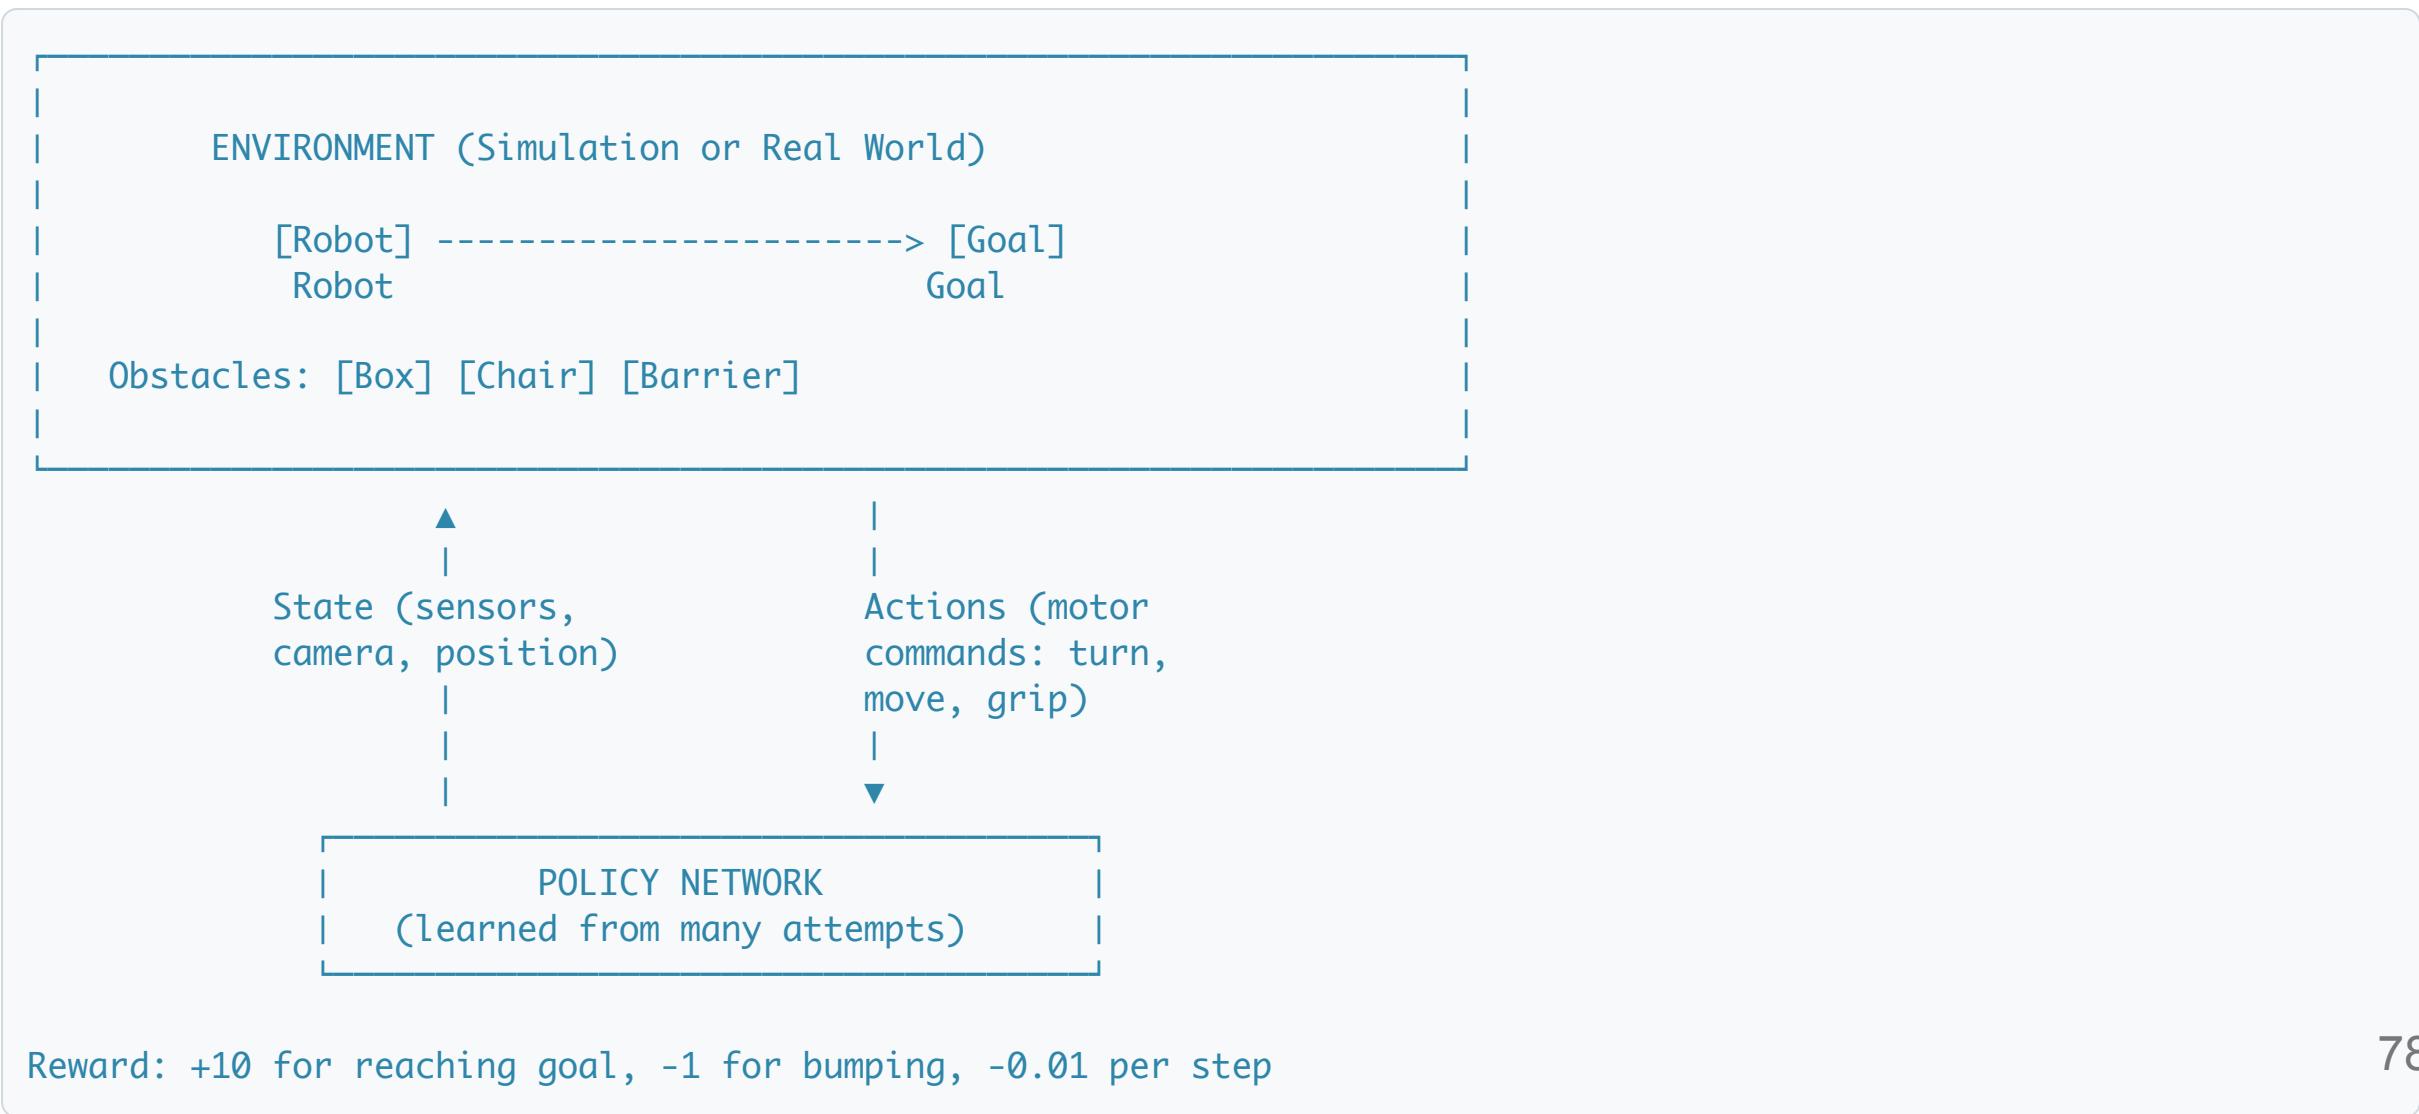
Beats world champion Lee Sedol at Go
Game with 10^{170} possible positions!

2019: AlphaStar

Grandmaster level at StarCraft II
Real-time, imperfect information, complex strategy

Task 30: Robot Control

What: Learn to move and interact in the physical world.



Real-World RL Applications

Robotics

- Robot manipulation (picking objects)
- Drone navigation
- Self-balancing robots
- Walking robots (Boston Dynamics)

Games

- Video game AI
- Board game engines
- Game testing automation

Beyond Games

- **Data center cooling** (Google: 40% energy savings)
- **Chip design** (Google, NVIDIA)
- **Trading** (quantitative finance)
- **Recommendations** (long-term engagement)
- **RLHF** (making LLMs helpful & safe!)

ChatGPT uses RLHF (Reinforcement Learning from Human Feedback) to learn to be helpful rather than just predicting text!

Domain 8: Multimodal Tasks

Combining Vision + Language

"The whole is greater than the sum of its parts." — Aristotle

The Multimodal Revolution

MULTIMODAL = MULTIPLE MODALITIES

MODALITY = Type of data

- Text (words, sentences)
- Images (photos, diagrams)
- Audio (speech, music)
- Video (sequences of images + audio)

MULTIMODAL MODELS:

- GPT-4V: Text + Images
- Gemini: Text + Images + Audio + Video
- CLIP: Connects text and images
- Claude: Text + Images + Documents

The trend: One model to rule them all!

Task 31: Visual Question Answering (VQA)

What: Answer questions about images using both visual and language understanding.

Image:



Questions & Answers:

Q: "How many people are in the image?"
A: "Two people"

Q: "What animal is there?"
A: "A dog"

Q: "What are they doing?"
A: "Walking their dog in a park"

Requires BOTH:

- Understanding image content (computer vision)
- Understanding question (NLP)
- Reasoning to connect them!

Task 32: Image Captioning

What: Generate natural language description of an image.

Image:

Generated Caption:

[Runners]

--> "A group of runners participating in a city marathon on a sunny day, with spectators cheering along the street and tall buildings in the background."

[Marathon scene
with crowds and
city buildings]

[Crowds][Buildings]

The inverse of VQA:

Instead of answering questions, generate descriptions!

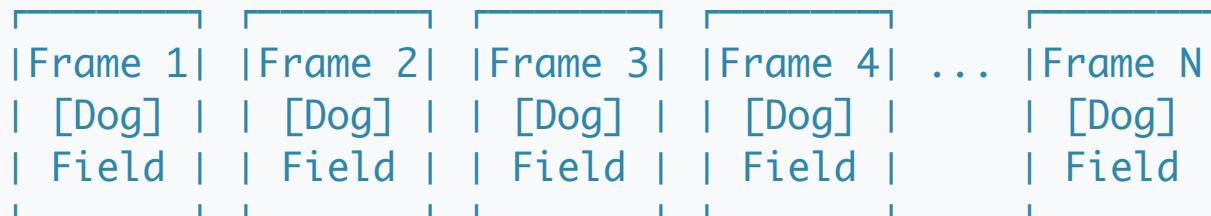
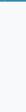
Task 33: Text-to-Video

What: Generate video from text description.

Prompt: "A golden retriever running through a field of sunflowers on a sunny day, slow motion"



Video Model
(Sora, Runway,
Pika, etc.)

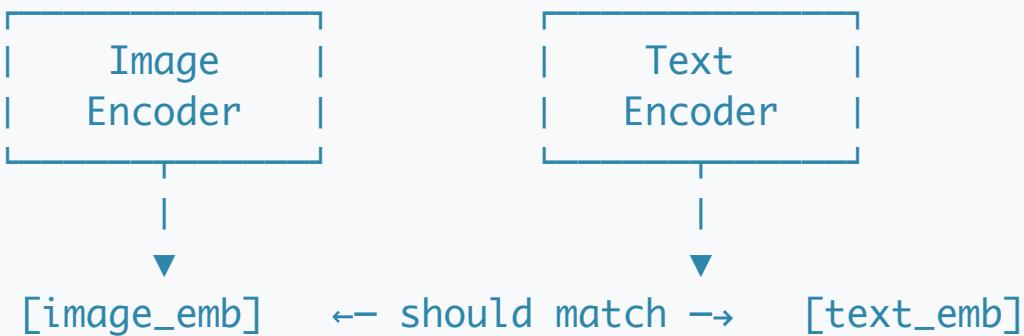


Temporally consistent video!

CLIP: The Foundation of Multimodal AI

CLIP: Connecting Text and Images

Training Data: 400 million (image, caption) pairs from internet



"A dog playing
in the snow" → [0.23, -0.41, 0.87, ...]
Shared embedding space!

Result: Can search images with text, or text with images!
Powers: DALL-E, Stable Diffusion, image search, and more

Domain 9: Tabular & Time Series

The Classic ML Tasks

"Not everything that counts can be counted, but data often helps."

Task 34-35: Regression & Classification on Tables

Tabular Regression:

Beds	SqFt	Price?
3	1500	???
4	2200	???
2	900	???

Predict: \$425,000

Output: Continuous number

Tabular Classification:

Age	Income	Default?
35	75K	???
52	120K	???
28	45K	???

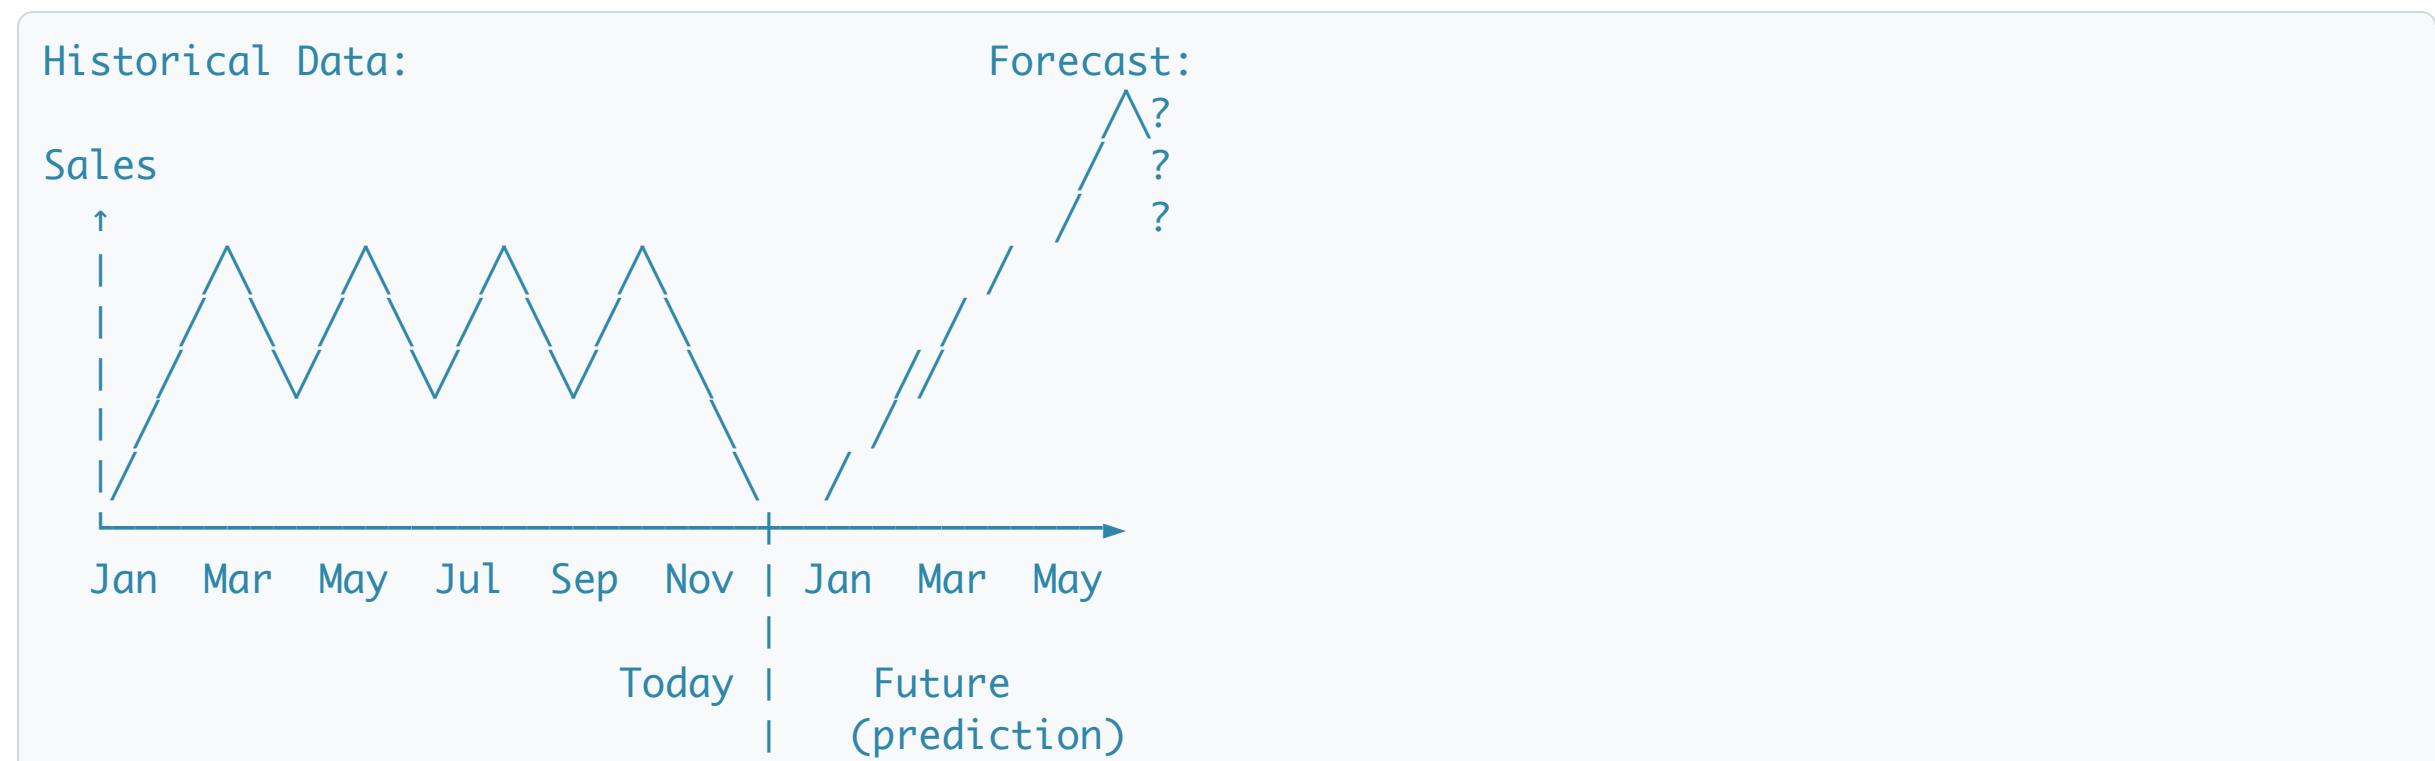
Predict: Yes / No

Output: Category

For tabular data, **gradient boosting (XGBoost, LightGBM)** often beats deep learning! Simpler, faster, and more interpretable.

Task 36: Time Series Forecasting

What: Predict future values from historical patterns.



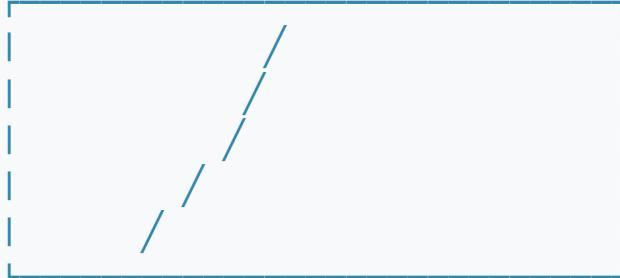
Components to model:

- Trend (overall direction)
- Seasonality (repeating patterns)
- Noise (random variation)

Time Series: Key Patterns

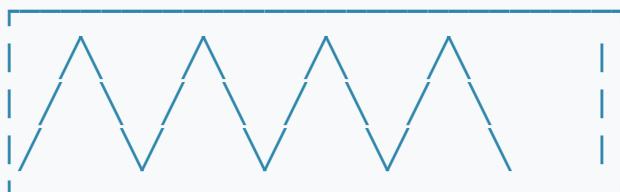
TIME SERIES COMPONENTS

TREND: Long-term direction



Upward trend
(e.g., company growth)

SEASONALITY: Repeating patterns



Weekly, monthly, yearly
(e.g., holiday shopping)

Task 37: Recommendation Systems

What: Predict what users will like based on their history.

User-Item Matrix:

	Movie1	Movie2	Movie3	Movie4		
User A	5	?	3	?	For User A:	
User B	4	5	?	2	* Movie2 (pred: 4.2)	
User C	?	4	5	3	* Movie4 (pred: 3.8)	

Recommendations:

| "Because you liked
| Movie1 and Movie3"

Two main approaches:

- Collaborative: "Users like you also liked..."
- Content-based: "Similar items to ones you liked..."

Recommendation: The Netflix Problem

RECOMMENDATION CHALLENGES

THE COLD START PROBLEM:

New user: No history → What to recommend?

New item: No ratings → Who might like it?

Solution: Ask preferences, use demographics, popular items

THE SPARSITY PROBLEM:

Netflix: 200M users × 15K movies = 3 trillion possible ratings

Actual ratings: ~5 billion → 0.17% filled!

Solution: Matrix factorization, embeddings

THE FILTER BUBBLE:

Only showing similar content → User misses diverse content

Solution: Exploration vs exploitation, diversity metrics

Summary: The ML Task Landscape

ML TASK FAMILIES

SUPERVISED

- Classification
- Regression
- Detection
- Segmentation
- Seq2Seq

UNSUPERVISED

- Clustering
- Dim. Reduction
- Anomaly Det.

SELF-SUPERVISED

- Masked LM (BERT)
- Next Token (GPT)
- Contrastive (CLIP)

GENERATIVE

- Image Gen
- Text Gen
- Video Gen
- Style Transfer

REINFORCEMENT

- Game Playing
- Robotics
- RLHF

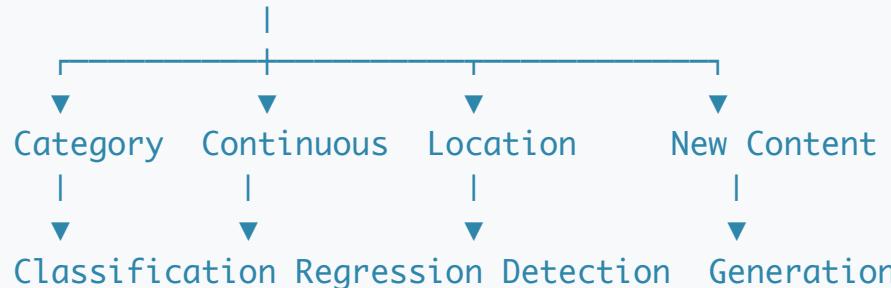
MULTIMODAL

- VQA
- Captioning
- Text-to-Image
- Text-to-Video

Choosing the Right Task

DECISION FLOWCHART

What do you want to predict?



Have labels?

- └ Yes → Supervised
- └ No → Unsupervised (clustering, anomaly)
- └ Can create from data → Self-supervised

Need to take actions?

- └ Yes → Reinforcement Learning

Multiple input types?

- └ Yes → Multimodal

Key Takeaways

The 5 Things to Remember

1. **Every task = Input type + Output type**

Define these clearly and you've defined your problem

2. **Same architectures work across domains**

Transformers power text, images, audio, and more

3. **Self-supervised learning powers modern AI**

GPT, BERT, CLIP — all learned from unlabeled data

4. **Start with the task → then choose the model**

Don't pick a model first and force it to fit

5. **Real-world ML often combines multiple tasks**

Self-driving cars: detection + segmentation + prediction + control

The ML Practitioner's Toolkit

Task Type	Go-To Models (2024)
Image Classification	ResNet, EfficientNet, ViT
Object Detection	YOLOv8, DETR, RT-DETR
Segmentation	SAM, Mask R-CNN
Text Classification	BERT, RoBERTa
Text Generation	GPT-4, Claude, Llama
Speech-to-Text	Whisper
Image Generation	Stable Diffusion, DALL-E
Tabular	XGBoost, LightGBM, TabNet
Time Series	Prophet, N-BEATS, TimeGPT
Recommendations	Two-Tower, DLRM

What's Next?

In the Labs:

- Lab 1-2: sklearn basics
- Lab 3: PyTorch & neural nets
- Lab 4-5: Build your own LLM
- Lab 6-7: Object detection

The Bigger Picture:

- Most tasks share core principles
- Transfer learning is key
- Start simple, add complexity
- The best model is the one you ship!

Pick a task, find a dataset, and start building! The best way to learn ML is by doing.

Thank You!

"The best way to predict the future is to invent it." — Alan Kay

Resources

- Papers With Code (paperwithcode.com) — State-of-the-art models
- Hugging Face (huggingface.co) — Pre-trained models
- Kaggle (kaggle.com) — Datasets and competitions

Questions?

