

# Next Token Prediction

Building ChatGPT from Scratch (Conceptually)

Nipun Batra

*IIT Gandhinagar*

# Overview of the Journey

Today we will build a Language Model from scratch.

1. **The "Autocomplete" Game:** Intuition.
2. **Level 1: Bigrams:** Just counting letters.
3. **Level 2: Embeddings:** Representing meaning.
4. **Level 3: Neural Networks (MLP):** Learning patterns.
5. **Level 4: RNNs:** Adding memory (briefly).
6. **Level 5: Transformers:** The "Attention" Revolution.
7. **Real World:** How ChatGPT actually works.

# **Part 1: The Intuition**

# The Core Problem

Deep Learning, at its heart, often answers one simple question:

**"Given what I have seen so far, what comes next?"**



This is a **Sequence Modeling** problem.

# The "Autocomplete" Game

Let's play. Guess the next word:

**Sequence:** "New Delhi is the capital of \_\_\_\_"

**Your Brain:** "India" (100% confidence)

**Sequence:** "I want to eat \_\_\_\_"

**Your Brain:** "Pizza"? "Apples"? (Uncertainty)

| **Key Insight:** Language modeling is about assigning **probabilities** to the next token.

# It's Just Prediction!

You might think ChatGPT "knows" physics or history.

Actually, it just predicts the next word.

- **Input:** "F = m"
- **Prediction:** "a" (Newton's Law)
- **Input:** "To be or not to"
- **Prediction:** "be" (Shakespeare)

If you predict well enough, you **appear** intelligent.

## Part 2: The "Counting" Era (Bigrams)

# The Simplest Model: Bigrams

Imagine we want to generate Indian names like `aabid`, `zeel`, etc.

Let's look at the dataset: `["aabid", "zeel"]`.

**Idea:** Only look at the **last character**.


If I see `a`, what usually comes next?

- In `aabid`: `a` -> `a`, `a` -> `b`. (50% each)

This is a **Bigram** (2-gram) model.

# Visualizing Bigram Probabilities

We can build a huge lookup table by just counting.



# Generating with Bigrams

1. Start with `.` (start token).
2. Look at row `.`. High prob for `a`, `z`.
3. Roll dice -> Get `a`.
4. Look at row `a`.
5. Roll dice -> Get `b`.
6. Repeat until `.` (end token).

**Result:** `abid`, `zel`, `aab`. (It makes decent syllables!)

# Why Counting Fails

What if we want to write a sentence?

"The quick brown fox jumps..."

If I only see `fox`, I might predict `runs` or `eats`.

I **forgot** that the sentence started with "The quick..." !

- **Bigram:** Context = 1 char. (Too small)
- **Trigram:** Context = 2 chars.
- **N-gram:** Context = N chars?

# The Curse of Dimensionality

Why not just count 10-grams? (Last 10 characters).

- **1-gram table:** 26 rows.
- **2-gram table:**  $26^2 = 676$  rows.
- **3-gram table:**  $26^3 = 17,576$  rows.
- **10-gram table:** 141,000,000,000,000 rows!

**Impossible!** We need a better way than counting. We need to **approximate**.

## Part 3: Representing Meaning (Embeddings)

# How do computers read?

Computers only understand numbers.

Option A: **One-Hot Encoding**



**Problem:** These vectors are **Orthogonal**.

Dot Product( `a` , `b` ) = 0.

The computer thinks `a` and `b` are completely different. It doesn't know they are both letters!

# Option B: Dense Embeddings

Let's represent each character as a list of **learnable attributes**.

Maybe:

- Dim 1: "Is it a vowel?"
- Dim 2: "Is it a common ending?"

a = [0.9, 0.1]


e = [0.8, 0.2]

b = [-0.5, 0.1]

Now, a and e are **mathematically close!**

# The Latent Space

If we train a model, these embeddings organize themselves automatically.



## **Part 4: The Neural Network (MLP)**

# Moving to "Deep Learning"

Instead of a lookup table, let's use a **Multi-Layer Perceptron (MLP)**.

**Goal:** Predict  $P(\text{next}|\text{context})$  using a function.

$$y = f(x; W)$$

- $x$ : Context (e.g., 3 previous characters).
- $W$ : Weights (The "Brain").
- $y$ : Probabilities.

# The Sliding Window

We create a dataset by sliding a window over text.



# The Architecture

1. **Lookup:** Get embeddings for `a`, `a`, `b`.
2. **Concat:** Merge them into one big vector.
3. **Hidden Layer:** Mix them (Feature extraction).
4. **Output:** Softmax (Convert to probability).

# The Architecture Diagram



# How does it learn? (Training)


We start with random weights. The model is stupid.

It predicts  $\hat{z}$  when the answer is  $i$ .

We calculate the **Loss** (Error).

Then we minimize it using **Gradient Descent**.

# Minimizing Loss



We "slide down the hill" to find the best weights.

This updates both the **Weights** and the **Embeddings**!

## Part 5: The "Context" Bottleneck

## The Problem with Fixed Windows

Our MLP has a fatal flaw.

It has a **Fixed Block Size** (e.g., 3 characters).

It can **never** see further back than 3 steps.

# The "Alice" Example



The model sees ...with the ?.

It has **no idea** Alice is holding a key. The key fell off the "Window".

# Attempted Solution: RNNs

**Recurrent Neural Networks (RNNs)** try to fix this by keeping a "Memory".



**Issue:** For long sentences, the memory gets "muddy".

It's like the "Telephone Game". By the end, the message is lost.

(Vanishing Gradient Problem).

## Part 6: The Revolution (Transformers)

# Enter Attention

What if, instead of a fixed window or a fading memory...

We could just **look back** at everything?

## **Attention Mechanism:**

"I am at step 10. I can 'attend' to step 1, 5, or 9 depending on what is important."

# The "Searchlight" Analogy

**Fixed Window (MLP):** Reading with tunnel vision. **RNN:** Reading while trying to memorize everything. **Attention:** Reading with a highlighter and a search engine.

When the model sees `"it"`, it asks: **"What object was mentioned before?"**

And it finds `"animal"`.

# Attention Diagram



This "connection" allows the model to handle contexts of thousands of words!

# How Attention Works ( $Q$ , $K$ , $V$ )

It uses a "Database Lookup" logic.

1. **Query ( $Q$ ):** What am I looking for? ("Holding object")
2. **Key ( $K$ ):** What does this word define? ("Alice", "Door")
3. **Value ( $V$ ):** The actual content.

**Score = Match( $Q$ ,  $K$ ).**

# QKV Visualized



In ChatGPT, this happens **billions** of times in parallel!

## Part 7: From Theory to ChatGPT

# Scaling Up

We built a character-level model.

ChatGPT is the **same thing**, just bigger.

- **Tokens:** Not characters, but sub-words.
- **Layers:** Not 1 hidden layer, but 96 layers.
- **Heads:** Not 1 attention head, but 96 heads.
- **Data:** The entire internet.

# Tokenization

We don't use characters (too slow) or words (too many).


We use **Tokens** (Sub-words).



# Sampling: The "Creativity" Knob

When ChatGPT generates, it picks from the probabilities.

We can control this with **Temperature**.



# The Sampling Tree

Because we sample, we can get different stories every time!



# Summary: The Recipe

To build a modern LLM:

1. **Input:** Massive text data.
2. **Tokenize:** Convert to numbers.
3. **Embed:** Learn meaning.
4. **Model:** Transformer (Attention) to predict next token.
5. **Train:** Minimize Cross-Entropy Loss on GPUs for months.
6. **RLHF:** Fine-tune it to be helpful (not covered today).

# Resources

1. **Andrej Karpathy:** "Neural Networks: Zero to Hero" (YouTube).
2. **Jay Alammar:** "The Illustrated Transformer".
3. **NanoGPT:** Karpathy's code repo to build this in 100 lines of Python.

# Thank You

"The best way to predict the future is to create it."

Questions?