

Object Detection Basics

Deep Learning for Computer Vision

Lecture by [Your Name]

Inspired by Andrew Ng's teaching style

What You Will Learn Today


1. **The Core Problem**: Classification vs. Detection
2. **The Tools**: Bounding Boxes, Confidence, IoU
3. **The Algorithm**: How NMS cleans up results
4. **The Architectures**: YOLO vs. Faster R-CNN
5. **The Metrics**: How do we measure success? (mAP)
6. **Practical Tips**: Training & Data Augmentation

1. Understanding the Goal

Classification answers "What is this?"

Detection answers "What is this **AND** Where is it?"


Object Detection: What & Where




The Bounding Box (BBox)

A BBox defines the object's location. It is crucial to understand the format because datasets vary!


Format 1: (x, y, w, h)
Top-Left Corner + Dimensions



Format 2: (x_l, y_l, x_r, y_r)
Two Corner Points



Format 3: (cx, cy, w, h)
Center Point + Dimensions



Format Conversions


2. Key Concept: IoU (Intersection over Union)

How do we know if a detection is "correct"? We compare it to the Ground Truth.

Intuition:

- **Intersection:** The shared area (overlap).
- **Union:** The total combined area.
- **Ratio:** 0 (No overlap) to 1 (Perfect match).

IoU Visualized



3. The Problem of Multiple Detections

Detectors often find the same object multiple times.


- "Is this a car?" -> Yes (90%)
- "Is this slightly shifted box a car?" -> Yes (85%)

We need a way to clean this up.

Non-Maximum Suppression (NMS)


The Algorithm to clean up duplicate boxes:

Non-Maximum Suppression (NMS) Flow



NMS in Action

Before NMS
Multiple Overlapping Detections



After NMS
Single Best Detection




4. How Detection Models Work

Most modern detectors follow a similar pipeline:



Two Main Approaches


Two-Stage vs. One-Stage Detectors




Deep Dive: YOLO (You Only Look Once)

YOLO divides the image into a **grid**. Each cell is responsible for detecting objects whose center falls inside it.

YOLO Grid (4×4)
Each cell predicts objects






YOLO Architecture Explained



5. Loss Functions (How the Model Learns)

The model needs to minimize errors in three things simultaneously:



6. Evaluation Metrics: mAP

Accuracy is not enough! We use **Mean Average Precision (mAP)**.

- **Precision:** When you predict a car, is it actually a car?
- **Recall:** Did you find all the cars?
- **mAP:** Area under the Precision-Recall curve, averaged over all classes.

		Confusion Matrix	
		Actual	Positive
Actual	Negative	False Negative (FN)	True Negative (TN)
	Positive	True Positive (TP)	False Positive (FP)
Predicted			
Positive	Negative		

Precision

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

"Of all detections, how many were correct?"

Example:
Model detects 100 cars
90 are actually cars (TP)
10 are not cars (FP)

$$\text{Precision} = \frac{90}{100} = 0.90 (90\%)$$

□ High precision:
Few false alarms
Important when FP is costly

Recall (Sensitivity)


$$\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

"Of all actual objects, how many did we find?"

Example:
Image has 100 cars
Model finds 85 cars (TP)
Misses 15 cars (FN)

$$\text{Recall} = \frac{85}{100} = 0.85 (85\%)$$

□ High recall:
Few missed detections
Important when FN is costly



Average Precision (AP)

AP = Area Under PR Curve


mean Average Precision (mAP)

mAP = Average of AP across all classes

7. Practical Training Tips

Training a robust detector requires more than just code.

Object Detection Training Workflow



Summary & Key Takeaways

1. **Object Detection** = Localization + Classification.
2. **IoU** measures how good a box is.
3. **NMS** removes duplicate boxes.
4. **YOLO** is fast (one-stage), **Faster R-CNN** is accurate (two-stage).
5. **mAP** is the gold standard metric.

Next Steps:

- Download a dataset (e.g., COCO128).
- Run YOLOv8 on your webcam.
- Inspect the "loss curves" to see learning in action!

Thank You

"AI is the new electricity." - Andrew Ng

Questions?

Slides generated with Marp, Matplotlib & Graphviz